

The Convex Origin of Fixed Costs[†]

Hanbaek Lee[‡]

University of Cambridge

April 2026

[\(click here for the latest version\)](#)

Abstract

Are models with fixed and convex adjustment costs different? This paper shows they are ex-ante equivalent via Fenchel duality: with fixed costs drawn from a distribution G , the threshold rule coincides with a convex cost over the adjustment probability — same value functions, same probabilities. Different shapes of G generate cost functions in menu costs, lumpy investment, rational inattention, discrete choice, and control costs, connecting these literatures in one G -indexed family. For pricing under the static-surplus approximation, a power law G indexed by γ traces a continuous spectrum between Calvo and Golosov-Lucas; when the gap distribution is supported inside the inaction band, non-neutrality decomposes exactly into frequency and a selection wedge $2\gamma\bar{\Lambda}$, and approximately so when the tail mass outside the band is negligible. Illustrative CPI calibration yields $\hat{\gamma} \approx 0.4$, with selection contributing 44% of the slope κ .

Keywords: Fixed costs, convex adjustment costs, Fenchel duality, menu costs, discrete choice, rational inattention, Phillips curve.

JEL codes: D81, E31, E22, E32.

[†]I am grateful for helpful comments and discussions from Alexander W. Bloedel, Philip Ushchev, Johannes Wieland, and Mindy Z. Xiaolan. All errors are my own.

[‡]Email: hl610@cam.ac.uk

1 Introduction

A central puzzle in macroeconomics is the gap between micro behaviour and aggregate dynamics. At the individual level, adjustment is often lumpy and discontinuous: firms change prices in sudden jumps, invest in infrequent spikes, and workers transition between labour-force states in discrete moves. Yet aggregate dynamics are smooth and gradual — inflation responds continuously to shocks, capital accumulates in humps, and labour markets absorb disruptions fluidly. How micro lumpiness aggregates into macro smoothness — and whether the micro structure still matters for aggregate conclusions — has shaped the literature for four decades. Two modelling traditions capture the tension. *Fixed* adjustment costs reproduce the micro jumps: agents stay inactive until a threshold is crossed, then act. *Convex* adjustment costs reproduce the aggregate smoothness: actions respond gradually and probability-weighted to the state. The two deliver sharply different predictions for monetary non-neutrality, the Phillips curve slope, and the welfare cost of adjustment frictions.

Are models with fixed and convex adjustment costs different in nature? This paper shows they are *ex-ante equivalent*: before the cost realisation, the two formulations are interchangeable. The key insight is that when fixed costs are heterogeneous — drawn from a distribution G — the optimal adjustment rule under G delivers the same expected payoff and the same adjustment probability as a convex cost $C(\varphi)$ over the adjustment probability, where C is derived from G via its quantile function and the payoff Π is the Fenchel conjugate of C . The realised action is binary in both formulations; what coincides is the smooth ex-ante extensive margin. The reason is economic: agents with low cost draws adjust first, so the marginal non-adjuster faces a higher cost draw than the last. This increasing marginal cost is convexity.

The paper makes three contributions. First, it establishes the ex-ante equivalence and its structural implications. The deterministic fixed cost — where every firm faces the same cost — is a knife-edge singularity; any absolutely continuous perturbation with positive density, however small, makes the induced cost over the adjustment probability strictly convex (Proposition 2). The realised action remains binary at the firm level; what becomes smooth and convex is the ex-ante extensive margin. At the aggregate level, the shape of G is the relevant dial: uniform G gives textbook quadratic costs (Khan and Thomas, 2008); concentrated G gives the Ss rule of Golosov and Lucas (2007); dispersed G approximates Calvo (1983) in the frequency-

calibrated limit of Proposition 9. These are not different kinds of models but different shapes of one distribution. This also provides a partial-equilibrium foundation for the irrelevance result of Khan and Thomas (2008): with uniform fixed costs, each firm’s ex-ante extensive margin is already convex, so aggregation is not the source of the smoothness.

Second, the equivalence connects classes of models with frictional adjustment through the shared mapping $G \mapsto C$. Different distributions G generate cost functions that coincide with well-known information-theoretic entropies (Proposition 4). Exponential G yields the generalised KL divergence from the no-adjustment prior — the cost function in the binary rational inattention model of Matějka and McKay (2015). Uniform G yields the quadratic attention penalty in Khan and Thomas (2008) and Gabaix (2014). Power law G yields a Tsallis q -divergence (Tsallis, 1988). In the multi-action extension, i.i.d. Gumbel cost shocks recover the multinomial logit of McFadden (1974) (with ξ reinterpreted as a generalised payoff shock to accommodate Gumbel’s full real-line support). Five literatures — menu costs, lumpy investment, discrete choice, rational inattention, and control costs (Mattsson and Weibull, 2002; Hofbauer and Sandholm, 2002) — are connected within the single G -indexed family; the contributions vary in degree, from genuine new connections (menu costs $\leftrightarrow$ lumpy investment via shared G) to identifications of objects already studied (discrete choice and control costs share the conjugate skeleton). The shared structure implies that the smooth first-order condition $S = C'(\varphi)$ replaces kink-based dynamic programming, and that $C(\varphi^*)$ has a concrete resource-cost interpretation under the fixed-cost primitive.

Third, the paper applies the equivalence to price setting under the standard static-surplus approximation $S(x) = Bx^2$, which is what makes the closed-form aggregation possible. The framework produces an analytical decomposition in which Calvo arises as a frequency-calibrated selection-equivalent limit and Golosov-Lucas as the opposite extreme. A power law family $G(\xi) = (\xi/\bar{\xi})^\gamma$ parametrises this spectrum through γ , the degree of cost concentration. When γ is low, menu costs are dispersed and repricing is close to random; when γ is high, only the most mispriced firms reprice and each repricing event moves the price level substantially. The aggregate price response to a nominal shock decomposes (Proposition 10) into *frequency* (how many firms reprice) and *selection* (how well repricing is targeted), with γ governing the balance. Alvarez, Le Bihan, and Lippi (2016) show that frequency and kurtosis of price changes are

sufficient statistics for non-neutrality within their model class; the present framework provides the structural layer from which those observable moments are derived, since G determines the repricing hazard and hence both moments. Illustrative calibration to CPI-based moments yields $\hat{\gamma} \approx 0.4$, with selection contributing 44% of κ . The calibrated value is one-moment matching at imposed heterogeneity ($\text{CV} = 1$); a Dominick’s scanner exercise at higher frequency does not tightly identify γ with the same moments and is consistent with the low- γ region of the spectrum.

Although the pricing application is developed in greatest depth, the duality applies to any extensive-margin decision with heterogeneous costs — firm entry, search and matching, labour force participation, technology adoption, and credit default — as discussed in Section 5.

Related literature The analysis connects several literatures. On macroeconomic frictions and pricing: [Caplin and Spulber \(1987\)](#); [Dotsey, King, and Wolman \(1999\)](#); [Goloso and Lucas \(2007\)](#); [Midrigan \(2011\)](#); [Nakamura and Steinsson \(2010\)](#); [Alvarez, Le Bihan, and Lippi \(2016\)](#); [Vavra \(2014\)](#); on lumpy investment, [Thomas \(2002\)](#); [Gourio and Kashyap \(2007\)](#); [Oskolkov and Lippi \(2026\)](#). The [Vavra \(2014\)](#) sufficient statistic is endogenous in G , which governs both the cross-sectional variance and the friction generating it. The repricing hazard $\Lambda(x)$ realises the “adjustment function” primitive of [Caballero and Engel \(1993, 1999\)](#): its shape is not assumed but derived from G , micro-founding the adjustment-function approach. [Caballero and Engel \(2007\)](#) show that the variance of idiosyncratic shocks relative to the inaction band governs amplification; Proposition 12 states an analogous monotonicity in γ .

The Williams-Daly-Zachary theorem ([McFadden, 1974, 1978, 1981](#)) establishes the same algebraic skeleton in discrete choice; the contributions here — the G -indexed taxonomy, the Khan-Thomas foundation, the pricing spectrum — are absent there. The fixed cost model is a control cost model ([Mattsson and Weibull, 2002](#); [Hofbauer and Sandholm, 2002](#); [Gabaix, 2014](#)) with penalty endogenously determined by G ; binary rational inattention ([Sims, 2003](#); [Matějka and McKay, 2015](#)) with a degenerate prior on inaction yields KL, recovered here from exponential G , and the multinomial logit’s negative Shannon entropy from i.i.d. Gumbel costs.

Roadmap Section 2 presents the duality theorem, characterises the induced cost function, establishes universality, develops the entropy correspondence, and applies the result to investment. Section 3 develops the pricing spectrum. Section 4 calibrates γ from micro price moments and quantifies the frequency-selection decomposition. Section 5 discusses further applications, interpretation, and scope. Section 6 concludes.

2 The duality

2.1 Primitives

Consider an agent who faces a binary decision: take an action (“adjust”) or maintain the status quo (“do nothing”). Adjusting yields a gross surplus $S \geq 0$ but requires paying a fixed cost $w\xi$, where $w > 0$ is a wage or unit cost and ξ is a random draw from a distribution G supported on $[0, \bar{\xi}]$ with $\bar{\xi} \in (0, \infty]$ (the case $\bar{\xi} = \infty$ covers unbounded supports such as the exponential), with density $g > 0$ on its support and finite mean $E_G[\xi] < \infty$. All integrals against dG are taken over the support of G . Not adjusting yields zero. The randomness in ξ captures heterogeneity in the cost of adjustment. This heterogeneity may be across agents (a cross-section of firms facing different menu costs) or across time (a single firm drawing a new cost each period, as in [Dotsey, King, and Wolman 1999](#)).

This setup is minimal. It strips away the details of particular applications to isolate the structure common to all of them. In the investment application, S is the surplus from investing at the optimal scale. In the pricing application of Section 3, S is the gain from closing the price gap. Because the setup is general, the duality theorem applies to any extensive-margin decision where the agent’s cost of acting is uncertain.

Definition 1 (Fixed cost formulation).

Given the payoff S , the following defines the fixed cost payoff $\Pi^{FC}(S)$ and the threshold rule ξ^ : The agent observes the realised cost ξ and adjusts if and only if $w\xi \leq S$:*

$$\Pi^{FC}(S) = \int_0^{\bar{\xi}} \max\{S - w\xi, 0\} dG(\xi) = \int_0^{\min\{S/w, \bar{\xi}\}} (S - w\xi) dG(\xi). \quad (1)$$

The agent adjusts for all cost draws below the threshold $\xi^ = S/w$ and does not adjust*

otherwise.

Definition 2 (Probability choice formulation).

Given the payoff S , the following defines the probability choice payoff $\Pi^{PC}(S)$ and the optimal adjustment probability φ^* : The agent chooses $\varphi \in [0, 1]$, paying a convex cost $C(\varphi)$:

$$\Pi^{PC}(S) = \max_{\varphi \in [0, 1]} \{\varphi \cdot S - C(\varphi)\}, \quad (2)$$

where $C : [0, 1] \rightarrow \mathbb{R}_+$ with $C(0) = 0$, C increasing and convex. The agent trades off the expected benefit of adjusting ($\varphi \cdot S$) against the cost of choosing a high adjustment probability ($C(\varphi)$).

The duality theorem establishes that these two formulations are ex-ante equivalent: same value functions and same adjustment probabilities, with C derived from G via the quantile integral. For each distribution G with a positive density, there exists a unique convex cost function C with $C(0) = 0$ that delivers this equivalence.

2.2 Main result

Theorem 1 (Convex duality of fixed costs).

The fixed cost formulation (Definition 1) and the probability choice formulation (Definition 2) are equivalent under the cost function

$$C(\varphi) = w \int_0^{G^{-1}(\varphi)} \xi dG(\xi). \quad (3)$$

Specifically:

- (i) **Payoff equivalence.** $\Pi^{FC}(S) = \Pi^{PC}(S)$ for all $S \geq 0$. The two formulations are Fenchel conjugate duals: $\Pi^{PC}(S)$ is the conjugate of C , and $C(\varphi)$ is the conjugate of Π^{FC} .
- (ii) **Policy equivalence.** The optimal adjustment probability is the same: $\varphi^* = G(S/w)$. Moreover, if the probability choice model randomises by drawing $\xi \sim G$ and adjusting whenever $\xi \leq G^{-1}(\varphi^*)$, the two models produce identical binary outcomes path by path.

Proof.

(i) **Payoff equivalence.** Substituting the threshold $\xi^* = S/w$ (assuming $S/w \leq \bar{\xi}$;

for $S/w > \bar{\xi}$, both formulations yield $\Pi(S) = S - wE[\xi]$ at $\varphi^* = 1$, and for unbounded support $\varphi^* \in (0, 1)$ for every finite S) into (1):

$$\Pi^{FC}(S) = S \cdot G(S/w) - w \int_0^{S/w} \xi dG(\xi). \quad (4)$$

The first term is the expected surplus from adjusting: the probability of adjustment $G(S/w)$ times the surplus S . The second term is the expected cost paid by adjusters. Setting $\varphi = G(S/w)$ so that $S/w = G^{-1}(\varphi)$:

$$\Pi^{FC}(S) = \varphi \cdot S - w \int_0^{G^{-1}(\varphi)} \xi dG(\xi) = \varphi S - C(\varphi).$$

This is the objective of (2) evaluated at $\varphi^* = G(S/w)$. Since $C'(\varphi) = wG^{-1}(\varphi)$ is non-decreasing, C is convex, so the first-order condition $S = C'(\varphi)$ is sufficient and $\varphi^* = G(S/w)$ attains the supremum. The conjugate relationship follows from the definition of the Fenchel conjugate: $\sup_{\varphi} \{\varphi S - C(\varphi)\} = \Pi^{PC}(S) = \Pi^{FC}(S)$.

(ii) Policy equivalence. The first-order condition of (2) is $S = C'(\varphi)$. Differentiating (3) via the fundamental theorem of calculus gives $C'(\varphi) = wG^{-1}(\varphi)$, so $G^{-1}(\varphi^*) = S/w$, i.e., $\varphi^* = G(S/w)$ — the same threshold probability as in Definition 1. For the path-by-path equivalence: in the fixed cost model, the agent draws $\xi \sim G$ and adjusts whenever $\xi \leq S/w$, which occurs with probability φ^* . If the probability choice model uses the same draw and adjusts whenever $\xi \leq G^{-1}(\varphi^*)$, the two rules coincide and the realised adjustment indicators are identical. The path-by-path coupling presupposes i.i.d. cost draws across time and agents; persistent-type heterogeneity in ξ preserves distributional but not trajectory equivalence. ■

The economic content is straightforward. In the fixed cost model, agents with low cost draws adjust and those with high draws do not. The expected payoff equals the adjustment probability times the surplus, minus the average cost paid by adjusters. In the probability choice model, the agent chooses an adjustment probability and pays a cost $C(\varphi)$. Proposition 1 below confirms that C is convex and characterises its shape. The theorem says these two descriptions yield exactly the same payoff and adjustment probability, provided C is constructed from G via (3).

The Fenchel conjugate structure places the two formulations on equal mathematical footing. The payoff function $\Pi^{PC}(S)$ and the cost function $C(\varphi)$ are conjugates of

each other: the Fenchel conjugate of C yields Π^{PC} , and the conjugate of Π^{FC} yields C . The two express the same ex-ante value function in dual variables — one in the surplus S , the other in the adjustment probability φ . Any result proved for one variable translates immediately to the other. At the ex-ante value/probability level, the two are not different theories but different parametrisations of a single theory; under different primitive interpretations of G (physical cost versus information-processing penalty), the welfare and identification implications can differ, as discussed in Section 2.6. When G is uniform, C is quadratic and the equivalence can be verified by direct calculation (see the investment application below).

Relation to discrete choice theory As noted in the introduction, the Williams-Daly-Zachary (WDZ) theorem (McFadden, 1978, 1981) establishes that the expected maximum utility under random taste shocks is the Fenchel conjugate of a generating function determined by the shock distribution. For binary choice, the WDZ identity and Theorem 1 share the same algebraic skeleton (Table F.1 in Online Appendix F). The binary conjugate identity is therefore not new; the contributions here lie in what it enables — the distribution-theoretic taxonomy, the structural foundation for Khan-Thomas, and the continuous pricing spectrum — all of which are absent from the discrete choice literature. The duality is the tool; the taxonomy and the pricing application are the results. Parallel recent work by Bloedel, Denti, and Pomatto (2025) develops a conjugate-duality analysis in the rational-inattention literature, using f -divergences to parametrise information costs over Blackwell experiments and their conjugates f^* to characterise optimal stochastic choice. Their duality runs between f and f^* on the space of information structures; the present duality runs between G (a distribution of physical fixed costs) and the induced C via the quantile integral (3). The two frameworks share Fenchel conjugacy as a tool but target different primitives and different economic questions — information acquisition versus extensive-margin adjustment frictions.

Two further distinctions are worth noting. First, the source of randomness differs: in WDZ, randomness enters payoffs via taste shocks ε_j ; here, randomness enters costs via the draw $\xi \sim G$. For binary choice the two are isomorphic under a sign change. Second, the coupling is new. WDZ establishes equivalence of expected payoffs and choice probabilities; Theorem 1 further constructs a natural randomisation device under which the realised adjustment decisions coincide path by path — not merely

in distribution.

Relation to control costs In the game theory literature, [Mattsson and Weibull \(2002\)](#) and [Hofbauer and Sandholm \(2002\)](#) study agents who choose mixed strategies subject to convex “control costs” that penalise departures from a default action. Deviating from habitual behaviour is costly, and the cost increases as the deviation becomes more extreme. The probability choice formulation (2) has exactly this structure: $C(\varphi)$ is the cost of choosing adjustment probability φ relative to the default of not adjusting ($\varphi = 0$). Theorem 1 therefore shows that the fixed cost model is a control cost model. The control cost is not an ad hoc penalty imposed by the modeller. It is the endogenous cost function determined by the distribution of frictions that the agent faces.

2.3 Properties of the induced cost function

The induced cost function $C(\varphi)$ inherits its shape entirely from the distribution G . The following proposition characterises its key properties and provides the economic interpretation that underlies all subsequent applications.

Proposition 1 (Cost function properties).

The induced cost function (3) satisfies:

- (i) $C(0) = 0$, $C(1) = w E_G[\xi]$.
- (ii) $C'(\varphi) = w G^{-1}(\varphi)$: *marginal cost equals the cost of the marginal adjuster.*
- (iii) $C''(\varphi) = w/g(G^{-1}(\varphi)) > 0$: *strictly convex whenever $g > 0$.*
- (iv) C is C^2 if g is C^1 and positive.

Proof. See Online Appendix A. ■

Part (ii) carries the core economic content. Agents are ranked by their cost draws: those with the lowest ξ adjust first. The marginal cost $C'(\varphi) = w G^{-1}(\varphi)$ is the cost of the agent just indifferent at the current threshold — the φ -th quantile of the cost distribution. Each increment in φ requires inducing a more expensive agent to adjust, which is the force that makes C convex. This is the economic origin of convexity: it

is not assumed but arises endogenously from the ranking of heterogeneous agents by their costs.

Part (iii) reveals how the density g governs the strength of this convexity. Where g is large, many agents are bunched near the threshold. A small increase in the surplus tips many of them into adjustment, so the marginal cost rises gently. Where g is small, few agents are near the threshold. The same increase tips few additional agents, and the marginal cost rises sharply. In short: crowded margins mean gentle curvature; sparse margins mean sharp curvature.

Corollary 1 (Shape of G determines curvature).

- (i) g increasing $\Rightarrow C''$ decreasing: there are relatively more agents at higher cost quantiles, so inducing additional adjustment becomes progressively easier. Adjustment is “cheap at the margin” for high φ .
- (ii) g decreasing $\Rightarrow C''$ increasing: agents thin out at higher cost quantiles, so inducing additional adjustment becomes progressively more expensive. Adjustment becomes disproportionately costly at high φ .
- (iii) g constant (uniform) $\Rightarrow C''$ constant: agents are equally spaced throughout the cost distribution, yielding the standard quadratic cost function with constant curvature.

Proof. See Online Appendix A. ■

The following corollary applies the mapping (3) to four canonical distributions, recovering well-known cost function specifications as special cases of the single formula $C(\varphi) = w \int_0^\varphi G^{-1}(u) du$. Each distribution represents a different assumption about the cross-sectional heterogeneity of adjustment costs, and the induced cost function inherits its functional form from the shape of G .

Corollary 2 (Canonical distributions).

- (i) **Uniform** $G = U[0, \bar{\xi}]$: $C(\varphi) = \frac{w\bar{\xi}}{2}\varphi^2$. The constant density produces constant curvature, so each unit increase in the adjustment probability costs the same marginal increment. This is the standard quadratic adjustment cost used throughout the macroeconomics and industrial organisation literatures.

(ii) **Power law** $G(\xi) = (\xi/\bar{\xi})^\gamma$, $\gamma > 0$: $C(\varphi) = \frac{\gamma}{\gamma+1}w\bar{\xi}\varphi^{1+1/\gamma}$. The cost function is iso-elastic with elasticity $1 + 1/\gamma$, and the shape parameter γ governs the curvature profile. For $\gamma > 1$, the density $g(\xi) = \gamma\xi^{\gamma-1}/\bar{\xi}^\gamma$ is increasing: most cost draws are high, so cheap adjusters are scarce and C is less curved than quadratic. For $\gamma < 1$, the density is decreasing: most cost draws are low, cheap adjusters are abundant, and C is more curved than quadratic. The uniform case is $\gamma = 1$. As $\gamma \rightarrow \infty$, the density concentrates near $\bar{\xi}$ and C approaches linearity (the degenerate limit). As $\gamma \rightarrow 0$, the density concentrates near zero and C approaches zero uniformly on $[0, 1]$ (free adjustment in the limit). The power law family thus spans the full range from lumpiness to frictionless adjustment, indexed by a single parameter. This family plays the central role in the pricing application of Section 3.

(iii) **Exponential** $G(\xi) = 1 - e^{-\lambda\xi}$ on $[0, \infty)$: $C'(\varphi) = -\frac{w}{\lambda}\ln(1 - \varphi)$, giving $C(\varphi) = \frac{w}{\lambda}[(1 - \varphi)\ln(1 - \varphi) + \varphi]$. The logarithmic marginal cost diverges as $\varphi \rightarrow 1$, severely penalising near-certain adjustment. This reflects the exponential distribution's unbounded support: to guarantee adjustment with probability close to one, the agent must be willing to pay arbitrarily high cost draws, which makes the marginal cost of the last increment unbounded, even though the total expected cost $C(1) = w/\lambda$ remains finite. As shown in Section 2.5, this cost function is the generalised KL divergence (I-divergence) from the no-adjustment prior, the Bregman divergence generated by $h(x) = x \ln x$.¹ It coincides with the mutual information cost in the binary rational inattention model under a degenerate default (Sims, 2003; Matějka and McKay, 2015).

(iv) **Degenerate** at $\bar{\xi}$: $C(\varphi) = w\bar{\xi}\varphi$. The cost function is linear, with no strict convexity. The only optimal choices are $\varphi = 0$ (never adjust) or $\varphi = 1$ (always adjust), depending on whether S falls below or above $w\bar{\xi}$. This is the canonical binary threshold rule of the deterministic fixed cost model — the singular case analysed in Proposition 2 below.

Proof. See Online Appendix A. ■

These are not four different models — they are four special cases of one model, distinguished solely by the shape of G .

¹Strict Shannon KL from a degenerate prior is infinite for any positive φ ; the cost here is the generalised KL / I-divergence, finite on the simplex interior.

2.4 Universality of the convex representation

The degenerate case — a single, known fixed cost — is the only setting in which the problem is not strictly convex. This section shows that this case is a knife-edge singularity: arbitrarily close perturbations of a point mass by absolutely continuous distributions with positive density restore strict convexity. The convex representation is the non-knife-edge case; the deterministic fixed cost is the singular exception.

Proposition 2 (Universality).

Let $\bar{\xi} > 0$ be a deterministic fixed cost. For any $\epsilon \in (0, \bar{\xi})$, perturb the cost to $\xi \sim G_\epsilon \equiv U[\bar{\xi} - \epsilon, \bar{\xi} + \epsilon]$. The induced cost function is

$$C_\epsilon(\varphi) = w(\bar{\xi} - \epsilon)\varphi + w\epsilon\varphi^2, \quad (5)$$

which is strictly convex with $C''_\epsilon(\varphi) = 2w\epsilon > 0$ for every $\epsilon > 0$. As $\epsilon \rightarrow 0$:

- (i) $C_\epsilon(\varphi) \rightarrow w\bar{\xi}\varphi$ pointwise: the linear (degenerate) cost.
- (ii) $C''_\epsilon(\varphi) \rightarrow 0$: curvature vanishes.
- (iii) The optimal $\varphi^*_\epsilon(S) = \max\left\{0, \min\left\{\frac{S - w(\bar{\xi} - \epsilon)}{2w\epsilon}, 1\right\}\right\}$ converges to $\mathbf{1}[S \geq w\bar{\xi}]$: the binary threshold rule.

Thus the deterministic fixed cost is the singular limit of a family of strictly convex problems.

Proof. For $G_\epsilon = U[\bar{\xi} - \epsilon, \bar{\xi} + \epsilon]$, applying (3) gives $C_\epsilon(\varphi) = w(\bar{\xi} - \epsilon)\varphi + w\epsilon\varphi^2$, with $C''_\epsilon = 2w\epsilon > 0$. The full derivation is in Online Appendix A. ■

The result recasts the standard view that convex costs are a “smooth approximation” to the “true” lumpy friction. The deterministic fixed cost is the knife-edge limit; arbitrarily close perturbations restore strict convexity. The curvature $C''_\epsilon = 2w\epsilon$ measures the “softness” of the adjustment margin: even an infinitesimally small ϵ produces a strictly convex problem.

Epistemic interpretation and continuity in G The deterministic fixed cost requires the agent to know the adjustment cost with infinite precision — an idealisation rarely satisfied in practice (staff time, managerial attention, installation logistics).

Any ϵ -band of noise makes the decision problem smooth and convex. The deterministic case is fragile under perturbation; the adjustment-probability representation varies continuously under standard regularity on G (g continuous and bounded away from zero on the relevant range). The set of distributions generating strictly convex C is dense under the weak topology: every neighbourhood of a point mass contains absolutely continuous distributions with strictly convex induced costs (Online Appendix A).

Mean-preserving spreads and convexity Proposition 2 shows that any perturbation of a point mass restores strict convexity. The next result records the integrated quantile order implied by MPS.

Proposition 3 (Dispersion and convexity).

Let G_1 and G_2 be distributions on $[0, \bar{\xi}]$ with the same mean and let G_2 be a mean-preserving spread of G_1 . The induced costs satisfy:

- (i) $C_1(1) = C_2(1) = w E[\xi]$.
- (ii) $C_2(\varphi) \leq C_1(\varphi)$ for all $\varphi \in [0, 1]$: greater dispersion lowers the cost function in the interior, via the integrated quantile order $\int_0^t G_2^{-1} \leq \int_0^t G_1^{-1}$ characterising MPS.
- (iii) Within the power law family, the marginal cost schedules $C'_\gamma(\varphi) = w\bar{\xi} \varphi^{1/\gamma}$ rotate around a single $\varphi_0 \in (0, 1)$ as γ varies. Outside parametric families, multiple crossings of $C'_2 - C'_1$ are possible without an additional single-crossing assumption.

The degenerate distribution maximises C on $(0, 1)$ and is the unique case in which C is linear; any non-degenerate G with positive density makes C strictly convex.

Proof. See Online Appendix A. ■

2.5 Fixed costs as generalised entropies

The cost function $C(\varphi)$ is fully determined by G and inherits its shape from the quantile function G^{-1} . Canonical distributions G yield cost functions that coincide with well-known information-theoretic entropies, revealing that fixed cost models in

macroeconomics, discrete choice, rational inattention, and sparse decision-making share the same duality viewed through different G .

Proposition 4 (Entropy correspondence).

(i) **Exponential distribution:** $G(\xi) = 1 - e^{-\lambda\xi}$ on $[0, \infty)$. Then

$$C'(\varphi) = -\frac{w}{\lambda} \ln(1 - \varphi), \quad C(\varphi) = \frac{w}{\lambda} [(1 - \varphi) \ln(1 - \varphi) + \varphi].$$

The adjustment probability under exponential costs is $\varphi^* = 1 - e^{-\lambda S/w}$, the CDF of the exponential distribution evaluated at S/w . Note that this is not the logistic function $1/(1 + e^{-S/\sigma})$ that arises in logit discrete choice; the two share a sigmoid shape but differ in functional form. The true logit correspondence arises in the multi-action extension (Section 2.7), where i.i.d. Gumbel-distributed costs across alternatives yield the multinomial logit.

(ii) **Power law distribution:** $G(\xi) = (\xi/\bar{\xi})^\gamma$ on $[0, \bar{\xi}]$. Then

$$C(\varphi) = \frac{\gamma}{\gamma + 1} w \bar{\xi} \varphi^{1+1/\gamma}.$$

This is proportional to φ^q with $q = 1 + 1/\gamma$, which shares the power-law exponent with the Tsallis q -divergence from the no-adjustment prior (Tsallis, 1988).² The parameter γ controls q : $\gamma = 1$ gives $q = 2$ (quadratic). As $\gamma \rightarrow \infty$, $q \rightarrow 1$ and $C(\varphi) \rightarrow w \bar{\xi} \varphi$, a linear function. Recovering the logarithmic (Shannon) form requires the exponential distribution, not the power law limit.

(iii) **Uniform distribution:** $G = U[0, \bar{\xi}]$. Then $C(\varphi) = \frac{w\bar{\xi}}{2} \varphi^2$, a quadratic penalty.

This is the simplest and most widely used convex cost specification, and it coincides with the L^2 penalty that appears in many regularisation frameworks.

Proof. See Online Appendix A. ■

²Bloedel, Denti, and Pomatto (2025) independently identify Csiszár- φ costs with power-law shape φ^q as a natural generalisation of mutual information in the rational-inattention setting, arriving at the same family of cost functions from the opposite side of the duality: their primitive is the convex transformation φ on likelihood ratios, while here the primitive is the power-law distribution $G(\xi) = (\xi/\bar{\xi})^\gamma$ of physical fixed costs. The identification $q = 1 + 1/\gamma$ connects the two.

The correspondences are of functional form, not of economic mechanism: exponential G yields the KL divergence used in binary rational inattention (Sims, 2003; Matějka and McKay, 2015); uniform G yields the quadratic attention penalty of Gabaix (2014); in the multi-action extension (Section 2.7), i.i.d. Gumbel costs yield negative Shannon entropy and the multinomial logit. The reduced-form $C(\varphi)$ is the same; welfare and identification implications, taken up in Section 2.6, differ.

The distribution G as a Rosetta Stone Table 1 summarises the mapping. Five literatures — menu costs, lumpy investment, discrete choice, rational inattention, and control costs — are unified by a single object: the distribution G . Each literature has independently developed models that are special cases of the duality, distinguished only by the distributional assumption on costs.

Table 1: The distribution G as a Rosetta Stone

Distribution G	Cost $C(\varphi)$	Form	Entropy	Literature
Exponential	$(1-\varphi)\ln(1-\varphi) + \varphi$	Log. MC	KL divergence	Rational inattention
Power law (γ)	$\varphi^{1+1/\gamma}$	Iso-elastic	Tsallis q -div.	Generalised discrete choice
Uniform ($\gamma=1$)	φ^2	Quadratic	L^2	Sparsity
Degenerate	$\xi \varphi$ on $\{0, 1\}$	Linear	—	Lumpy adjustment
Type I EV (multinomial)	$\sum \varphi_j \ln \varphi_j$	Neg. entropy	Shannon (mv)	Logit discrete choice

Notes: Each row specifies a distribution G of fixed costs. The second column gives the induced convex cost function $C(\varphi)$ via equation (3) (up to scale). The fourth column identifies the information-theoretic entropy to which C corresponds. The last column indicates the literature where this cost function appears. “Log. MC” = logarithmic marginal cost. “mv” = multivariate. The first four rows are binary choice; the last row is the multi-action extension (Proposition 5).

2.6 Welfare content of the duality

The welfare cost of adjustment at a given equilibrium is $C(\varphi^*)$. The identity has two uses: a smooth closed-form welfare measure for fixed cost models, and a physical-cost interpretation for the abstract penalty functions of rational inattention and discrete choice.

From fixed to convex: smooth computation and welfare analysis The probability choice formulation replaces $\max\{S - w\xi, 0\}$ with the concave objective $\max_{\varphi}\{\varphi S - C(\varphi)\}$, with smooth first-order condition $S = C'(\varphi) = wG^{-1}(\varphi)$ and

policy function $\varphi^* = G(S/w)$. The value function is C^2 when g is C^1 , so standard perturbation methods apply without special treatment of the threshold; the convex formulation is directly pluggable into linearisation-based heterogeneous-agent toolkits (Reiter, 2009; Auclert, Rognlie, and Straub, 2020). For the power law family, all aggregate objects follow in closed form: $\Lambda(x) = (Bx^2/(w\bar{\xi}))^\gamma$, $\kappa = (1 + 2\gamma)\bar{\Lambda}$ (exact when f_{ss} is supported in the interior of the inaction band; approximate otherwise — Proposition 12), and the welfare cost $C(\varphi^*)$. The total welfare cost at a given equilibrium is

$$\text{Welfare cost} = w \int_0^{G^{-1}(\varphi^*)} \xi dG(\xi) = C(\varphi^*), \quad (6)$$

a compact representation given φ^* . Note that φ^* itself still depends on G through the full equilibrium, so (6) summarises welfare without shortcutting the equilibrium computation.

From convex to fixed: welfare interpretation of information costs In rational inattention, the mutual-information cost $\lambda \sum \varphi_j \ln \varphi_j$ is measured in nats; λ is the shadow price of attention with no independent resource interpretation. Under the fixed-cost primitive in the binary exponential- G case, $C(\varphi^*)$ equals the expected physical cost paid by adjusters — and λ becomes the rate parameter of the cost distribution, in principle estimable from micro data. The resource-cost reading applies only under non-negative physical-cost primitives; the Gumbel/multi-action logit case (Proposition 5(ii)) is a reduced-form conjugacy with ξ a real-valued payoff shock and does not carry the resource interpretation. Bloedel, Denti, and Pomatto (2025) independently emphasise that the tractable object is the conjugate f^* , not f ; the present duality complements this with a primitive structural reading of the conjugate as expected physical cost via the quantile integral (3).

2.7 Beyond binary choice

The duality extends to multi-action settings via the same conjugate construction. Suppose the agent chooses among $\{a_0, a_1, \dots, a_J\}$ ($a_0 = \text{inaction}$), with a_j ($j \geq 1$) yielding surplus S_j at random cost $w\xi_j$, $\xi_j \sim G_j$ independent across actions, and selects the action with the highest net payoff.

Proposition 5 (Multi-action duality).

Let $J \geq 1$ and suppose $\xi_j \sim G_j$ independently for $j = 1, \dots, J$, with each G_j absolutely continuous with positive density on $[0, \bar{\xi}_j]$. Define $S_0 = 0$, $\xi_0 = 0$. Then:

(i) The expected payoff $\Pi^{FC}(\mathbf{S}) = E[\max_{j \in \{0, \dots, J\}} \{S_j - w\xi_j\}]$ equals

$$\Pi^{PC}(\mathbf{S}) = \max_{\boldsymbol{\varphi} \in \Delta^J} \sum_{j=0}^J \varphi_j S_j - C(\boldsymbol{\varphi}), \quad (7)$$

where $C : \Delta^J \rightarrow \mathbb{R}$ is jointly convex and determined by the joint distribution of $(\xi_1, \dots, \xi_J)$, which under independence is pinned down by the marginals $(G_1, \dots, G_J)$.

(ii) (Payoff-shock analogue.) Reinterpret ξ as a real-valued payoff shock rather than a non-negative physical cost (so $\eta_j \equiv -w\xi_j$ plays the role of a Gumbel taste shock in the standard random-utility setup). When $\eta_j \stackrel{i.i.d.}{\sim} \text{Gumbel}(0, \sigma)$ for all $j \geq 0$, the conjugate construction yields $C(\boldsymbol{\varphi}) = \sigma \sum_{j=0}^J \varphi_j \ln \varphi_j \leq 0$ (negative Shannon entropy on Δ^J). The optimal choice probabilities are

$$\varphi_j^* = \frac{e^{S_j/\sigma}}{\sum_{k=0}^J e^{S_k/\sigma}}, \quad (8)$$

the multinomial logit of [McFadden \(1974\)](#).

Proof.

Apply the Fenchel conjugate construction to the multi-alternative maximum and specialise to the Gumbel distribution. The detailed proof is given in Online Appendix A. ■

When the surplus depends on a state x , the duality applies state by state: $\varphi^*(x) = G(S(x)/w)$ is the endogenous state-dependent hazard. In dynamic settings, the probability choice formulation makes the Bellman equation smooth, while the fixed cost formulation has kinks at adjustment thresholds.

2.8 Application to investment

The duality has a direct implication for the lumpy investment literature. In the [Thomas \(2002\)](#) model, [Gourio and Kashyap \(2007\)](#) show that the shape of the fixed cost distribution is the key determinant of whether lumpy micro-level investment

spikes matter for aggregates. The duality makes this precise: each G maps to a unique convex cost C , and the shape of G governs the curvature of C and hence the aggregate investment elasticity. Under uniform G , C is quadratic (Corollary 2), so each firm’s investment problem is individually convex — providing a partial-equilibrium foundation for the irrelevance result of [Khan and Thomas \(2008\)](#). The GE channel (endogenous factor price adjustment dampening investment spikes) provides another route to smoothness; the PE convexity established here provides an independent sufficient condition that does not rely on GE feedbacks. More broadly, any G with a positive density delivers smoothness at the individual level. The dispersion of G provides a continuous index of lumpiness: concentrated G (high γ) means costs cluster near $\bar{\xi}$ (similar across firms), so adjustment occurs in sharp spikes when surpluses cross the common threshold; dispersed G (low γ) spreads costs widely, producing smoother adjustment. As dispersion shrinks to a point mass, the model recovers the binary threshold rule. [Oskolkov and Lippi \(2026\)](#) recover G from plant-level investment data, confirming that the shape of the cost distribution varies across sectors. Formal conditions for smoothness and the investment elasticity formula are in Online Appendix B.

3 The menu cost–Calvo spectrum

The deepest application of the duality is to price setting, where the distinction between time-dependent (Calvo) and state-dependent (menu cost) pricing is central to New Keynesian economics. The idea that these models sit on a spectrum is not new: [Dotsey, King, and Wolman \(1999\)](#) draw menu costs from a distribution, and [Caballero and Engel \(1993\)](#) parametrise the adjustment hazard directly. Without a tractable parametric family, exploring this spectrum requires solving a dynamic program numerically for each candidate G . What the duality plus the power-law parametrisation add is a closed-form analytical decomposition: the repricing hazard $\Lambda(x)$ is derived from G rather than assumed, and the frequency-selection decomposition is indexed by a single parameter. The distribution G governs the repricing hazard, the selection effect, the Phillips curve slope, and the degree of monetary non-neutrality.

3.1 Setup

Consider a monopolistically competitive firm with current log price p and desired log price p^* . The *price gap* is $x \equiv p^* - p$. When the firm reprices (closing the gap to zero), it gains a surplus that depends on the distance from the optimum. Based on the standard second-order approximation of the profit function $S(x)$ around $p = p^*$, I assume

$$S(x) = Bx^2, \quad (9)$$

where $B > 0$ depends on the demand elasticity and the discount factor.³ The quadratic form reflects the approximate symmetry of profits around the optimum. The loss from mispricing grows with the square of the price gap: a firm with a 2% gap loses roughly four times as much as one with a 1% gap.

The firm draws a menu cost $\xi \sim G$ each period and reprices if and only if the gain from repricing exceeds the cost $w\xi$, where $w > 0$ is the wage (as in Section 2):

$$\text{Reprice if and only if } w\xi \leq Bx^2. \quad (10)$$

Remark 1 (Static surplus approximation). *In the full dynamic problem, the surplus from repricing is $V^*(0) - V^*(x)$, where V^* is the firm's value function. This depends on the option value of waiting, which in turn depends on G and the distribution of future shocks. The quadratic approximation $S(x) = Bx^2$ is first-order accurate when idiosyncratic shocks are small relative to the inaction band: the option value correction is $O(\sigma_z^2)$ relative to Bx^2 (Alvarez, Le Bihan, and Lippi, 2016). This is the same approximation used in the sufficient-statistic literature and is standard in the analytical menu cost tradition. The duality itself (Theorem 1) holds for any surplus function S ; the quadratic specialisation is adopted here for tractability, not as a limitation of the duality. The closed-form decomposition $\kappa = (1 + 2\gamma)\bar{\Lambda}$ (Corollary 5) depends essentially on the quadratic surplus, however: option-value corrections of order $\sigma_z^2/(B\bar{x}^2)$ propagate into the aggregate decomposition. The 44% selection share reported in Section 4 is therefore the static-surplus value; a numerical comparison to the dynamic model in the spirit of Golosov and Lucas (2007) is left for future work.*

³This is standard; see Golosov and Lucas (2007) and Alvarez, Le Bihan, and Lippi (2016). The quadratic loss approximation is $S(x) \approx \frac{\varepsilon}{2}x^2$ where ε is the demand elasticity, up to discounting.

3.2 The repricing hazard

By Theorem 1, the firm’s repricing decision under the fixed cost formulation (Definition 1) is equivalent to the probability choice formulation (Definition 2) with convex cost $C(\varphi)$. The first-order condition $Bx^2 = C'(\varphi) = wG^{-1}(\varphi)$ yields the optimal repricing probability as a function of the price gap:

Definition 3 (Repricing hazard).

The repricing hazard — the probability that a firm with price gap x adjusts its price — is

$$\Lambda(x) \equiv G\left(\frac{Bx^2}{w}\right). \quad (11)$$

This is the central object of the pricing application. It is not a numerical output of a dynamic program; it is an explicit formula, given directly by the duality and the choice of G .⁴ The repricing hazard $\Lambda(x)$ summarises everything about the firm’s pricing behaviour that matters for aggregate dynamics. Conditional on (B, w) and the steady-state gap distribution f_{ss} , all aggregate pricing implications — inflation, the Phillips curve slope, the selection effect, monetary non-neutrality — flow through Λ . The distribution G is therefore the structural primitive from which aggregate pricing dynamics are derived.

Proposition 6 (Hazard properties).

The repricing hazard $\Lambda(x)$ satisfies:

- (i) $\Lambda(0) = G(0) = 0$: *a firm at its optimal price never reprices. The gain from doing so is zero, and every cost draw is positive.*
- (ii) Λ *is symmetric in x and increasing in $|x|$: the repricing probability depends only on the magnitude of the gap, not its sign. The surplus Bx^2 depends only on x^2 . Larger gaps produce larger surpluses, which exceed more cost draws.*
- (iii) $\Lambda'(x) = g(Bx^2/w) \cdot 2Bx/w$: *the sensitivity depends on the density g evaluated at the firm’s surplus. Where g is high, many firms are near the threshold, and a small change in the gap tips many across it. Where g is low, few firms are near the threshold, and the same change has little effect.*

⁴The repricing hazard $\Lambda(x)$ is the “adjustment function” in the language of Caballero and Engel (1993, 1999). Those papers take the adjustment function as a primitive and characterise how its shape determines aggregate dynamics. The present contribution is to derive $\Lambda(x)$ from first principles as a function of the menu cost distribution G , establishing G as the deeper structural primitive.

(iv) If G has bounded support $[0, \bar{\xi}]$, then $\Lambda(x) = 1$ for $|x| \geq \sqrt{w\bar{\xi}/B}$: when the gap is sufficiently large, the surplus exceeds even the highest cost draw, and the firm reprices with certainty. This defines the “inaction band” beyond which all firms adjust. If G has unbounded support (e.g., exponential), $\Lambda(x) < 1$ for all finite x and $\Lambda(x) \rightarrow 1$ only as $|x| \rightarrow \infty$.

Together, these properties paint a simple picture of the repricing decision. A firm that is well priced ($x \approx 0$) has little to gain from repricing and almost never does. As its price drifts further from the optimum, the gain from repricing grows and exceeds more cost draws, so the repricing probability rises smoothly. Eventually, for a sufficiently large gap, the gain exceeds every possible cost draw, and the firm reprices with certainty. The shape of this transition — how quickly the probability rises from zero to one — is governed entirely by the density g of menu costs. A dense region of g near a given surplus level means many firms are close to indifferent there, producing a steep slope in Λ : a small change in the gap tips many across the threshold. A sparse region produces a gentle transition. This is the microeconomic mechanism through which G shapes aggregate price dynamics.

3.3 Stationary distribution of price gaps

To close the model, the cross-sectional density f_{ss} of price gaps must be characterised. Let z denote the idiosyncratic desired-price shock with symmetric density ϕ_z , and assume zero steady-state inflation: a firm at gap x adjusts with probability $\Lambda(x)$ and resets to $z' \sim \phi_z$; otherwise it drifts to $x + z'$.

Proposition 7 (Stationary distribution).

Under $\Lambda(x) = G(Bx^2/w)$ and symmetric ϕ_z , the stationary density f_{ss} satisfies the fixed-point equation

$$f_{ss}(x) = \underbrace{\bar{\Lambda} \cdot \phi_z(x)}_{\text{reset flow}} + \underbrace{\int (1 - \Lambda(\tilde{x})) \cdot \phi_z(x - \tilde{x}) \cdot f_{ss}(\tilde{x}) d\tilde{x}}_{\text{diffusion flow}}, \quad (12)$$

with $\bar{\Lambda} = \int \Lambda(x) f_{ss}(x) dx$, and f_{ss} is symmetric around zero. The proof is in Online Appendix C.

The shape of G governs the balance between the two flows: concentrated G (steep hazard) yields a tight cross-section dominated by resets, while dispersed G (flat haz-

ard) yields a diffuse cross-section. With the hazard and the stationary distribution in hand, the next question is how the composition of repricing firms shapes aggregate outcomes.

3.4 The selection effect

Under state-dependent pricing, firms with the largest price gaps are most likely to reprice, amplifying the aggregate price response relative to Calvo (Goloso and Lucas, 2007). The duality provides a precise characterisation of how G governs this selection. I first define a measure of selection intensity, then show how it depends on G .

Definition 4 (Selection intensity).

Given a cross-sectional density $f(x)$ of price gaps, the average absolute adjusted price change is defined as

$$\overline{|x|}^{adj} := \frac{\int |x| \cdot \Lambda(x) \cdot f(x) dx}{\int \Lambda(x) \cdot f(x) dx} = \frac{E[\Lambda(x) \cdot |x|]}{E[\Lambda(x)]}. \quad (13)$$

This is the mean absolute price change among repricing firms. The selection intensity is the proportional excess of the average adjuster's absolute price change over the population average:

$$\mathcal{S} \equiv \frac{\overline{|x|}^{adj}}{E[|x|]} - 1. \quad (14)$$

When Λ is constant (Calvo), $\mathcal{S} = 0$: every firm is equally likely to reprice, so adjusters look just like the population. When Λ is a step function (Goloso-Lucas), $\mathcal{S}$ is at its polar value within the power law family at fixed $\bar{\Lambda}$: only firms at the boundary of the inaction band reprice, so every adjuster has a large gap. In between, selection intensity is governed by how sharply Λ increases with $|x|$. The next proposition makes the link to G precise: since $\Lambda(x) = G(Bx^2/w)$, the shape of the cost distribution directly determines how strong selection is.

Proposition 8 (Selection and the menu cost distribution).

Fix the cross-sectional density f . The adjustment-weighted mean $\overline{|x|}^{adj}$ admits the decomposition

$$\overline{|x|}^{adj} = E[|x|] + \frac{\text{Cov}(\Lambda(x), |x|)}{\bar{\Lambda}},$$

so that selection intensity $\mathcal{S}$ has the same sign as $\text{Cov}(\Lambda(x), |x|)$. As a consequence:

- (i) $\mathcal{S} = 0$ in the Calvo limit (constant Λ), since the covariance vanishes.
- (ii) $\mathcal{S} > 0$ whenever Λ is strictly increasing in $|x|$. The size of $\mathcal{S}$ scales with the local elasticity $\partial \ln \Lambda / \partial \ln |x| = 2(g/G)(Bx^2/w) \cdot (Bx^2/w)$, i.e., twice the hazard rate of G at the firm's surplus.
- (iii) In the Golosov-Lucas limit, Λ is a step function and adjusters are concentrated at the inaction-band boundary; this is the polar case of the spectrum, with $\mathcal{S}$ the largest attained within the power law family at fixed $\bar{\Lambda}$. A general extremality statement across all hazards calibrated to the same average frequency would require a rearrangement-style argument not pursued here.

The proof is given in Online Appendix C.

Proof. See Online Appendix C. ■

A concentrated G (high density near $\bar{\xi}$) means most firms face large menu costs. Only firms with very large gaps find it worthwhile to reprice, so the hazard is steep and selection is strong. A dispersed G (density spread across $[0, \bar{\xi}]$) means many firms face small costs. Even firms with small gaps reprice frequently, the hazard is flat, and selection is weak. The density g at the threshold is the key: it determines how many additional firms are tipped into adjustment by a small increase in the gap.

3.5 The Calvo–Golosov-Lucas spectrum

The results so far hold for any G with a positive density. To make the framework quantitatively useful, one needs a tractable parametric family that can span the range of pricing behaviour observed in the data. The two canonical models of price rigidity in the New Keynesian literature — Calvo and Golosov-Lucas — correspond to polar extreme shapes of the hazard $\Lambda(x)$. The goal of this subsection is to connect them through a single continuous family indexed by one parameter.

Calvo (1983) In the Calvo (1983) model, each firm reprices with a fixed probability $\lambda \in (0, 1)$ regardless of its price gap. Since $\Lambda(x) = G(Bx^2/w)$ is always increasing in $|x|$, a strictly constant hazard cannot arise from any non-degenerate absolutely continuous G . The framework therefore does not literally nest Calvo; the following proposition makes the limiting sense precise.

Proposition 9 (Calvo approximation).

For the power law family $G(\xi) = (\xi/\bar{\xi})^\gamma$ with $\bar{\xi}(\gamma)$ calibrated to maintain a fixed average frequency λ , and assuming f_{ss} is supported in the interior of the inaction band, the selection wedge $\kappa - \lambda = 2\gamma\lambda$ vanishes as $\gamma \rightarrow 0^+$. In this limit:

- (i) The Phillips curve slope converges to the Calvo slope: $\kappa \rightarrow \lambda$.
- (ii) Non-neutrality converges to the Calvo level: $1 - \kappa \rightarrow 1 - \lambda$.

The hazard $\Lambda(x) = (Bx^2/(w\bar{\xi}))^\gamma$ does not converge to a constant as $\gamma \rightarrow 0^+$ with fixed $\bar{\xi}$: it converges to near-certain adjustment for any $x \neq 0$ (since $a^\gamma \rightarrow 1$ for any $a > 0$). The Calvo approximation arises only through the frequency-calibrated limit, in which $\bar{\xi} \rightarrow \infty$ as $\gamma \rightarrow 0^+$ so that the hazard remains small and approximately linear in x^2 over the support of f_{ss} . In this regime, the hazard is state-dependent but its variation is small: selection contributes $2\gamma\lambda$ to the Phillips curve slope, vanishing relative to the frequency component λ .

Proof. See Online Appendix C. ■

Golosov and Lucas (2007): step-function hazard When G is degenerate at $\bar{\xi}$, $\Lambda(x) = \mathbf{1}[Bx^2 \geq w\bar{\xi}]$ — the canonical Ss rule with inaction band $\bar{x} = \sqrt{w\bar{\xi}/B}$. This is the polar case of the spectrum: the step-function hazard generates the largest selection within the power law family at fixed $\bar{\Lambda}$ and dramatically reduces non-neutrality relative to Calvo.

The power law family The power law $G(\xi) = (\xi/\bar{\xi})^\gamma$, $\gamma > 0$, provides a tractable one-parameter family that spans a selection spectrum whose low- γ limit is Calvo-equivalent in the frequency-calibrated sense (Proposition 9) and whose high- γ limit is Golosov-Lucas. The parameter γ is the degree of cost concentration. When γ is low, the density g is decreasing: most firms draw small menu costs, adjustment is cheap, and the hazard rises gently with the gap — close to Calvo. When γ is high, the density concentrates near $\bar{\xi}$: most firms draw large menu costs, only firms with very large gaps find it worthwhile to reprice, and the hazard rises steeply — close to Golosov-Lucas. The induced hazard and cost function are:

$$\Lambda(x) = \left(\frac{Bx^2}{w\bar{\xi}} \right)^\gamma, \quad C(\varphi) = \frac{\gamma}{\gamma + 1} w\bar{\xi} \varphi^{1+1/\gamma}. \quad (15)$$

The hazard expression applies for $|x| \leq \sqrt{w\xi/B}$; for larger gaps $\Lambda(x) = 1$, defining the inaction band boundary. Λ is continuous but not C^1 at $|x| = \sqrt{w\xi/B}$; boundary sets have measure zero under absolutely continuous f_{ss} . This family nests the uniform distribution ($\gamma = 1$) and extends to right-skewed cost distributions ($\gamma < 1$).

Example 1 (Power law menu costs).

For $G(\xi) = (\xi/\bar{\xi})^\gamma$, the following relations hold:

- (i) $G^{-1}(\varphi) = \bar{\xi}\varphi^{1/\gamma}$.
- (ii) $C(\varphi) = \frac{\gamma}{\gamma+1}w\bar{\xi}\varphi^{(\gamma+1)/\gamma}$.
- (iii) $C'(\varphi) = w\bar{\xi}\varphi^{1/\gamma}$, $C''(\varphi) = \frac{w\bar{\xi}}{\gamma}\varphi^{(1-\gamma)/\gamma}$.
- (iv) $\Lambda(x) = \left(\frac{Bx^2}{w\xi}\right)^\gamma$.

The parameter γ governs the entire spectrum: $\gamma = 1$ gives uniform costs and quadratic C (Dotsey, King, and Wolman, 1999; Khan and Thomas, 2008); $\gamma \rightarrow \infty$ gives the step-function hazard of Golosov and Lucas (2007); $\gamma < 1$ gives right-skewed costs with a concave hazard and weak selection; $\gamma \rightarrow 0$ gives the “cheap adjustment” limit where nearly all firms reprice. Note that the $\gamma \rightarrow 0$ limit with fixed $\bar{\xi}$ yields near-certain adjustment, not a constant interior hazard. The Calvo approximation requires the frequency-calibrated limit of Proposition 9, which adjusts $\bar{\xi}$ with γ to hold $\bar{\Lambda}$ fixed.

3.6 Aggregate inflation and the Phillips curve

The repricing hazard $\Lambda(x)$ governs individual pricing behaviour. Because the duality gives $\Lambda(x) = G(Bx^2/w)$ in closed form, aggregate inflation can be written directly in terms of G — enabling analytical decompositions that would require numerical computation without the duality. Let $f_t(x)$ denote the cross-sectional density of price gaps at time t . Aggregate inflation is

$$\pi_t = E_{f_t}[\Lambda(x_t) \cdot x_t] = \int \Lambda(x) x f_t(x) dx. \quad (16)$$

Each firm contributes in proportion to $\Lambda(x) \cdot x$. Firms with large gaps and high repricing probabilities contribute the most; firms near the optimum contribute almost nothing.

To understand how aggregate inflation responds to nominal disturbances, consider a small aggregate shock μ_t that shifts the desired price of all firms by μ_t , moving the cross-sectional distribution of gaps from $f_{ss}(x)$ (the zero-inflation steady state) to $f_{ss}(x - \mu_t)$.

Proposition 10 (Nonlinear Phillips curve).

Under the quadratic surplus approximation $S(x) = Bx^2$ and menu cost distribution G , aggregate inflation satisfies

$$\pi_t = E_t \left[G \left(\frac{Bx_t^2}{w} \right) \cdot x_t \right]. \quad (17)$$

For small aggregate disturbances around a zero-inflation steady state with gap distribution f_{ss} :

$$\pi_t \approx \kappa(G) \cdot \mu_t, \quad (18)$$

where the distribution-dependent slope is

$$\kappa(G) := \underbrace{\int G \left(\frac{Bx^2}{w} \right) f_{ss}(x) dx}_{\text{extensive margin (frequency)}} + \underbrace{\int g \left(\frac{Bx^2}{w} \right) \frac{2Bx^2}{w} f_{ss}(x) dx}_{\text{intensive margin (selection)}}. \quad (19)$$

Proof.

Write $f_t(x) = f_{ss}(x - \mu_t)$ and differentiate $\pi_t = \int G(Bx^2/w) x f_{ss}(x - \mu_t) dx$ with respect to μ_t at $\mu_t = 0$, using $\partial_\mu f_{ss}(x - \mu) = -\partial_x f_{ss}(x - \mu)$. Integrating by parts (assuming boundary terms vanish) and expanding $\frac{d}{dx}[G(Bx^2/w) \cdot x]$ yields the two terms: the first from G itself (extensive margin) and the second from $g \cdot 2Bx^2/w$ (intensive margin). The full derivation is in Online Appendix C. ■

The extensive margin (first term) captures the average repricing frequency — the only margin under Calvo. The intensive margin (second term) captures selection: a nominal shock tips firms near the threshold into adjustment, amplifying the price response. It is absent under Calvo and dominant under Golosov-Lucas. The slope $\kappa(G)$ is derived under a partial-equilibrium one-time shift of f_{ss} . GE feedbacks — strategic complementarities, endogenous wages (which enter the hazard via $\Lambda(x) = G(Bx^2/w)$), and non-translate evolution of f_t after a shock — modify the quantitative slope; the present results characterise the PE component, complementing the GE treatment in [Alvarez, Le Bihan, and Lippi \(2016\)](#).

The slope $\kappa(G)$ varies across the spectrum because G governs the relative strength of these two margins. The following corollary traces out the implications for the canonical cases.

Corollary 3 (How G governs the Phillips curve).

- (i) **Calvo limit:** When G is highly dispersed, g is small across the region where f_{ss} has mass, and the intensive margin is negligible. The slope approaches $\kappa \approx \bar{\Lambda}$, approximating the textbook New Keynesian Phillips Curve.
- (ii) **Golosov-Lucas:** When G is degenerate, $G(Bx^2/w) = \mathbf{1}[|x| \geq \bar{x}]$ and $g = \delta_{\bar{x}}$. The intensive margin dominates: inflation is driven by selection of firms at the threshold. The slope is very steep, implying low non-neutrality.
- (iii) **Uniform:** When $G = U[0, \bar{\xi}]$, the slope simplifies to

$$\kappa^U = \frac{3B}{w\bar{\xi}} E_{ss}[x^2], \quad (20)$$

which equals $\frac{3B}{w\bar{\xi}} \text{Var}_{ss}(x)$ at the zero-inflation symmetric steady state (where $E_{ss}[x] = 0$). The slope depends on the steady-state dispersion of price gaps. Economies with more dispersed gaps have steeper Phillips curves, because the selection margin is stronger when more firms have large gaps near the threshold.

- (iv) **Power law γ :** The slope κ is strictly increasing in γ . Higher γ means more concentrated menu costs, a steeper hazard, stronger selection, and a steeper Phillips curve.

3.7 Non-neutrality decomposition

The degree of monetary non-neutrality is inversely related to the price flexibility index κ . When κ is high, prices adjust quickly and real effects are small. When κ is low, prices are sticky and real effects are large. Because the duality expresses both κ and $\bar{\Lambda}$ as explicit functions of G and g , it yields a clean decomposition of κ into two forces. For the power law family, this decomposition is available in closed form.

Proposition 11 (Frequency vs. selection).

The on-impact real effect of a monetary shock dm is $(1 - \kappa(G)) dm$, well-defined

provided $\kappa(G) \leq 1$; more generally, cumulative output is inversely related to $\kappa(G)$ (see Table 2 note for the $\kappa > 1$ regime). The flexibility index decomposes as

$$\kappa(G) = \underbrace{\bar{\Lambda}}_{\text{frequency}} + \underbrace{(\kappa - \bar{\Lambda})}_{\text{selection wedge}}, \quad (21)$$

where $\bar{\Lambda} = E_{f_{ss}}[\Lambda(x)]$ is the average adjustment frequency. Non-neutrality is governed by the selection wedge $\kappa - \bar{\Lambda}$.

The decomposition separates two conceptually distinct channels. Frequency ($\bar{\Lambda}$) measures how often firms reprice. The selection wedge ($\kappa - \bar{\Lambda}$) measures how well those repricings are targeted at the firms that need them most. Under Calvo, $\kappa = \bar{\Lambda} = \lambda$ and selection contributes nothing: repricing is random. For any absolutely continuous G , provided f_{ss} has non-degenerate mass on the interior region where Λ is strictly increasing, the selection wedge is strictly positive: firms with larger gaps are more likely to reprice, amplifying the aggregate price response beyond what frequency alone would produce.

This distinction has a striking implication. Two economies can share the same repricing frequency yet differ dramatically in non-neutrality. The difference is entirely attributable to the selection wedge, which depends on the shape of G .

Corollary 4 (Ordering of non-neutrality within the power law family).

For the power law family $G(\xi) = (\xi/\bar{\xi})^\gamma$ with $\bar{\xi}(\gamma)$ calibrated to a fixed average frequency $\bar{\Lambda}$, and assuming f_{ss} is supported in the interior of the inaction band:

- (i) **Calvo benchmark** ($\gamma \rightarrow 0^+$): selection wedge = 0; largest non-neutrality.
- (ii) **Interior** $\gamma > 0$: selection wedge = $2\gamma\bar{\Lambda} > 0$, strictly increasing in γ (Proposition 12); intermediate non-neutrality.
- (iii) **Golosov-Lucas** ($\gamma \rightarrow \infty$): selection wedge attains its supremum within the power law family; smallest non-neutrality.

The interior-support hypothesis is approximate under typical normal-shock calibrations; under that approximation the ordering above holds to the same approximation. A general rearrangement-style ordering across all absolutely continuous G requires additional regularity beyond the present framework and is not pursued here.

Within the power law family, the ordering is exact: γ alone determines selection, non-neutrality, and the Phillips curve slope.

Proposition 12 (Monotonicity in γ).

For the power law family $G(\xi) = (\xi/\bar{\xi})^\gamma$ with B/w fixed and $\bar{\xi}(\gamma)$ chosen so that the average adjustment frequency is constant across γ , and assuming f_{ss} is supported in the interior of the inaction band $|x| < \sqrt{w\bar{\xi}/B}$:

- (i) *The normalised selection wedge $(\kappa - \bar{\Lambda})/\bar{\Lambda} = 2\gamma$ is strictly increasing in γ .*
- (ii) *Monetary non-neutrality is strictly decreasing in γ .*
- (iii) *The Phillips curve slope at $\mu = 0$ is strictly increasing in γ .*

The interior-support hypothesis holds approximately under typical normal-shock calibrations, where the inaction band is wide relative to the shock variance; in that regime the bullets above hold to the same approximation. Thus γ indexes “how Calvo-like” the pricing environment is. Low γ (dispersed costs, weak selection, large non-neutrality) resembles Calvo. High γ (concentrated costs, strong selection, small non-neutrality) resembles Golosov-Lucas. The distribution G determines the economy’s position on this spectrum.

Proof. See Online Appendix C. ■

A natural question is whether repricing frequency alone determines non-neutrality. The answer is no. Frequency is only half the story; what matters equally is how that repricing is allocated across firms. The following proposition makes this precise.

Proposition 13 (Frequency does not determine non-neutrality).

Fix f_{ss} and B/w . There exist absolutely continuous distributions G_1, G_2 with:

- (i) *The same average repricing frequency: $\bar{\Lambda}_1 = E_{f_{ss}}[\Lambda_1(x)] = E_{f_{ss}}[\Lambda_2(x)] = \bar{\Lambda}_2$.*
- (ii) *Different flexibility indices: $\kappa(G_1) \neq \kappa(G_2)$, hence different selection wedges and different real effects of nominal shocks.*

Proof. See Online Appendix C. ■

For the power law family, the decomposition takes a particularly clean form. The ratio of total flexibility to frequency is pinned down by γ alone — independent of the steady-state distribution, the surplus parameter, or the wage.

Corollary 5 (Selection-frequency ratio for power laws).

For $G(\xi) = (\xi/\bar{\xi})^\gamma$, provided the tail mass of f_{ss} outside the inaction band is negligible (exact when f_{ss} is supported in the interior $|x| < \sqrt{w\xi/B}$):

$$\kappa = (1 + 2\gamma) \bar{\Lambda}, \quad \text{selection wedge} = 2\gamma \bar{\Lambda}. \quad (22)$$

The ratio $\kappa/\bar{\Lambda} = 1 + 2\gamma$ depends only on γ , not on f_{ss} , B , or w .

Proof.

For $G(\xi) = (\xi/\bar{\xi})^\gamma$, the density is $g(\xi) = (\gamma/\bar{\xi})(\xi/\bar{\xi})^{\gamma-1}$. In the interior of the band, $\Lambda(x) = (Bx^2/(w\bar{\xi}))^\gamma < 1$, and evaluating the intensive-margin integrand at $\xi = Bx^2/w$:

$$g\left(\frac{Bx^2}{w}\right) \frac{2Bx^2}{w} = \frac{\gamma}{\bar{\xi}} \left(\frac{Bx^2}{w\bar{\xi}}\right)^{\gamma-1} \frac{2Bx^2}{w} = 2\gamma \left(\frac{Bx^2}{w\bar{\xi}}\right)^\gamma = 2\gamma \Lambda(x).$$

Substituting into (19): $\kappa = \int \Lambda(x) f_{ss}(x) dx + \int 2\gamma \Lambda(x) f_{ss}(x) dx = (1 + 2\gamma) \bar{\Lambda}$. The identity requires that the uncapped power law expression applies f_{ss} -a.s., i.e., that f_{ss} has negligible mass outside the inaction band. In the calibrated model with normal shocks this condition is satisfied: the band is wide relative to the shock variance. ■

Table 2 illustrates. Six distributions calibrated to the same repricing frequency $\bar{\Lambda} = 0.10$ produce selection wedges ranging from 0 to 0.60. The implied cumulative output effects differ by a factor of three.

Figure 2 complements the table with the dynamic counterpart. Four power law economies — from near-Calvo ($\gamma = 0.20$) through uniform ($\gamma = 1$) to near-Golosov-Lucas ($\gamma = 5$) — are calibrated to the same quarterly repricing frequency. Each economy is hit with a one-time 1% monetary shock. The normalised output gap $y(t)/\delta$ is tracked as prices catch up to the new money supply.

The mechanism is selection: low γ means nearly random repricing and large persistent output effects; high γ means targeted repricing and small transient effects. The impulse responses are plotted in Figure 2 (Section 4.4).

3.8 Summary of the spectrum

Table 3 collects the pricing results: a single object, the menu cost distribution G , indexes the entire spectrum from Calvo (dispersed G , vanishing selection wedge)

Table 2: Same repricing frequency, different selection wedges

Distribution G	$\bar{\Lambda}$	κ	$\kappa - \bar{\Lambda}$	$1 - \kappa$
Calvo	0.10	0.10	0	0.90
Power law ($\gamma = \frac{1}{2}$)	0.10	0.20	0.10	0.80
Exponential	0.10	0.27	0.17	0.73
Uniform ($\gamma = 1$)	0.10	0.30	0.20	0.70
Power law ($\gamma = 2$)	0.10	0.50	0.40	0.50
Power law ($\gamma = 3$)	0.10	0.70	0.60	0.30

Notes: Each row is calibrated to $\bar{\Lambda} = 0.10$. The last column, $1 - \kappa$, is proportional to the cumulative real output effect of a monetary shock. For high γ and high $\bar{\Lambda}$, κ can exceed one; the linearisation $\pi \approx \kappa\mu$ then overstates the first-period price response, and the non-neutrality measure $1 - \kappa$ loses its direct interpretation. The formula remains valid as a local slope but should be interpreted with caution when $\kappa > 1$. *Calvo entry:* The Calvo row is the $\gamma \rightarrow 0^+$ limit of the power-law family with $\bar{\xi}(\gamma) \rightarrow \infty$ to hold $\bar{\Lambda}$ fixed (Proposition 9). At interior γ the hazard remains state-dependent; what vanishes is the selection wedge, not the state dependence of the individual hazard. *Power law entries:* $\kappa = (1 + 2\gamma)\bar{\Lambda}$, exact when f_{ss} is supported in the interior of the inaction band; approximate when the tail mass outside the band is negligible (the typical case under normal idiosyncratic shocks calibrated to standard menu cost frequencies). *Exponential entry:* $G(\xi) = 1 - e^{-\lambda\xi}$ with $f_{ss} = N(0, \sigma_x^2)$. Setting $\rho = B\sigma_x^2/w$, the χ^2 moment generating function gives $\bar{\Lambda} = 1 - (1 + 2\lambda\rho)^{-1/2}$ and intensive margin $\lambda\rho(1 + 2\lambda\rho)^{-3/2}$. Calibrating $\lambda\rho$ to match $\bar{\Lambda} = 0.10$ gives $\kappa \approx 0.27$.

through uniform G (quadratic costs, linear hazard, moderate selection) to Golosov-Lucas (point mass, step-function hazard, polar selection within the power law family). The power law γ traces this spectrum continuously; the empirical question — where real economies sit — is taken up in Section 4.

4 Quantitative analysis

The theoretical results of Sections 2–3 establish that the distribution G is the structural primitive governing adjustment dynamics. The closed-form hazard $\Lambda(x) = G(Bx^2/w)$ suggests an immediate identification route: for the power law family, $\log \Lambda$ is linear in $\log |x|$ with slope 2γ (equivalently, linear in $\log(x^2)$ with slope γ), so γ can in principle be read off from the hazard shape alone, without solving the full model. This section asks: what does G look like in the data? The answer disciplines the

Table 3: The state-dependence spectrum

	Calvo	Uniform G	Golosov-Lucas
Distribution G	Dispersed	$U[0, \bar{\xi}]$	Point mass at $\bar{\xi}$
$C(\varphi)$	Highly curved	Quadratic	Linear on $\{0, 1\}$
$C'''(\varphi)$	$\rightarrow \infty$	$w\bar{\xi}$	0 or ∞
Hazard $\Lambda(x)$	$\approx$ constant	$Bx^2/(w\bar{\xi})$	Step function
Selection	$\approx$ none	Moderate	Maximum
NKPC slope	$\approx$ constant	State-dependent	Very steep
Non-neutrality	Maximum	Intermediate	Minimum
Power law γ	$\rightarrow 0^+$	1	$\rightarrow \infty$

theory and quantifies the frequency-selection decomposition.

The quantitative analysis proceeds in two parts. Section 4.1 describes the hazard-based identification and notes that a model-free direct implementation on standard micro price datasets (scanner or CPI) is infeasible without observing the price gap x directly; model-based proxies for x exist but are not pursued here. Section 4.2 therefore identifies γ by matching CPI-based price-change moments (kurtosis and the fraction of small changes in a heterogeneous- σ_z simulation), with a robustness exercise on Dominick’s scanner data summarised in Section 4.4 and detailed in Online Appendix. The hazard-based route remains a testable prediction of the framework for datasets that permit constructing x directly.

4.1 A testable prediction: hazard-based identification

Before turning to the implemented estimator, I record one prediction the framework makes for future empirical work. The theoretical hazard for the power law family is $\Lambda(x) = (Bx^2/(w\bar{\xi}))^\gamma$. Taking logs gives

$$\log \Lambda(x) = 2\gamma \log |x| + \text{const.} \quad (23)$$

The slope of $\log \Lambda$ against $\log |x|$ is 2γ (equivalently, against $\log(x^2)$ the slope is γ). This identifies γ from the hazard shape alone, independently of $\bar{\xi}$, B , and w . A Monte Carlo exercise (Figure D.1 in the Online Appendix) confirms recovery to mean absolute error 0.012 and $R^2 > 0.99$.

Implementing this on actual data requires observing the price gap x at each deci-

sion point, which is not recorded in standard micro price datasets (scanner or CPI) — only price changes Δp . Model-based proxies for x exist but are not pursued here. Datasets with posted reference prices or shelf tags permit constructing x directly; that is the natural setting for the hazard-based test. The remainder of this section uses moment matching as the implemented estimator.

4.2 Calibration and moment matching

The baseline calibration sets $B = w = 1$ (normalisation) and $\sigma_z = 0.04$ (quarterly standard deviation of idiosyncratic desired-price shocks, consistent with Golosov and Lucas 2007 and Nakamura and Steinsson 2008). For each γ , the upper bound $\bar{\xi}$ is calibrated to match a quarterly repricing frequency of 24%, corresponding to approximately 8–9% monthly excluding temporary sales (Nakamura and Steinsson, 2008); under constant monthly hazards, $1 - (1 - 0.085)^3 \approx 0.234$. Throughout, φ is interpreted as a per-period probability.

In the baseline model with homogeneous σ_z , simulated kurtosis ranges from 1.33 (at $\gamma = 2$) to 2.79 (at $\gamma = 0.15$) across the γ grid (Table 4, left panel) — well below the empirical kurtosis of 4–6, a well-known limitation of single-product menu cost models with a homogeneous shock process. Following Midrigan (2011), the extended model draws per-product σ_z from a lognormal distribution with mean 0.04 and coefficient of variation 1.0. The CV is imposed, not estimated; this generates kurtosis in the empirical range while preserving the identification of γ from the hazard shape. Sensitivity to $CV \in \{0.5, 1.5\}$ is left for future quantitative work.

4.3 Calibrated value of γ from CPI moments

Illustrative calibrated value Matching kurtosis ≈ 5.0 in the heterogeneous- σ_z model yields

$$\hat{\gamma} \approx 0.4. \tag{24}$$

This is a single-moment calibrated value, not a fully identified estimate: kurtosis varies non-monotonically across the γ grid and the imposed $CV = 1$ for σ_z heterogeneity is not estimated. A right-skewed G corresponds to $\gamma < 1$, with the density $g(\xi) = \gamma \xi^{\gamma-1} / \bar{\xi}^\gamma$ decreasing: most firms face small menu costs while a few face large ones. The simulated fraction-of-small-changes at $\gamma = 0.4$ in the heterogeneous-

Table 4: Model moments: baseline and heterogeneous σ_z

γ	Baseline (homogeneous σ_z)					Heterogeneous σ_z (CV = 1)				
	Freq	$E[\Delta p]$	Kurt	Sm%	κ	Freq	$E[\Delta p]$	Kurt	Sm%	κ
0.15	0.24	7.2%	2.79	17%	0.31	0.24	7.9%	7.38	28%	0.31
0.25	0.24	7.5%	2.47	14%	0.36	0.24	8.8%	6.18	21%	0.36
0.40	0.24	7.8%	2.11	10%	0.43	0.24	9.7%	4.86	14%	0.44
0.60	0.24	8.1%	1.80	6%	0.53	0.24	10.3%	4.49	9%	0.53
0.80	0.24	8.3%	1.61	4%	0.63	0.24	10.6%	4.44	5%	0.63
1.00	0.24	8.4%	1.51	2%	0.72	0.24	10.8%	4.37	3%	0.73
1.50	0.24	8.6%	1.38	1%	0.96	0.24	11.0%	4.31	1%	0.97
2.00	0.24	8.7%	1.33	0%	1.20	0.24	11.1%	4.28	0%	1.21
Data targets	8–11%	4–6	25–40%			8–11%	4–6	25–40%		

Notes: All economies calibrated to quarterly repricing frequency $\bar{\Lambda} = 0.24$. “Freq” is the simulated frequency (confirming calibration). “ $E[|\Delta p|]$ ” is the mean absolute (non-sale) price change. “Kurt” is the raw kurtosis. “Sm%” is the fraction of price changes smaller than 2.5% in absolute value. $\kappa = (1 + 2\gamma)\bar{\Lambda}$ holds exactly when f_{ss} is supported in the interior of the inaction band and approximately to the extent that tail mass falls outside it; in the calibrated economies the approximation is tight under the standard normal-shock specification. Data targets from [Nakamura and Steinsson \(2008\)](#), [Midrigan \(2011\)](#), and [Alvarez, Le Bihan, and Lippi \(2016\)](#). 15,000 firms, 150 quarters per simulation.

σ_z model (14%) undershoots the empirical 25–40% target — a known limitation of single-product menu cost models with moderate σ_z heterogeneity ([Midrigan, 2011](#)) that absorbs into multi-product or sectoral structure not modelled here.

Implied decomposition At $\hat{\gamma} = 0.4$, the flexibility index is

$$\kappa = (1 + 2 \times 0.4) \times 0.24 = 0.43.$$

The frequency component is $\bar{\Lambda} = 0.24$, and the selection wedge is $\kappa - \bar{\Lambda} = 0.19$. Selection contributes 44% of κ :

$$\frac{\kappa - \bar{\Lambda}}{\kappa} = \frac{0.19}{0.43} = 0.44.$$

Non-neutrality is proportional to $1 - \kappa = 0.57$. Relative to the same-frequency Calvo benchmark ($1 - \bar{\Lambda} = 0.76$), this is $(1 - \kappa)/(1 - \bar{\Lambda}) = 0.75$: selection reduces non-neutrality by 25% relative to Calvo, a substantial but not dominant effect. [Table 5](#)

summarises.

Table 5: Frequency-selection decomposition at $\hat{\gamma} = 0.4$

	Value	Share of κ
Frequency ($\bar{\Lambda}$)	0.24	56%
Selection wedge ($\kappa - \bar{\Lambda}$)	0.19	44%
Flexibility (κ)	0.43	100%
Non-neutrality ($1 - \kappa$)	0.57	—

Notes: $\kappa = (1 + 2\hat{\gamma})\bar{\Lambda}$ with $\hat{\gamma} = 0.4$ and $\bar{\Lambda} = 0.24$. Non-neutrality as a fraction of Calvo: $(1 - \kappa)/(1 - \bar{\Lambda}) = 0.57/0.76 = 0.75$. Selection reduces non-neutrality by 25% relative to the Calvo benchmark at the same frequency.

Sensitivity of the selection share Since the selection share $2\gamma/(1 + 2\gamma)$ is an analytic function of γ , its sensitivity is transparent. At $\gamma = 0.15$, the selection share is 23%. At $\gamma = 0.25$, it is 33%. At the baseline $\gamma = 0.4$, it is 44%. At $\gamma = 0.8$, it would be 62%. The share is monotonically increasing in γ but bounded above by one. Kurtosis alone is consistent with γ in roughly $[0.15, 0.6]$ at the calibrated frequency, with the corresponding selection share spanning 23–54%. The headline 44% figure corresponds to the central $\hat{\gamma} \approx 0.4$ from kurtosis matching; tighter identification would require additional moments. Across this range, selection is a non-trivial but not dominant component of price flexibility. The formula $\kappa = (1 + 2\gamma)\bar{\Lambda}$, which is exact when f_{ss} is supported in the interior of the inaction band and approximately so otherwise, is independent of $\bar{\Lambda}$, so the selection share does not depend on the targeted repricing frequency.

Comparison with the literature The calibrated value $\hat{\gamma} \approx 0.4$ is consistent with several independent findings. [Midrigan \(2011\)](#) estimates a menu cost model with leptokurtic price changes and finds that matching the size distribution requires a multi-product structure that generates many small price changes alongside occasional large ones — precisely the pattern produced by a right-skewed G ($\gamma < 1$). [Alvarez, Le Bihan, and Lippi \(2016\)](#) report a kurtosis-frequency ratio of approximately 3.5, implying substantial but not maximal non-neutrality; the present calibration yields a CIR ratio of $\text{Kurt}/(6 \cdot \text{Freq}) \approx 4.86/(6 \times 0.24) = 3.4$, closely matching their sufficient statistic. This is a within-class consistency check, not an independent test, since both

calculations rest on the same sufficient-statistic formula. [Vavra \(2014\)](#) documents that the cross-sectional variance of price changes comoves with inflation — a prediction that follows directly from the state-dependent hazard $\Lambda(x) = G(Bx^2/w)$ whenever G generates a non-constant hazard.

4.4 Hazard functions and impulse responses

Figure 1 plots the simulated repricing hazard for four values of γ , along with the theoretical prediction (dashed). The hazard is concave in $|x|$ for $\gamma < 1$ (firms with moderate gaps adjust frequently, but the marginal increase for very large gaps is small) and convex for $\gamma > 1$ (adjustment is concentrated sharply at the threshold). The calibrated $\hat{\gamma} \approx 0.4$ produces a gently increasing, concave hazard — qualitatively consistent with the smooth adjustment functions estimated by [Caballero and Engel \(1993\)](#).




Figure 1: Repricing hazard across the spectrum

Notes: Repricing hazard $\Lambda(x)$ for selected values of γ . All economies calibrated to $\bar{\Lambda} = 0.24$ quarterly. Low γ produces a nearly flat hazard (close to Calvo); high γ produces a steep hazard (close to Golosov-Lucas).

Figure 2 shows the output impulse response to a one-time 1% monetary shock. All

economies share the same repricing frequency; they differ only in γ . The Calvo benchmark (zero selection) decays exponentially at rate $\bar{\Lambda}$. State-dependent economies converge faster because repricing is targeted at the most mispriced firms. The gap between Calvo and the state-dependent economy widens with γ : at $\hat{\gamma} = 0.4$, the output gap is approximately 75% of the Calvo path — selection matters, but frequency remains the dominant force.




Figure 2: Output impulse response to a 1% monetary shock

Notes: All economies calibrated to $\bar{\Lambda} = 0.24$ quarterly; they differ only in γ . Higher γ implies stronger selection, faster price-level convergence, and smaller cumulative output effects. CIR values computed via [Alvarez, Le Bihan, and Lippi \(2016\)](#): $\text{CIR} = \text{Kurt}/(6 \cdot \text{Freq})$. The Calvo benchmark (dashed) has zero selection.

Figure D.2 in the Online Appendix shows the distribution of price changes for several values of γ ; the calibrated $\hat{\gamma} = 0.4$ produces a moderately leptokurtic distribution with kurtosis around 5, matching the empirical evidence.

Relation to [Alvarez, Le Bihan, and Lippi \(2016\)](#) The ABLL sufficient statistic characterises non-neutrality in terms of the kurtosis and frequency of price changes. Within the power law family, both are endogenous functions of γ and $\bar{\xi}$. The present

framework therefore provides the structural layer from which ABLL’s observable diagnostics are derived: γ determines the kurtosis-frequency ratio, and hence the CIR, through the shape of the repricing hazard. The practical value of γ over the ABLL moments is that γ is a single structural parameter with a clear interpretation (the shape of the cost distribution), whereas the ABLL ratio combines two moments that can move independently in response to changes in the economic environment (e.g., changes in trend inflation shift both frequency and kurtosis, but the mapping depends on G).

Robustness on Dominick’s scanner data A moment exercise on the Dominick’s Finer Foods scanner dataset (Midrigan, 2011; Chevalier, Kashyap, and Rossi, 2003) (Online Appendix E) yields a non-sale kurtosis of 7.6 at a quarterly repricing frequency of 0.72 — substantially higher than CPI moments. Recalibrated to Dominick’s frequency, simulated kurtosis is essentially flat across $\gamma \in [0.10, 0.40]$: the heterogeneity-driven kurtosis variation flattens at high frequency, so kurtosis bounds $\gamma \lesssim 0.4$ but does not pin it down further. The CPI/Dominick’s gap is consistent with sectoral heterogeneity in G , paralleling Oskolkov and Lippi (2026)’s plant-level investment evidence; modelling sector-specific $(\gamma_s, \bar{\xi}_s)$ explicitly is a natural next step.

5 Discussion

Further applications Beyond pricing (and investment, in Online Appendix B), the duality applies to any extensive-margin decision with heterogeneous costs. Heterogeneous entry costs micro-found smooth aggregate entry: the free-entry condition $S = wG^{-1}(\varphi)$ equates the marginal entrant’s cost to the marginal cost of raising the entry rate. Heterogeneous vacancy-creation costs micro-found convex posting costs (Pissarides, 2000), with C'' reflecting the cost distribution rather than a technology. Heterogeneous participation costs generate smooth labour supply without Rogerson (1988) lotteries, with elasticity $g(G^{-1}(\varphi))/w$. The same logic extends to technology adoption and credit default (Bernanke, Gertler, and Gilchrist, 1999).

Structural discipline and identifiability The G framing makes distributional assumptions transparent: quadratic costs require uniform G ; iso-elastic costs require power law G . It supplies a common language — the Calvo–Golosov–Lucas debate,

selection in non-neutrality, and the [Khan and Thomas \(2008\)](#) irrelevance result are all statements about G — and an identifiable primitive: G is fully recovered from the adjustment hazard (Section 4.1) and disciplined by price-change size distributions and investment spike frequencies (model-dependent). [Oskolkov and Lippi \(2026\)](#) recover G from plant-level data, where the duality implies the same G simultaneously determines C , the investment hazard, and capital misallocation. The aggregate smoothness emphasised by [Auclert, Rognlie, and Straub \(2020\)](#) has an individual-level structural counterpart here: whenever G has a positive density, each agent’s problem is smooth even though realised actions remain binary; the cross-sectional aggregation step is separate.

Scope and limitations The equivalence is an identity of value functions and adjustment probabilities, not of welfare, interpretation, or causality. The identifiable object is G (equivalently C), not the “type” of friction. Random-menu-cost and probability-choice economies make different predictions about how G responds to changes in the environment; endogenising G is a natural next step.

6 Conclusion

The equivalence established in this paper changes how one should think about adjustment frictions. The distinction between fixed and convex costs is not a distinction between two kinds of economic friction. It is a distinction between two representations of the same friction, linked ex-ante by Fenchel conjugacy. The fixed cost distribution G is the structural primitive that governs both.

Three implications follow. First, the lumpy–smooth divide dissolves at the ex-ante extensive margin. Realised actions remain binary at the firm level; what becomes smooth and convex is the ex-ante extensive margin, whenever fixed costs are absolutely continuous with positive density. The deterministic fixed cost — the textbook case — is the singular exception. This recasts the standard view that convex costs approximate the “true” lumpy friction: at the ex-ante extensive margin, the deterministic fixed cost is the knife-edge limit and arbitrarily close perturbations restore strict convexity. The same observation provides a partial-equilibrium foundation for the irrelevance result of [Khan and Thomas \(2008\)](#): with uniform fixed costs, each firm’s ex-ante extensive margin is already convex, so aggregation is not the source of

smoothness.

Second, the equivalence connects literatures with shared structure: menu costs, lumpy investment, discrete choice, rational inattention, and control costs all share the same mapping $G \mapsto C$, with the contributions varying in degree (genuine new connections for menu costs and lumpy investment; identifications of objects already studied for discrete choice and control costs). Under the fixed-cost primitive, $C(\varphi^*)$ has a concrete resource-cost interpretation. The smooth first-order condition $S = C'(\varphi)$ replaces kink-based dynamic programming, accelerating computation.

Third, in the pricing application under the static-surplus approximation, the equivalence gives the repricing hazard in closed form and enables an analytical decomposition of non-neutrality into frequency and selection. This decomposition rests on the explicit functional form $\Lambda(x) = G(Bx^2/w)$.

Several directions remain open: GE embedding to quantify inflation welfare across the spectrum; endogenising G in response to the inflation environment, connecting to the Lucas critique and to [Vavra \(2014\)](#); and multi-sector identification testing whether γ varies across markets.

References

- Alvarez, Fernando, Hervé Le Bihan, and Francesco Lippi. 2016. “The Real Effects of Monetary Shocks in Sticky Price Models: A Sufficient Statistic Approach.” *American Economic Review* 106 (10):2817–2851.
- Auclert, Adrien, Matthew Rognlie, and Ludwig Straub. 2020. “Micro Jumps, Macro Humps: Monetary Policy and Business Cycles in an Estimated HANK Model.” Working Paper 26647, National Bureau of Economic Research.
- Bernanke, Ben S., Mark Gertler, and Simon Gilchrist. 1999. “The Financial Accelerator in a Quantitative Business Cycle Framework.” In *Handbook of Macroeconomics*, vol. 1. Elsevier, 1341–1393.
- Bloedel, Alexander W., Tommaso Denti, and Luciano Pomatto. 2025. “Modeling Information Acquisition via f -Divergence and Duality.” Working paper, UCLA, NYU Stern, and Caltech. Working paper, October 2025.

- Caballero, Ricardo J. and Eduardo M. R. A. Engel. 1993. "Microeconomic Rigidities and Aggregate Price Dynamics." *European Economic Review* 37 (4):697–711.
- . 1999. "Explaining Investment Dynamics in U.S. Manufacturing: A Generalized (S, s) Approach." *Econometrica* 67 (4):783–826.
- . 2007. "Price Stickiness in Ss Models: New Interpretations of Old Results." *Journal of Monetary Economics* 54 (Supplement 1):100–121.
- Calvo, Guillermo A. 1983. "Staggered Prices in a Utility-Maximizing Framework." *Journal of Monetary Economics* 12 (3):383–398.
- Caplin, Andrew S. and Daniel F. Spulber. 1987. "Menu Costs and the Neutrality of Money." *Quarterly Journal of Economics* 102 (4):703–725.
- Chevalier, Judith A., Anil K. Kashyap, and Peter E. Rossi. 2003. "Why Don't Prices Rise During Periods of Peak Demand? Evidence from Scanner Data." *American Economic Review* 93 (1):15–37.
- Dotsey, Michael, Robert G. King, and Alexander L. Wolman. 1999. "State-Dependent Pricing and the General Equilibrium Dynamics of Money and Output." *Quarterly Journal of Economics* 114 (2):655–690.
- Gabaix, Xavier. 2014. "A Sparsity-Based Model of Bounded Rationality." *Quarterly Journal of Economics* 129 (4):1661–1710.
- Golosov, Mikhail and Robert E. Lucas, Jr. 2007. "Menu Costs and Phillips Curves." *Journal of Political Economy* 115 (2):171–199.
- Gourio, François and Anil K. Kashyap. 2007. "Investment Spikes: New Facts and a General Equilibrium Exploration." *Journal of Monetary Economics* 54 (Supplement 1):1–22.
- Hofbauer, Josef and William H. Sandholm. 2002. "On the Global Convergence of Stochastic Fictitious Play." *Econometrica* 70 (6):2265–2294.
- Khan, Aubhik and Julia K. Thomas. 2008. "Idiosyncratic Shocks and the Role of Nonconvexities in Plant and Aggregate Investment Dynamics." *Econometrica* 76 (2):395–436.

- Mattsson, Lars-Göran and Jörgen W. Weibull. 2002. “Probabilistic Choice and Procedurally Bounded Rationality.” *Games and Economic Behavior* 41 (1):61–78.
- Matějka, Filip and Alisdair McKay. 2015. “Rational Inattention to Discrete Choices: A New Foundation for the Multinomial Logit Model.” *American Economic Review* 105 (1):272–298.
- McFadden, Daniel. 1974. “Conditional Logit Analysis of Qualitative Choice Behavior.” In *Frontiers in Econometrics*, edited by Paul Zarembka. Academic Press, 105–142.
- . 1978. “Modelling the Choice of Residential Location.” In *Spatial Interaction Theory and Planning Models*, edited by A. Karlqvist et al. North-Holland, 75–96.
- . 1981. “Econometric Models of Probabilistic Choice.” In *Structural Analysis of Discrete Data*, edited by Charles F. Manski and Daniel McFadden. MIT Press, 198–272.
- Midrigan, Virgiliu. 2011. “Menu Costs, Multiproduct Firms, and Aggregate Fluctuations.” *Econometrica* 79 (4):1139–1180.
- Nakamura, Emi and Jón Steinsson. 2008. “Five Facts about Prices: A Reevaluation of Menu Cost Models.” *Quarterly Journal of Economics* 123 (4):1415–1464.
- . 2010. “Monetary Non-Neutrality in a Multisector Menu Cost Model.” *Quarterly Journal of Economics* 125 (3):961–1013.
- Oskolkov, Aleksei and Francesco Lippi. 2026. “Misallocation with Lumpy Investment.” Discussion Paper 21212, Centre for Economic Policy Research.
- Pissarides, Christopher A. 2000. *Equilibrium Unemployment Theory*. MIT Press, 2nd ed.
- Reiter, Michael. 2009. “Solving Heterogeneous-Agent Models by Projection and Perturbation.” *Journal of Economic Dynamics and Control* 53 (6):649–665.
- Rogerson, Richard. 1988. “Indivisible Labor, Lotteries and Equilibrium.” *Journal of Monetary Economics* 21 (1):3–16.

- Sims, Christopher A. 2003. “Implications of Rational Inattention.” *Journal of Monetary Economics* 50 (3):665–690.
- Thomas, Julia K. 2002. “Is Lumpy Investment Relevant for the Business Cycle?” *Journal of Political Economy* 110 (3):508–534.
- Tsallis, Constantino. 1988. “Possible Generalization of Boltzmann-Gibbs Statistics.” *Journal of Statistical Physics* 52 (1–2):479–487.
- Vavra, Joseph. 2014. “Inflation Dynamics and Time-Varying Volatility: New Evidence and an Ss Interpretation.” *Quarterly Journal of Economics* 129 (1):215–258.