

Dancing on the Saddles: A Geometric Framework for Stochastic Equilibrium Dynamics[†]

Hanbaek Lee[‡]

University of Cambridge

July 2026

[\(click here for the latest version\)](#)

Abstract

Extending saddle-path analysis to stochastic environments, I introduce *conditional saddle paths*: equilibrium trajectories under frozen exogenous regimes, traced by stochastic-equilibrium decision rules. Their *geometry* diagnoses state dependence: shock responses depend on position along the saddle and on slope gaps across regimes. Their *monotonicity* delivers dimension reduction: aggregate capital exactly indexes the equilibrium state within a frozen-regime spell in [Krusell and Smith \(1998\)](#), and a synchronization theorem gives sufficient conditions under which the distributional gap across histories shrinks with their shared recent regime path. Finally, the framework turns policy comparative statics in stochastic general equilibrium into geometry: in a search-and-matching economy, unemployment volatility is hump-shaped in the benefit level, because high benefits pull the good- and bad-regime conditional steady states together, leaving little room for unemployment to fluctuate between them.

Keywords: Conditional saddle path, business cycles, state-dependent dynamics, sufficient statistics, heterogeneous agents, search and matching.

JEL codes: C62, D58, E32, J64.

[†]I am grateful for helpful comments and discussions from Martin M. Andreasen, Yunus Aksoy, Marlon Azinovic-Yang, Adrien Bilal, Timo Boppart, Francisco Buera, Christopher Carroll, Vasco Carvalho, Minsu Chang, Rachel Leah Childers, Miguel H. Ferreira, Elisa Giannone, Matthias Hänsel, Tomohiro Hirano, Dirk Krueger, Andrea Lanteri, Benjamin Larin, Danial Lashkari, Tim Lee, Benjamin Moll, Juan Carlos Parra-Alvarez, Tommaso Porzio, Michael Reiter, Itay Saporta-Eksten, Yucheng Yang, Yuki Yao, and Jan Žemlička, as well as the seminar participants at Birkbeck, University of London, the University of Manchester, Aarhus University, QMUL E1 macro conference, and SED at Athens. All errors are my own.

[‡]Email: hl610@cam.ac.uk

1 Introduction

Understanding equilibrium dynamics in models with aggregate uncertainty remains a central challenge in macroeconomics. Unlike deterministic models, where saddle-path diagrams provide immediate geometric intuition, stochastic equilibrium models lack comparable geometric frameworks. This makes it difficult to develop intuition about how economies respond to shocks, how different states interact, and why certain computational methods work. The challenge becomes particularly acute in heterogeneous-agent models, where the natural state variable is an infinite-dimensional wealth distribution—yet computational work routinely achieves dimension reduction to low-dimensional aggregates. Can we extend geometric saddle-path analysis to stochastic environments? What does such a framework reveal about equilibrium dynamics, state-dependent responses, and the success of computational approximations?

The paper makes three contributions, all drawn from a single geometric object. First, I define *conditional saddle paths*: equilibrium trajectories under frozen exogenous regimes, traced by the decision rules of the stochastic equilibrium. Freezing the regime isolates the economy’s internal dynamics while agents continue to price in future switches, and each regime’s trajectory draws a policy graph in the aggregate phase plane. These objects decompose business-cycle fluctuations into movements *along* a conditional saddle (endogenous propagation) and *across* conditional saddles (exogenous shocks), providing the stochastic analogue of the deterministic phase diagram. The geometry carries the economics. Shock impacts are state-independent exactly when the saddles are parallel (vertical translations of one another); relative tilt makes the same shock hit harder in some states than in others; and even with parallel saddles, the *shape* of the response, its peak timing and persistence, depends on the economy’s position along the saddle. I develop these results in a canonical RBC model (Sections 2–3).

Second, the same geometry speaks to a long-standing computational puzzle: why tracking a single aggregate can stand in for an entire wealth distribution. The answer is a property of the saddle. When an aggregate variable moves monotonically along the conditional saddle, it exactly indexes the equilibrium state within a frozen-regime spell, so that one number recovers the distribution and every price. I show that aggregate capital plays this role in [Krusell and Smith \(1998\)](#), exactly near the conditional steady state and, under conditions checkable from calibrated parameters, on an explicit region away from it. The result is candid about its edge: it must rule out a stall in which capital stops moving while the distribution keeps rearranging, and across spells, where the same capital can sit behind different distributions, a synchronization argument controls the discrepancy once the frozen-

regime dynamics contract. A perfect-foresight diagnostic then carries the key condition from a single deterministic transition path, the kind practitioners already compute, to the full rational-expectations equilibrium (Section 4.1).

Third, the same geometry turns policy comparative statics in stochastic general equilibrium into a picture. In a search-and-matching economy (Section 4.2), the saddle geometry captures how micro- and macro-level policy parameters reshape equilibrium volatility. Conditional saddles trace upward-sloping, regime-specific vacancy loci whose across-regime transitions generate the empirical Beveridge-curve loop, and the surplus compression behind the Shimer puzzle manifests as differential saddle steepness. Comparative statics on unemployment benefits yield a new result: volatility is hump-shaped in the benefit level. The reason is geometric: each regime pulls unemployment toward its own conditional steady state, and regime switches move the economy back and forth between them, so volatility is governed by the *gap between the conditional steady states*, not by slope asymmetry alone. At high benefit levels, the conditional steady states converge despite large slope gaps, leaving little room for unemployment to fluctuate between regimes. Across the three, the conclusions are read from one object: the slope of the conditional saddle gives state dependence, its monotonicity gives dimension reduction, and the spacing of its resting points gives comparative statics.

Related literature This paper contributes to three strands of literature in macroeconomics. The first is the literature studying global equilibrium dynamics and solution methods under aggregate uncertainty. The challenge of characterizing equilibrium dynamics in models with aggregate shocks has motivated extensive methodological development. [Marcet \(1988\)](#), [Den Haan and Marcet \(1990\)](#), and [Krusell and Smith \(1998\)](#) pioneered the use of bounded rationality approximations. In particular, [Krusell and Smith \(1998\)](#) discovered that a simple linear forecasting rule in aggregate capital achieves remarkable accuracy ($R^2 > 0.9999$) despite the infinite-dimensional state space. Subsequent work has sharpened accuracy and efficiency through perturbation and projection methods ([Reiter, 2010](#); [Ahn et al., 2018](#); [Levintal, 2018](#); [Winberry, 2018](#); [Bayer and Luetticke, 2020](#); [Bhandari et al., 2023](#)), simulation and aggregation approaches ([Den Haan, 1996](#); [Den Haan and Rendahl, 2010](#); [Judd et al., 2011](#); [Maliar and Maliar, 2015](#)), adaptive sparse grids ([Brumm and Scheidegger, 2017](#)), sequence-space methods ([Auclert et al., 2021](#); [Azinovic-Yang and Žemlička, 2025](#)), and deep learning ([Maliar et al., 2021](#); [Azinovic et al., 2022](#); [Fernández-Villaverde et al., 2023](#); [Han et al., 2026](#)). Rather than reducing the state, [Childers \(2018\)](#) linearizes in function space, treating the entire distribution as the state, with consistency guarantees for the resulting approximation. On the theoretical side of this literature, [Miao \(2006\)](#) and [Cao \(2020\)](#) establish existence of competitive and recursive equilibria in this class of economies; [Proehl \(2025\)](#) provides conditions for the

existence and uniqueness of a recursive equilibrium under aggregate and idiosyncratic risk, and [Proehl \(2024\)](#) develops a companion global solution method. The framework here is complementary: it takes the recursive equilibrium as given and characterizes the geometry of its dynamics.

The sequence-space Jacobian (SSJ) method of [Auclert et al. \(2021\)](#) approaches dimension reduction from a different angle. SSJ linearizes around a steady state; dimension reduction follows from the additive structure of linearized IRFs. By contrast, conditional saddle paths operate globally and nonlinearly: dimension reduction follows from monotonicity of an aggregate coordinate along the frozen-regime equilibrium path. Near a conditional steady state the two coincide to first order: the tangent of the policy graph carries the same information as the SSJ linearization. Away from the steady state, the conditional saddle can bend, capturing state-dependent dynamics that first-order methods cannot, such as the regime-dependent slope asymmetries that frictions generate (Section 3).

This paper differs from these computational contributions by providing geometric foundations for *why* dimension reduction works. Aggregate capital’s within-spell sufficiency follows from geometric properties (nullcline invariance and monotonicity) rather than numerical happenstance. The result also has computational implications: along each conditional saddle, aggregate capital provides a theoretically justified metric for the repeated transition method of [Lee \(2026a\)](#).¹

Second, this paper extends the geometric tradition of phase-diagram analyses in growth theory ([Solow, 1956](#); [Swan, 1956](#); [Ramsey, 1928](#); [Koopmans, 1963](#); [Cass, 1965](#); [Barro and Sala-i-Martin, 1992](#)) to stochastic environments. Within this tradition, [Buera et al. \(2024\)](#) pursue a complementary generalization: their Stable Transformation Path extends the balanced growth path, and its turnpike property, to non-stationary growth models, defining a stable path independent of initial conditions. Conditional saddle paths extend the same classical geometry along a different dimension, from deterministic to stochastic settings. They provide a visual tool for understanding how microfounded frictions generate state-dependent dynamics, as documented empirically ([Kaplan and Violante, 2014](#); [Vavra, 2014](#); [Berger and Vavra, 2015](#); [Basu and Bundick, 2017](#); [Bloom et al., 2018](#); [Kaplan et al., 2018](#); [Petrosky-Nadeau et al., 2018](#); [Rognlie et al., 2018](#); [Baley and Blanco, 2019](#); [Pizzinelli et al., 2020](#); [Berger et al., 2021](#); [Winberry, 2021](#); [Lee, 2026b](#); [Melcangi, 2024](#)).

¹The idea of solving heterogeneous-agent models by freezing the aggregate state and analyzing transitions appears in several computational approaches, including [Bourany \(2018\)](#), Section 4.3 of a working paper version of [Achdou et al. \(2021\)](#), and [Lee \(2026a\)](#). In continuous time, the master-equation literature ([Achdou et al., 2021](#); [Bilal, 2023](#)) treats the distribution itself as the state and asks when finite-dimensional projections suffice—the rigorous mathematical home of the question studied here. The present paper provides a theoretical foundation: conditional saddle paths formalize the underlying geometric objects, and the sufficiency results explain when dimension reduction is exact along frozen-regime paths.

Finally, in the random dynamical systems literature ([Arnold, 1998](#); [Schenk-Hoppé, 1998, 2001](#)), the focus is on history-dependent stochastic saddle paths ([Yannacopoulos, 2011](#)), which vary with the realized shock sequence.² By contrast, conditional saddle paths are history-invariant geometric objects that provide economists with the direct analogue of deterministic phase diagrams.

2 Conditional saddle path

2.1 Definitions and assumptions

I consider a generic dynamic stochastic model where the corresponding recursive competitive equilibrium (RCE) is characterized by the following aggregate state S and the endogenous and exogenous laws of motion ($\Gamma_{\text{endo}}, \Gamma_{\text{exo}}$):

$$S = [\Phi, A] \tag{1}$$

where Φ is the endogenous aggregate state variable, and A is the exogenous aggregate state variable. The latter may be a multivariate vector that follows a stochastic process. I assume the exogenous aggregate law of motion Γ_{exo} is a Markov chain. For simplicity in the illustration, I assume Γ_{exo} is a two-state Markov chain where $A \in \{B, G\}$, and $\Gamma_{\text{exo}}(A'|A) > 0 \forall A', A$.

I consider a distributional state space $\mathcal{X}$, whose elements $\Phi \in \mathcal{X}$ summarize the cross-sectional distribution of idiosyncratic household states that are payoff-relevant for equilibrium (e.g., assets and employment/productivity types) together with any endogenous objects needed to evaluate equilibrium decision rules. Formally, one can take $\mathcal{X}$ to be a subset of a metric space of probability measures augmented, if needed, by a finite-dimensional vector of aggregate variables. Throughout, $\Gamma_{\text{endo}}(\cdot, A) : \mathcal{X} \rightarrow \mathcal{X}$ denotes the recursive equilibrium law of motion mapping the current distributional state into the next-period state under frozen regime A .

The central object of this paper formalizes a natural thought experiment. Suppose the economy is in state (Φ_0, A) today and the exogenous regime happens to remain at A

²In Random Dynamical Systems language, the conditional steady state is a deterministic realization of the random fixed point ([Schenk-Hoppé and Schmalfuß, 2001](#)), and the conditional saddle path is the invariant manifold of the frozen-regime dynamics (the curve that those dynamics map into itself). The connection is put to work in Section 4.1: the synchronization theorem (Theorem 2), which underlies the across-spell bound (Proposition 9), uses exactly the pullback random-fixed-point construction of [Schenk-Hoppé and Schmalfuß \(2001\)](#): the current state is recovered as the long-history limit of trajectories started ever further in the past under the same regime history.

indefinitely. The resulting equilibrium trajectory, computed under decision rules that fully internalize future regime changes, is a frozen-regime orbit. Its closure (the set of states the orbit visits, together with their limit points) is the conditional saddle path. In aggregate phase diagrams, the corresponding regime-specific locus is a conditional policy graph; slope and curvature statements below refer to these graph projections.

Definition 1 (Frozen-regime orbit and conditional saddle path).

Fix a regime $A \in \{B, G\}$ and an initial state $\Phi_0 \in \mathcal{X}$. Let $\{\Phi_t\}_{t \geq 0}$ denote the frozen-regime continuation under A , defined recursively by

$$\Phi_{t+1} = \Gamma_{\text{endo}}(\Phi_t, A), \quad t \geq 0. \quad (2)$$

The frozen-regime orbit is

$$\mathcal{O}(\Phi_0, A) := \{\Phi_t : t \geq 0\}. \quad (3)$$

The conditional saddle path under A from Φ_0 is defined as the closure of the frozen-regime orbit:

$$\mathcal{M}(\Phi_0, A) := \overline{\mathcal{O}(\Phi_0, A)}. \quad (4)$$

Throughout, the distributional state space $(\mathcal{X}, d)$ is a metric space whose elements are the distributional states just described: two states are close in d when the averages they assign to any well-behaved function of the individual state are close, with any finite-dimensional components compared in the Euclidean norm.³ The closure in Definition 1 is taken with respect to this metric.

Two projection ingredients recur throughout: a *coordinate* $e : \mathcal{X} \rightarrow \mathbb{R}$, a predetermined aggregate variable such as aggregate capital K or inherited unemployment u_{-1} ; and an *aggregate object* $g(\Phi, A)$, an equilibrium outcome such as consumption or vacancies.

Definition 2 (Conditional policy graph).

For a regime A , a coordinate e , and an aggregate object g , the conditional policy graph in the (e, g) projection is

$$\mathcal{P}_A^g := \{(e(\Phi), g(\Phi, A)) : \Phi \in \mathcal{D}_A\}, \quad (5)$$

where $\mathcal{D}_A$ is the set of distributional states reached by the frozen- A continuations under study.

³Formally, d is any metric for which $d(\Phi_n, \Phi) \rightarrow 0$ if and only if $\int f d\Phi_n \rightarrow \int f d\Phi$ for every bounded continuous function f ; this is the standard notion of weak convergence.

When e is injective (one-to-one) on a given conditional saddle path, the restriction of $\mathcal{P}_A^g$ to that path is the graph of a single-valued function of e .

Based on the conditional saddle path, I define the conditional steady state as follows:

Definition 3 (Conditional steady state).

Given (Φ_0, A) , if the frozen-regime continuation $\{\Phi_t\}_{t \geq 0}$ converges, I denote its limit by

$$\Phi^{cs}(\Phi_0, A) := \lim_{t \rightarrow \infty} \Phi_t. \quad (6)$$

Definition 1 defines the conditional saddle path directly from the frozen-regime continuation. Fixing A makes the equilibrium law of motion deterministic on the distribution space: starting from Φ_0 , the sequence $\{\Phi_t\}_{t \geq 0}$ is generated by repeated application of the endogenous transition operator $\Gamma_{\text{endo}}(\cdot, A)$. In this sense, $\mathcal{M}(\Phi_0, A)$ is the natural analogue of the saddle arm in deterministic saddle-path analysis. It isolates the equilibrium states visited (and their limit points) along the convergence dynamics under regime A .⁴

For the conditional saddle to serve as a useful geometric object (a one-dimensional curve that the economy traverses during each frozen-regime spell), the frozen-regime dynamics must converge to a unique limit and vary smoothly. The following assumption collects these regularity conditions.

Assumption 1 (Regularity of conditional saddle paths).

Fix a regime $A \in \{B, G\}$ and an initial state $\Phi_0 \in \mathcal{X}$. Let $\{\Phi_t\}_{t \geq 0}$ be defined by $\Phi_{t+1} = \Gamma_{\text{endo}}(\Phi_t, A)$.

- (i) (Unique existence) The conditional steady state $\Phi^{cs}(\Phi_0, A)$ uniquely exists on $\mathcal{M}(\Phi_0, A)$.
- (ii) (Continuity) The map $\Gamma_{\text{endo}}(\cdot, A) : \mathcal{X} \rightarrow \mathcal{X}$ is continuous. Moreover, the aggregate variables (e.g., aggregate capital K and consumption C) vary continuously in Φ when restricted to $\mathcal{M}(\Phi_0, A)$.

Assumption 1 collects the regularity properties needed for the paper's geometric arguments. Part (i) requires that, conditional on a regime A and an initial state Φ_0 , the frozen-regime continuation converges to a unique conditional steady state. This does not assert global uniqueness across all initial conditions, nor does it rule out multiplicity across

⁴Strictly speaking, $\mathcal{M}(\Phi_0, A) = \overline{\mathcal{O}(\Phi_0, A)}$ in discrete time is an orbit-closure, not necessarily a smooth manifold (in this one-dimensional setting, a smooth curve). I use saddle-path language to emphasize its one-dimensional, path-like role. Slope and curvature language applies to the conditional policy graph $\mathcal{P}_A^g$ in a chosen aggregate projection; in continuous time, the corresponding frozen-regime object is a one-dimensional invariant manifold; see Appendix L for the full continuous-time formulation.

regimes.⁵ Part (ii) imposes continuity of the frozen-regime law of motion and aggregate observables on the conditional saddle, ensuring that small perturbations of the distributional state produce small changes in aggregate outcomes. Assumption 1(i) thus restricts attention to convergent frozen-regime continuations, ruling out periodic orbits or divergent paths. In the models studied below, convergence follows from standard arguments: diminishing returns in the RBC benchmark and contractivity of the Bewley–Huggett–Aiyagari operator (Bewley, 1986; Huggett, 1993; Aiyagari, 1994) in the heterogeneous-household economy.

With this regularity, the conditional saddle admits a natural notion of position along the frozen-regime orbit. Theorem 1 (Section 4.1) formalizes the key implication: if some aggregate equilibrium variable moves strictly one-way along this path, it provides a valid coordinate for the entire conditional saddle, yielding an exact one-dimensional representation on $\mathcal{M}(\Phi_0, A)$.

2.2 Representative-agent economy: A canonical RBC

For illustrative purposes, I consider a representative household with temporal log utility. Given initial condition (a_0, A_0) , the household maximizes lifetime utility under stochastic aggregate TFP A_t . I present the economy in recursive form and work with a recursive competitive equilibrium. Under standard regularity conditions, this formulation is equivalent to the sequential equilibrium; the sequential formulation is in Appendix A⁶. The recursive form is as follows:

$$v(a; K, A) = \max_{c, a'} \log(c) + \beta \mathbb{E}v(a'; K', A') \quad (7)$$

$$c + a' = a(1 + r(K, A)) + w(K, A) \quad (8)$$

$$a' \geq -\bar{a} \quad (9)$$

$$K' = \Gamma_{\text{endo}}(K, A), \quad A' \sim \Gamma_{\text{exo}}(A'|A), \quad (10)$$

where v is the household's value function; a is wealth at the beginning of a period; K is aggregate capital; A is aggregate TFP; Γ_{endo} is the law of motion for K . For illustrative purposes, I assume that TFP A follows a Markov-switching process between the levels B

⁵In particular, Proehl (2025) establishes global existence and uniqueness of the recursive equilibrium under explicit conditions on primitives, whereas Assumption 1 (i) allows for global multiplicity as in Walsh and Young (2024). Nevertheless, conditional on a given initial state and a fixed regime, the equilibrium path is assumed to be uniquely determined.

⁶All appendices (A–N) are collected in a separate Online Appendix.

and G , where $G > B$:

$$\Gamma_{\text{exo}} = \begin{bmatrix} \pi_{BB} & \pi_{BG} \\ \pi_{GB} & \pi_{GG} \end{bmatrix} \quad (11)$$

I consider competitive factor prices given a CRS Cobb–Douglas production function:

$$r(K, A) = A\alpha K^{\alpha-1} - \delta \quad (12)$$

$$w(K, A) = A(1 - \alpha)K^\alpha, \quad (13)$$

where K is capital stock, which satisfies $K = a$ in equilibrium. The recursive competitive equilibrium (RCE, hereafter) is as follows:

Definition 4 (Recursive competitive equilibrium).

$(c, a', v, r, w, \Gamma_{\text{endo}})$ is a recursive competitive equilibrium if:

1. given prices (r, w) , the policies (c, a') attain the value v in the household problem;
2. prices satisfy the competitive factor-price conditions above;
3. consistency holds:

$$a'(K; K, A) = K' = \Gamma_{\text{endo}}(K, A). \quad (14)$$

Based on the endogenous law of motion Γ_{endo} defined in the recursive competitive equilibrium of Definition 4, I define the conditional K -nullcline.

Definition 5 (Conditional K -nullcline in RBC).

Fix a regime $A \in \{B, G\}$. The conditional K -nullcline is the zero-net-investment locus in the (K, C) phase space,

$$\mathcal{N}(A) := \{(K, C) \in \mathbb{R}_{++} \times \mathbb{R}_{++} : \Delta K(K, C; A) = 0\}, \quad (15)$$

where $\Delta K(K, C; A)$ is the one-period change in capital implied by the resource constraint at (K, C) under regime A . When $\mathcal{N}(A)$ is the graph of a function, I write $\mathcal{N}(A) = \{(K, C) : C = C_A^{K\text{null}}(K)\}$.

Along a strictly monotone frozen-regime continuation, the orbit meets $\mathcal{N}(A)$ only in the limit, at the conditional steady state. The locus $\mathcal{N}(A)$ is therefore distinct from the realized nullcline crossings along the continuation, $\{(K_t, C_t) : t \geq 0, K_{t+1} = K_t\}$: for a

strictly monotone continuation, this set contains no point at all before the limit. In the RBC model, the aggregate resource constraint $Y = C + I$ with $I = K_{t+1} - (1 - \delta)K_t$ gives $\Delta K = AK^\alpha - \delta K - C$, so $\mathcal{N}(A)$ is the graph of an explicit function.

Proposition 1 (Characterizing the conditional K -nullcline in RBC).

The conditional nullclines of aggregate capital K for $A \in \{B, G\}$ are

$$C_A^{Knull}(K) = AK^\alpha - \delta K, \quad (16)$$

independent of any initial condition.

Proof. From $\Delta K(K, C; A) = 0$ and $\Delta K = AK^\alpha - \delta K - C$, the nullcline is $C = AK^\alpha - \delta K$ for each A . ■

The derivation invokes only $\Delta K = AK^\alpha - \delta K - C$ (the technology and the aggregate resource constraint), so the nullcline bears no imprint of preferences, beliefs, or the distribution. The same holds for the entire *direction field* it induces: at each (K, C) the sign of net investment, $\Delta K = C_A^{Knull}(K) - C$, is fixed by accounting alone. Beliefs and heterogeneity therefore relocate the equilibrium *orbit* within a phase portrait whose boundary and sign structure they cannot move. This is the common root of two results below: belief-invariance under aggregate uncertainty (Proposition 3), which underlies the perfect-foresight diagnostic of Section 4.1.1, and distribution-invariance in the heterogeneous-household model (Proposition 7), which underlies within-spell dimension reduction.

In discrete time, capital can jump by a finite amount in one period, so the projected (K, C) trajectory can develop folds or loops. These pathologies are absent in continuous time, where capital moves continuously and cannot skip past the $\dot{K} = 0$ locus. Appendix L develops the continuous-time formulation and shows that monotonicity holds without any segment-crossing restriction. In the discrete-time setting, I rule out these pathologies with the following geometric restriction.

Assumption 2 (No segment crossing of the conditional K -nullcline).

For each $A \in \{B, G\}$, along the frozen- A continuation: for every pre-limit transition (every $t \geq 0$ with $K_{t+1} \neq K_t$), the open straight line segment in (K, C) connecting (K_t, C_t) and (K_{t+1}, C_{t+1}) does not intersect the active regime's conditional K -nullcline $\mathcal{N}(A)$:

$$((K_t, C_t), (K_{t+1}, C_{t+1})) \cap \mathcal{N}(A) = \emptyset, \quad \forall t \geq 0 \text{ with } K_{t+1} \neq K_t. \quad (17)$$

Endpoints are not required to lie off the nullcline; in particular, the conditional steady state is allowed to lie on $\mathcal{N}(A)$, and the assumption imposes no condition at the degenerate terminal transition $K_{t+1} = K_t$.

Assumption 2 ensures consecutive points cannot strictly switch sides of the K -nullcline. Its companion condition is *unique K -stationarity*: $K_{t+1} = K_t$ only at $\Phi_t = \Phi^{cs}(\Phi_0, A)$. In the representative-agent model this follows from Assumption 1(i) (K is Markov under the frozen regime); otherwise it is maintained as a pathwise hypothesis. Together the two conditions deliver monotonicity in three short steps. First, net investment equals minus the nullcline gap: writing $H_t := C_t - C_A^{Knull}(K_t)$, the resource constraint gives $\Delta K_t = -H_t$. Second, the gap cannot change sign: $H_t H_{t+1} \geq 0$ for every transition t with $K_{t+1} \neq K_t$, and $H_t \neq 0$ on any such transition (since $\Delta K_t = -H_t$); $H_{t+1} = 0$ can occur only when the transition lands on a flat step, which unique K -stationarity confines to terminal arrival at the conditional steady state. Third, net investment therefore keeps a constant sign along nondegenerate transitions, so capital moves one way until it stops at the conditional steady state: strict K -monotonicity up to convergence. Economically, Assumption 2 rules out single-period overshooting so dramatic that capital jumps past its stationary level and reverses direction: the discrete-time analogue of the continuous-time property that trajectories cannot cross nullclines.⁷

Although Assumption 2 is stated here for the RBC benchmark, it plays a central role in the heterogeneous-household economy (Section 4.1). The assumption can be verified from model primitives without solving the full stochastic equilibrium, conditional on a simulated frozen-regime continuation (or certified envelope bounds on the relevant moments along it); Sections 4.1.1 and 4.1.2 develop two complementary strategies.

Proposition 2 (K -monotonicity in RBC).

Fix $A \in \{B, G\}$. Under Assumptions 1 and 2, together with unique K -stationarity ($K_{t+1} = K_t$ only at $\Phi_t = \Phi^{cs}(\Phi_0, A)$ along the frozen- A continuation), aggregate capital converges to K_A^{cs} strictly monotonically.

Proof. Fix A and let $H_t := C_t - C_A^{Knull}(K_t)$ denote the nullcline gap. The resource constraint gives $\Delta K_t = -H_t$. By Assumption 2, no transition strictly switches sides of the nullcline: $H_t H_{t+1} \geq 0$ for every transition t with $K_{t+1} \neq K_t$, and $H_t \neq 0$ on any such transition. Unique K -stationarity confines $H_{t+1} = 0$ to terminal arrival at Φ_A^{cs} , so prior to convergence $\{\Delta K_t\}$ has constant nonzero sign. Unique K -stationarity rules out nonterminal flat steps ($\Delta K_t = 0$ at any $\Phi_t \neq \Phi_A^{cs}$), so every pre-convergence transition is nondegenerate and $\{K_t\}$ is strictly monotone. Convergence $K_t \rightarrow K_A^{cs}$ follows from Assumption 1. ■

Unique K -stationarity as a no-plateau condition Economically, unique K -stationarity rules out *plateau equilibria*: distributional states off the conditional steady state at which

⁷For related geometric restrictions in dynamic optimization, see Reddy et al. (2020) and, on threshold (Skiba) phenomena, Skiba (1978), Dechert and Nishimura (1983), and Wagener (2003).

aggregate capital stops moving ($\Delta K_t = 0$) while the underlying distribution continues to evolve. It excludes the kind of degenerate fixed points of K that arise in models with poverty traps, occasionally binding constraints that fully bind over a range, or policy kinks that neutralize net investment at non-steady-state capital levels. In the representative-agent RBC, K is Markov under the frozen regime, so unique K -stationarity reduces to uniqueness of the conditional steady state (Assumption 1(i)). In the heterogeneous-agent economy, it holds locally via the stable eigenvalue $\lambda_A^s \in (0, 1)$ of the linearized frozen-regime dynamics (Proposition 10). A globally certified version is more delicate: the bounded price-income MPC of Assumption 3 (Appendix G) constrains the dynamics off the conditional steady state, but I treat unique K -stationarity as a maintained pathwise condition, directly checkable along any simulated frozen-regime continuation by inspecting whether $\Delta K_t = 0$ occurs at any point before convergence to the conditional steady state.

In the representative-agent benchmark these geometric restrictions are not free-standing assumptions but consequences of the scalar concave structure, given two explicitly checkable inputs: a curvature-margin bound and, for stochastic beliefs, a single-crossing condition on the conditional Euler residual, both verified in the computed equilibrium. Because capital is the autonomous scalar state, the frozen- A continuation is the iteration of a single recursive policy $g_A(K)$. Each restriction then follows from a primitive of the benchmark:

1. strict concavity of the one-sector problem delivers a unique conditional steady state (under deterministic beliefs; the stochastic-beliefs case adds the verified single-crossing input, on the capital range over which it is verified);
2. the same strict concavity delivers strictly monotone convergence to that steady state, with the same qualification;
3. the same strict concavity, under an explicit curvature-margin bound satisfied with wide slack in standard calibrations, delivers no segment crossing of the nullcline;
4. feasibility bounds capital, so explosion never arises;
5. the transversality condition enters only in selecting the stable policy.


Appendix N gives the proof. The conditions are stated as assumptions because the heterogeneous-agent and search economies of Sections 4.1–4.2 replace this scalar map by an operator on the distribution. There K is no longer autonomous, and “monotone map with a real stable eigenvalue” becomes a spectral property (a statement about the eigenvalues of the linearized dynamics on the distribution space) that cannot be read off primitives. Finally,

the non-concave illustration of Appendix G shows the concavity used in the benchmark is essential, not decorative.

Figure 1 depicts the conditional saddle paths of the calibrated RBC model in the (K, C) plane under $A \in \{B, G\}$, solved globally via the repeated transition method of Lee (2026a).⁸ Panel (a) zooms into the conditional steady states and overlays a histogram of the simulated capital outcomes K_t ; panel (b) shows the global partition with conditional K -nullclines and direction arrows. In the long run, all equilibrium dynamics are confined to the region bounded by the two conditional saddles: within each spell, the economy converges along the active saddle, and regime switches move it across, so the entire ergodic distribution is nested within the saddle boundaries.

MIT shock on the saddle This decomposition yields a sharp characterization of unexpected transitory shocks (MIT shocks), whose connection to first-order perturbation in heterogeneous-agent models is established by Boppart et al. (2018). I restrict attention here to shocks representable as displacements of the predetermined state under unchanged equilibrium decision rules. For this class of shocks, there exists a shift along the conditional saddle path that replicates the shock’s impact, and all subsequent dynamics are along-path convergence. Consequently, the slope and curvature of the conditional policy graph at the point of displacement fully determine the economy’s propagation of the shock, including its persistence, amplification, and state dependence.

Figure 1: Conditional saddle paths and the phase diagram



Notes: Conditional K -nullclines ($\Delta K = 0$) partition the (K, C) plane. Direction arrows show frozen-regime convergence along each conditional saddle. Panel (a) zooms into the neighborhood of the conditional steady states shown in panel (b). The histogram plots the time-series distribution of aggregate capital.

⁸For each regime A , I hold productivity fixed for 2,500 periods and iterate the equilibrium law of motion from initial capital stocks on both sides of K_A^{CS} , recovering the full conditional saddle in the phase diagram.

Generalized transition function The conditional saddle paths capture all equilibrium transitions, along saddles (endogenous propagation) and across them (exogenous regime shifts), enabling generalized transition function (GTF) analysis (Lee, 2026a), which encompasses generalized impulse response functions (Koop et al., 1996; Andreasen et al., 2017). The economy’s response to a sequence of exogenous shocks is not confined to local dynamics around a steady state: its trajectory is sharply traced in the phase diagram as a sequence of along-saddle and across-saddle movements.

Figure 2 illustrates the equilibrium dynamics in the phase diagram (panel (a)) and in the time domain (panels (b) and (c)) for a 20-quarter subsample. Negative TFP shocks produce downward jumps in consumption across conditional saddle paths, followed by endogenous decline along them; positive shocks reverse the pattern. Compared to time-domain representations, the phase diagram illustrates the full equilibrium dynamics more concisely in a single figure and enables immediate consideration of counterfactual scenarios. For the GIRF illustration, I replace the two-state Markov chain with a finely discretized AR(1) process for log productivity: $\ln A' = \rho \ln A + \epsilon$, $\epsilon \sim N(0, \sigma^2)$. In this case, there is effectively a continuum of frozen- A conditional saddles indexed by A , and the GIRF path moves both along a given conditional saddle (endogenous propagation) and across the family of saddles as A evolves.

Figures 3 and 4 illustrate post-shock dynamics following a positive TFP shock. Upon impact, consumption jumps upward across conditional saddles. In subsequent periods, forces along and across saddles jointly generate a bow-shaped trajectory in the phase diagram: capital exhibits a hump-shaped response as the shock mean-reverts, while consumption initially rises then falls as TFP crosses the threshold where the direction of along-saddle adjustment reverses.

Comparison with the perfect-foresight saddles In the *perfect-foresight* (PF) economy, agents believe the current regime persists forever; the PF steady states solve $1 = \beta(1 - \delta + \alpha A(K_A^{pf})^{\alpha-1})$. Along the *conditional saddle* (CS), the regime is held fixed ex post, but agents correctly anticipate regime switches under Γ_{exo} , so the conditional steady states K_A^{cs} are pinned by the probability-weighted Euler equations.⁹ PF and CS economies share the same K -nullcline (Proposition 3), but differ in their consumption dynamics.

Proposition 3 (K -nullcline invariance over beliefs).

Conditional K -nullclines are identical between the RBC model with aggregate uncertainty and the perfect-foresight counterpart.

⁹Specifically, $\beta\pi_{AA}(1 + \alpha A(K_A^{cs})^{\alpha-1} - \delta) + \beta\pi_{AA'}(c(K_A^{cs}, A')/c_A^{cs})^{-1}(1 + \alpha A'(K_A^{cs})^{\alpha-1} - \delta) = 1$ for each $A \in \{B, G\}$, where $A' \neq A$.




Figure 2: Equilibrium dynamics: saddle vs. time domain

Notes: Stochastic equilibrium dynamics on the conditional saddle paths (panel (a)) and in the time domain (panels (b) and (c)). A 20-quarter window from the simulated path (quarters 830–849).

Proof.

From the stationary condition $\delta K = I$ and the aggregate resource constraint $Y = C + I$, equation (16) is immediate. Neither condition involves expectations or belief structures, so the nullcline is identical whether or not there is aggregate uncertainty. ■

This invariance implies that PF–CS differences are driven entirely by consumption dynamics, not by the feasibility locus.¹⁰ Proposition 4 formalizes this via steady-state orderings. The same invariance extends to the heterogeneous-household environment (Section 4.1), where it underlies both K -sufficiency and the perfect-foresight diagnostic (Section 4.1.1).

¹⁰There is no consumption nullcline under aggregate uncertainty; the Euler equations at the conditional steady states play that role.




Figure 3: Generalized impulse response function (GIRF) at (K_0, A_0)

Notes: Generalized impulse responses of consumption and capital to a positive TFP shock at initial state (K_0, A_0) , where A_0 is the pre-shock productivity level. Upon impact, the economy jumps across conditional saddles to the conditional saddle indexed by $A_0 + \Delta$ and subsequently propagates along and across saddles as TFP mean-reverts.

Proposition 4 (Aggregate uncertainty and the conditional steady states).

Let $R(A, K) := 1 - \delta + \alpha AK^{\alpha-1}$ denote the gross return and $c(K, A')$ the stochastic-equilibrium consumption function. Suppose that (i) at each regime's PF steady state, the consumption response to a regime switch dominates the return response (jump dominance),

$$\frac{c(K_B^{pf}, G)}{c_B^{pf}} > \frac{R(G, K_B^{pf})}{R(B, K_B^{pf})}, \quad \frac{c(K_G^{pf}, B)}{c_G^{pf}} < \frac{R(B, K_G^{pf})}{R(G, K_G^{pf})}; \quad (18)$$

and (ii) each conditional Euler residual F_A , the residual of the conditional steady-state Euler equation viewed as a function of K (Appendix C), is strictly positive below K_A^{cs} and strictly negative above it. Then

$$K_B^{cs} < K_B^{pf} < K_G^{pf} < K_G^{cs}, \quad c_B^{cs} < c_B^{pf} < c_G^{pf} < c_G^{cs}. \quad (19)$$

Proof.

See Appendix C. ■

Both hypotheses are explicit and directly checkable on the computed equilibrium. In words, a persistent regime switch barely moves the gross return, because only the marginal-product term shifts, while consumption moves roughly one-for-one with the switch. Formally, the return ratio in (18) deviates from one only at first order in the regime gap and with a small coefficient, $R(G, K)/R(B, K) = 1 + (G - B) \alpha K^{\alpha-1}/R(B, K)$. Equilibrium consumption, by




Figure 4: Impulse responses to a positive TFP shock: saddle vs. time domain

Notes: Impulse responses on conditional saddle paths (panel (a)) and in the time domain (panels (b) and (c)). Sample: 50 quarters.

contrast, responds to a persistent regime switch at first order with an $O(1)$ coefficient, one that does not shrink as the regime gap closes. Condition (i) therefore fails only in degenerate parameterizations. Indeed, the conditions are tight. Under full depreciation ($\delta = 1$), decision rules collapse to the belief-independent Brock–Mirman form. In that corner, condition (i) holds with equality and the conditional and PF steady states exactly coincide, so the strict ordering cannot hold unconditionally (Appendix C). Away from this corner, the orderings are confirmed in the computed equilibrium of Figure 5.

PF steady states are nested within the CS steady states: hedging against regime switches pushes the conditional steady states apart (Deaton, 1991; Carroll, 1997). Conversely, the conditional saddle *paths* are nested within the PF paths: at any given K , hedging compresses the consumption gap across regimes. The two nestings reflect the same mechanism: compressed

consumption responses cumulate into amplified long-run capital divergence (Figure 5).




Figure 5: Conditional saddle path comparison: with and without uncertainty

Notes: Conditional saddle paths for $A = B$ (solid) and $A = G$ (dashed) and perfect-foresight saddle paths for $A = B$ (dash-dotted) and $A = G$ (dotted) under standard quarterly calibration.

Conditional boundary condition I work directly with the recursive formulation, which implicitly imposes a stochastic analogue of the deterministic transversality condition.¹¹ Given conditional saddle paths, the natural counterpart is *conditional transversality*:

$$\lim_{t \rightarrow \infty} \beta^t u'(c_t) K_t = 0 \quad \text{under a frozen aggregate regime } A \in \{B, G\} \quad (20)$$

(along the conditional saddle for each $A \in \{B, G\}$).

Condition (20) rules out explosive continuations along each conditional saddle. Under standard assumptions, it is equivalent to the usual stochastic transversality conditions (Kamihi-gashi, 2003, 2005) and selects the same bounded RCE. The conditional formulation makes explicit that the boundary condition operates regime-by-regime on $\mathcal{M}(\Phi_0, A)$.

Economies without conditional saddle paths A conditional saddle requires a predetermined endogenous state. Without one, the economy jumps each period to the allocation implied by the current exogenous state, so the conditional “saddle” under frozen A collapses

¹¹This condition is imposed at the aggregate level and should be distinguished from individual no-Ponzi constraints.

to a singleton. I call such an economy *endogenously stateless*.¹² Models with predetermined states (capital, habits, interest-rate smoothing, or distributional states in HANK) are endogenously stateful, and conditional saddle paths exist under mild regularity conditions.

3 State dependence

In this section, I analyze nonlinear shock responsiveness through the lens of conditional saddle paths. In any stochastic dynamic model that admits conditional saddle paths, the response of an aggregate variable to an exogenous shock is represented by a shift across different conditional policy graphs. The natural benchmark is *state independence*: the impact gap between policy graphs is constant in the endogenous coordinate. Departures from this benchmark show up locally as slope differences between the graphs, which serve as the first-order diagnostic of state-dependent responses.

A large empirical literature documents that macroeconomic responses to the same impulse vary with the phase of the business cycle: unemployment dynamics are asymmetric (Neftci, 1984), fiscal multipliers are larger in recessions (Auerbach and Gorodnichenko, 2012), and monetary policy is less powerful in recessions (Tenreyro and Thwaites, 2016). The following result formalizes the benchmark and its local failure.

Lemma 1 (Vertical-separation benchmark).

Let I be an interval over which the regime-specific conditional policy graphs $\mathcal{P}_{A_j}^g = \{(e, g_j(e)) : e \in I\}$, $j \in \{0, 1\}$, are single-valued. Define the impact gap between the two graphs as $\Delta_g(e; A_1, A_0) := g_1(e) - g_0(e)$. Then conditions (i) and (ii) below are equivalent; if g_0, g_1 are differentiable on I , all three are equivalent:

- (i) Δ_g is constant on I ;
- (ii) $\mathcal{P}_{A_1}^g$ is a vertical translation of $\mathcal{P}_{A_0}^g$;
- (iii) the slope gap $g'_1 - g'_0$ vanishes on I .

Proof. The first equivalence is definitional; the second is the mean value theorem applied to Δ_g . ■

Lemma 1 is the elementary benchmark; the substantive statement, that the slope gap reflects model primitives, is Proposition 5 below. Sharp state independence is a knife-edge.

¹²The textbook three-equation New Keynesian model (Gali, 2008) under determinacy is endogenously stateless: $(x_t, \pi_t, i_t) = \Psi s_t$ for exogenous state $s_t := (r_t^n, u_t, \nu_t)$, so freezing s_t produces no nontrivial dynamics.

Differential tilt across conditional policy graphs—whether from real (Winberry, 2021; Lee, 2026a), financial (Melcangi, 2024), or labor market frictions (Petrosky-Nadeau et al., 2018; Pizzinelli et al., 2020)—generates state-dependent shock responses. Lemma 1 provides the qualitative characterization; the following proposition translates it into a structural diagnostic computable from model primitives. In the log-linearized frozen-regime RBC dynamics near K_A^{cs} , let $R_A^{cs} := 1 + f'(K_A^{cs}; A) - \delta$ denote the gross return at the conditional steady state, $\lambda_A^s \in (0, 1)$ the stable eigenvalue, and $\psi_A := dC/dK|_{\text{saddle}, A}$ the level saddle slope at K_A^{cs} .

Proposition 5 (Slope ratio as a structural diagnostic).

The level saddle slope admits the closed form

$$\psi_A = R_A^{cs} - \lambda_A^s, \quad (21)$$

and hence the cross-regime slope ratio is

$$\mathcal{R} := \frac{\psi_B}{\psi_G} = \frac{R_B^{cs} - \lambda_B^s}{R_G^{cs} - \lambda_G^s}. \quad (22)$$

Proof. Log-linearize the frozen-regime budget constraint $K_{t+1} = f(K_t; A) + (1 - \delta)K_t - C_t$ around (K_A^{cs}, C_A^{cs}) :

$$\hat{k}_{t+1} = R_A^{cs} \hat{k}_t - \frac{C_A^{cs}}{K_A^{cs}} \hat{c}_t. \quad (23)$$

On the stable manifold (the saddle-path branch of the linearized dynamics), $\hat{c}_t = \psi_A^{\log} \hat{k}_t$ for some log slope $\psi_A^{\log}$. Substituting into (23) yields $\hat{k}_{t+1} = [R_A^{cs} - \frac{C_A^{cs}}{K_A^{cs}} \psi_A^{\log}] \hat{k}_t$, which by definition equals $\lambda_A^s \hat{k}_t$. Solving, $\psi_A^{\log} = (K_A^{cs}/C_A^{cs})(R_A^{cs} - \lambda_A^s)$. Converting from log to level units via $dC/dK|_{K^{cs}} = (C_A^{cs}/K_A^{cs}) \psi_A^{\log}$ gives (21). Equation (22) follows by taking the ratio.

Under perfect-foresight expectations, the steady-state Euler condition gives $R_A^{cs} = 1/\beta$ and the formula reduces to $\psi_A = 1/\beta - \lambda_A^s$. Under stochastic expectations over regime transitions, R_A^{cs} is pinned down by $1 = \beta \mathbb{E}_A[(C_A^{cs}/C(\Phi_A^{cs}, A'))^\sigma (1 + f'(K_A^{cs}; A') - \delta)]$ (the return is evaluated at tomorrow's capital K_A^{cs} under the realized regime A') and generally differs from $1/\beta$; the formula remains valid with R_A^{cs} evaluated at the (stochastic) conditional steady state. ■

In log units the slope is $(K_A^{cs}/C_A^{cs}) \psi_A$; both R_A^{cs} and λ_A^s are computable from a local linearization at the regime-specific conditional steady states, without solving the full global policy. In plain words, the conditional saddle path is a discrete set of points, not a smooth curve. Near the conditional steady state, however, these points can be treated as lying on a smooth curve, and the slope and tangent language I use refers to that neighborhood

alone. Formally, Definition 1 introduces the conditional saddle path $\mathcal{M}(\Phi_0, A)$ as an orbit closure (generically a countable orbit together with its limit point). The slope, curvature, and “tangent” language used in Proposition 5 (and subsequently in Proposition 6, the tangent-coherence discussion of Section 4.1, and Proposition 10) refers to the local C^1 parametrization of $\mathcal{M}(\Phi_0, A)$ near K_A^{cs} . This parametrization exists under the regularity conditions of Proposition 10 (Appendix F). Outside this neighborhood, the orbit is treated as the discrete object on which Theorem 1 guarantees injectivity: the orbit never revisits a capital level.

Proposition 5 gives the slope ratio $\mathcal{R}$ a structural interpretation: each regime’s saddle slope is the gap between the gross return on capital at that regime’s conditional steady state and the corresponding stable eigenvalue. Local departures from $\mathcal{R} = 1$ are the geometric signature of state-dependent impact responses, and can be computed from a local computation, without solving the full global policy: one linearizes the frozen-regime equilibrium dynamics around each conditional steady state, reads off (R_A^{cs}, λ_A^s) , and applies (22).

This diagnostic connects the geometric framework to the empirical state-dependence literature: $\mathcal{R}$ is a model-implied moment that could in principle be identified from state-dependent local projections (Auerbach and Gorodnichenko, 2012; Tenreyro and Thwaites, 2016).

A minimal example: asymmetric adjustment cost I modify the budget constraint by adding $\mathcal{C}(K', K) = \frac{\tilde{\mu}}{2}((K' - K)/K)^2 K$, with $\tilde{\mu} = \mu_+$ when $K' - K \geq 0$ and $\tilde{\mu} = \mu_-$ otherwise. Thus the curvature changes with the sign of *net capital accumulation*, not with the sign of gross investment $I = K' - (1 - \delta)K$. The specification is C^1 at $K' = K$ but has different second derivatives on the two sides. I interpret $\mu_+ > \mu_-$ as a reduced-form bottleneck to expanding installed capacity—installation, ramp-up, or congestion—rather than as a resale wedge.¹³

This direction and threshold should be distinguished from costly reversibility. Abel and Eberly (1996) formalize costly reversibility as a purchase price above the resale price, and Cooper and Haltiwanger (2006) estimate a richer mixture of convex, nonconvex, and resale-price adjustment costs. Zhang (2005) instead uses asymmetric quadratic curvature with $\theta^- > \theta^+$ and switches at gross investment $I = 0$. Because $K' = I + (1 - \delta)K$, Zhang’s threshold is $K' = (1 - \delta)K$, whereas mine is $K' = K$ and puts high curvature on the opposite

¹³The calibration maintains the arithmetic-mean curvature $\bar{\mu} = (\mu_+ + \mu_-)/2 = 3.05$ and sets the ratio $r = \mu_+/\mu_- = 3$, hence $(\mu_-, \mu_+) = (1.53, 4.57)$. The remaining annual parameters are $\beta = 0.96$, $\delta = 0.10$, and $\alpha = 0.33$; $A \in \{0.98, 1.02\}$ has transition matrix with diagonal entries 0.90. Appendix ?? details the calibration: the ratio is disciplined by an outlier-robust asymmetry moment (the downside/upside semideviation ratio of investment growth) rather than the raw skewness, and $\mathcal{R}$ is reported across the robustness band $r \in [2, 4]$. This is a conditional discipline on the asymmetry plus an external anchor on the level, not a structural estimate of the two curvature parameters.

side. The present specification is a parsimonious laboratory for how an asymmetric real friction tilts equilibrium policy graphs, not a structural model of capital resale.

All same-capital comparisons below use interior common-support domains. Within a given specification, define the pointwise diagnostic $\mathcal{R}(K) := C'_B(K)/C'_G(K)$ at the same K over the intersection of the regimes' 5th–95th percentile policy supports. Under the calibrated $r = 3$ specification, its median is 1.265. Holding $\bar{\mu} = 3.05$ fixed, equal curvature gives 0.976, while reversing the pair to $(\mu_-, \mu_+) = (4.57, 1.53)$ gives 0.753. The sign reversal therefore reverses the orientation of the saddle fan. I do not report the earlier conditional-steady-state estimate of 1.83: neither regime-specific fixed point lies inside the simulated support, so that number was generated by clamped extrapolation rather than an interior policy derivative.

Figure 6 examines whether this policy tilt is robust to preferences. Every case keeps $\mu_+/\mu_- = 3$, equivalently $(\mu_-, \mu_+) = (1.53, 4.57)$ —a 1:3 contraction-to-expansion curvature ratio at the maintained mean curvature $\bar{\mu} = 3.05$ —as well as the same technology, shock process, and simulated shock path. Its left column reports the consumption saddles $C_A(K)$, and its right column reports marginal $q_A^K(K) = 1 + \mathcal{C}_{K'} = 1 + \tilde{\mu}(K' - K)/K$. The top row varies the CRRA coefficient γ ; under CRRA this changes relative risk aversion and the EIS, $1/\gamma$, jointly. As γ rises from 1 to 10, the median consumption-slope ratio C'_B/C'_G falls from 1.265 to 1.047. Marginal q exhibits a useful distinction between within-regime tilt and across-regime dispersion. Both q saddles flatten, so the median absolute slope gap $|q_G^{K'} - q_B^{K'}|$ falls from 0.0577 to 0.0307. Yet their conditional level gap, $\mathbb{E}[q^K | G] - \mathbb{E}[q^K | B]$, rises from 0.0153 to 0.0315, and the simulated unconditional $\text{sd}(q^K)$ rises from 0.00932 to 0.01723. Thus q becomes less sensitive to capital position within a regime but more volatile because regime switches traverse more widely separated saddles.

The bottom row provides the controlled preference comparison. Epstein–Zin utility holds the EIS fixed at one while relative risk aversion rises from 1 to 10. The curves then barely move: the median consumption-slope ratio remains approximately 1.265, the absolute marginal- q slope gap changes only from 0.0577 to 0.0576, and $\text{sd}(q^K)$ changes only from 0.00932 to 0.00933. The strong CRRA comparative statics therefore come predominantly from the accompanying fall in the EIS, not from higher risk aversion by itself.

The preference exercise uses the same friction parameters as the empirical calibration of Appendix ???. I maintain the mean curvature $\bar{\mu} = 3.05$ and choose the ratio $r = \mu_+/\mu_-$ to match the asymmetry of annual investment growth over 1967–2019.¹⁴ Because the third-moment skewness of $\Delta \log(\text{PNFICA}_t)$ is dominated by the 2009 collapse in a short annual sample—it falls from -1.03 to -0.54 once 2009 is excluded, and is -0.32 in the broader

¹⁴The investment series is real private nonresidential fixed investment (PNFICA), downloaded through FRED.




Figure 6: Consumption and marginal- q saddles across preferences

Notes: Panels (a) and (c) plot the bad- and good-regime consumption policy $C_A(K)$; panels (b) and (d) plot marginal $q_A^K(K)$. Colors and markers denote relative risk aversion; solid and dashed lines denote the bad and good regimes. Panels (a) and (b) vary relative risk aversion under CRRA, which simultaneously changes the EIS because $EIS = 1/\gamma$. Panels (c) and (d) use Epstein–Zin utility to vary relative risk aversion while holding the EIS fixed at one; their near overlap is the substantive result that risk aversion alone has little effect on these saddles in this calibration. Every panel retains $\mu_+/\mu_- = 3$, equivalently $(\mu_-, \mu_+) = (1.53, 4.57)$, a 1:3 contraction-to-expansion curvature ratio at $\bar{\mu} = 3.05$. Every curve is evaluated at the same K on the intersection of the 5th–95th percentile solved supports across all regime–preference cases. Horizontal lines in the marginal- q panels mark $q^K = 1$, or zero net capital accumulation. Marginal q directly displays the assumed adjustment wedge and is not independent calibration evidence. Same-mean symmetric and sign-reversed cases remain in Table 1 and the replication outputs.

gross-private-investment series—I target the outlier-robust downside/upside semideviation ratio, which is 1.32 in the data and rises monotonically with r in the model. This yields $r = 3$, hence $(\mu_-, \mu_+) = (1.53, 4.57)$; the exercise is a conditional discipline on the asymmetry, not a structural estimate of either curvature parameter, and $\mathcal{R}$ is robust across $r \in [2, 4]$.

Relative investment volatility is reported only as an untargeted diagnostic. As Appendix ?? explains, the two-state $\pm 2\%$ TFP process caps the model’s $\text{sd}(I^{cyc})/\text{sd}(Y^{cyc})$ below the data value even at negligible friction, so it cannot identify the level; the level is

anchored instead on the investment– q elasticity, and the calibration matches the asymmetry of investment growth rather than its cyclical amplitude. At the same mean curvature, the symmetric and sign-reversed specifications give median same-capital slope ratios of 0.976 and 0.753, compared with 1.265 under the targeted specification, so the orientation of the asymmetric friction drives the saddle fan.

Table 1: Asymmetric adjustment costs: robust-asymmetry calibration and same-capital slopes

Case	μ_-	μ_+	$\sigma(I^{cyc})/\sigma(Y^{cyc})$	σ_-/σ_+	skew($\Delta \log I$)	med $\mathcal{R}(K)$	$\mathcal{C}/Y$
Data	–	–	3.21	1.32	-1.03	–	–
Targeted ($r = 3$)	1.53	4.57	1.85	1.14	-0.63	1.265	0.0032%
Same-mean symmetric ($r = 1$)	3.05	3.05	1.79	1.01	-0.07	0.976	0.0033%
Same-mean sign reversal ($r = 1/3$)	4.57	1.53	1.81	0.90	0.51	0.753	0.0032%

Notes: The asymmetry ratio $r = 3$ is disciplined by the outlier-robust downside/upside semideviation ratio of $\Delta \log(\text{PNFICA})$ over 1967–2019 (Appendix ??); the moment skewness is reported for reference only, as it is fragile to the 2009 outlier (it falls from -1.03 to -0.54 once 2009 is excluded). The exercise is a conditional discipline on the asymmetry, not a structural estimate of the two curvatures. I^{cyc} and Y^{cyc} are annual HP-filtered log investment and output with $\lambda = 6.25$; their relative volatility is untargeted and cannot identify the level (Appendix ??). $\mathcal{R}(K) = C'_B(K)/C'_G(K)$ is summarized by its median on the common interior capital support. $\mathcal{C}/Y$ is the mean resource cost.

The more direct external motivation is state-dependent investment responsiveness. Aggregate investment responds more strongly in expansions in [Bachmann et al. \(2013\)](#), and [Winberry \(2021\)](#) obtains procyclical responsiveness through the extensive margin. These findings motivate an untargeted qualitative state-dependence comparison: $\mathcal{R}(K)$ is the paper’s model-side geometric diagnostic, not a directly comparable empirical elasticity. By Lemma 1, state-independent impacts correspond to constant vertical separation, while differential tilt reveals the endogenous-state dependence generated by the friction.

Having used the asymmetric friction to isolate slope-driven state dependence, I now return to the frictionless RBC benchmark to study a distinct channel: variation in response shape generated by initial position along the saddle.

Response shape in the frictionless benchmark Although $\mathcal{R} \approx 1$ in this benchmark, the *shape* of the response trajectory is highly state dependent. A positive productivity shock displaces the economy across saddles, and the subsequent trajectory depends on K_0 ’s position relative to K_G^{cs} . Far below K_G^{cs} , the marginal product of capital is high, tilting the intertemporal tradeoff toward accumulation and producing a pronounced, late consumption hump ([Barro and Sala-i-Martin, 1992](#)). Near K_G^{cs} , there is little room to accumulate; the across-saddle reversion dominates quickly and the peak arrives earlier. This channel is orthogonal to the slope asymmetry of Proposition 5. Slope differences govern

impact magnitudes (cross-saddle dimension); position-dependent convergence speed governs response *shapes* (along-saddle dimension).

Crucially, this state-dependent shape does *not* arise from nonlinearity of the saddle path itself. Even in the log-linearized economy, where the conditional saddle is a straight line, peak timing varies with K_0 . The mechanism is the *speed of along-saddle convergence*. Far below K_G^{cs} , the capital gap is large relative to the shock-induced across-saddle displacement, so the along-saddle force (capital accumulation) dominates the across-saddle force (TFP reversion) for many periods before the balance tips, producing a later and more pronounced consumption hump. The shape asymmetry is thus a property of the *race between two convergence rates* (the endogenous along-saddle rate λ and the exogenous across-saddle rate ρ) and how initial distance from steady state weights the former relative to the latter. Nonlinearity (concavity of $f(K)$) amplifies the effect quantitatively but is not its source.

The mechanism can be made precise via the log-linearized dynamics. Near the non-conditional steady state, the Blanchard–Kahn decomposition (Blanchard and Kahn, 1980) on the saddle branch yields

$$\hat{c}_t = \underbrace{\mathcal{P}(\hat{k}_0)}_{\text{along saddle}} \lambda^t + \underbrace{\mathcal{Q}}_{\text{across saddle}} \rho^t, \quad (24)$$

where $\hat{c}_t$ and $\hat{k}_0$ are log-deviations from steady state, $\lambda \in (0, 1)$ is the stable eigenvalue of the capital dynamics, and $\rho \in (0, 1)$ is the TFP persistence. The initial shock size $\varepsilon > 0$ does not appear in (24) itself; it enters through the coefficients $\mathcal{P}$ and $\mathcal{Q}$ defined below.¹⁵ The coefficient $\mathcal{P} = \psi_k[\hat{k}_0 + \theta\varepsilon/(\lambda - \rho)]$ depends on initial capital through the saddle slope $\psi_k > 0$ and the capital-TFP elasticity $\theta > 0$. The coefficient $\mathcal{Q} = [\psi_a - \psi_k\theta/(\lambda - \rho)]\varepsilon$ does not. The hump shape arises from the tension between two decaying exponentials. The λ^t term captures along-saddle propagation, the endogenous capital accumulation effect, while the ρ^t term captures the direct TFP reversion. The peak occurs where these forces balance. A lower initial capital stock reduces $\mathcal{P}(\hat{k}_0)$, weakening the slowly decaying along-saddle term relative to the TFP-reversion term, so the balance tips later in time: exactly the comparative static $\partial t^*/\partial \hat{k}_0 < 0$ established below.

Proposition 6 (State-dependent peak timing).

¹⁵This λ and the regime-specific λ_A^s of Proposition 5 are linearized at different steady states: λ comes from the linearization at the *non-conditional* steady state, while λ_A^s is linearized at the *conditional* steady state K_A^{cs} . In the frictionless RBC, to first order in the TFP perturbation, $\lambda_B^s = \lambda_G^s = \lambda$, consistent with $\mathcal{R} \approx 1$. Friction-induced departures from this equality are what generate $\mathcal{R} \neq 1$. Proposition 6 operates in the $\mathcal{R} \approx 1$ regime, so its state dependence is driven by initial location $\hat{k}_0$, not by regime-varying slopes. This is an orthogonal channel.

Suppose $\lambda > \rho$, $\mathcal{P}(\hat{k}_0) > 0$, and the consumption response (24) admits a peak at $t^* > 0$. Then

$$\frac{\partial t^*}{\partial \hat{k}_0} = \frac{-\psi_k}{\mathcal{P}(\hat{k}_0) \cdot \ln(\lambda/\rho)} < 0, \quad (25)$$

where $\hat{k}_0$ is the log-deviation of the initial capital stock from steady state.

Proof.

See Appendix D. ■

The consumption peak arrives strictly later when initial capital is farther below the steady state. The first two hypotheses are mild: $\lambda > \rho$ holds under standard quarterly RBC calibrations ($\lambda \approx 0.96$, $\rho \leq 0.95$), and $\mathcal{P}(\hat{k}_0) > 0$ requires only $\hat{k}_0 > -\theta\varepsilon/(\lambda - \rho)$ —initial capital not farther below steady state than the shock’s along-saddle displacement $\theta\varepsilon/(\lambda - \rho)$ —which the ergodic distribution satisfies with substantial margin. Proposition 6 is local by construction; Figure 7 confirms that the mechanism operates globally across the ergodic capital distribution.


Figure 7 confirms the logic of Proposition 6 quantitatively from the 20th and 80th percentiles of the ergodic capital distribution. From low K_0 , both C and K rise longer before reversion dominates, producing a later and larger consumption peak. From high K_0 , the peak arrives earlier: the same saddle geometry, the same shock, the same preferences; only the inherited capital stock differs.

4 Applications


The conditional-saddle framework applies across model classes, and the two applications below each exercise a different geometric property. In the heterogeneous-household economy (Section 4.1), the *monotonicity* of the saddle delivers dimension reduction; in the search-and-matching economy (Section 4.2), the *spacing* of its resting points delivers comparative statics. The spacing principle already surfaced in the asymmetric-adjustment-cost example of Section 3, where preference-driven separation of the price saddles governed asset-price volatility; it reappears here as policy-driven separation of the conditional steady states governing unemployment volatility.

4.1 Heterogeneous households: dimension reduction


In a heterogeneous-household economy, the natural state variable is the entire cross-sectional distribution of wealth, an infinite-dimensional object. Classical aggregation (Gorman, 1961)



(a) On the saddles



(b) C : time domain



(c) K : time domain

Figure 7: Shape of the impulse responses: low K_0 vs. high K_0

Notes: Impulse responses to a positive productivity shock from K_0 at the 20th percentile and 80th percentile of the ergodic distribution. Panel (a): post-shock trajectories in (K, C) space on the conditional saddle paths; numbers indicate quarters since impact. Panels (b) and (c): time-domain responses of consumption and capital; vertical markers indicate the peak quarter. Sample: 50 quarters.

reduces dimensionality by restricting preferences, but these restrictions fail under uninsurable idiosyncratic risk and borrowing constraints. Yet practitioners routinely approximate the state by a single number, aggregate capital. The conditional saddle framework explains why this works without preference restrictions: if the distribution evolves monotonically along a one-dimensional equilibrium path, then aggregate capital is an exact coordinate for the distribution along that path. This section formalizes the argument.

I consider a continuum of ex-ante homogeneous households of unit measure. The recursive

formulation of the households' problem is as follows:

$$v(a, z; \Phi, A) = \max_{c, a'} \log(c) + \beta \mathbb{E} v(a', z'; \Phi', A') \quad (26)$$

$$c + a' = a(1 + r(\Phi, A)) + w(\Phi, A)z \quad (27)$$

$$a' \geq 0 \quad (28)$$

$$\Phi' = \Gamma_{\text{endo}}(\Phi, A), \quad A' \sim \Gamma_{\text{exo}}(A'|A). \quad (29)$$

The problem is the same as in the representative-household economy except for (i) uninsurable idiosyncratic labor productivity, which follows a Markov process $z \sim \Gamma_z(z'|z)$; (ii) inclusion of the distribution of individual states Φ in the aggregate endogenous state; and (iii) the corresponding change in the law of motion for the aggregate endogenous state. The RCE is defined as in [Krusell and Smith \(1998\)](#).

The model features two exogenous stochastic processes: idiosyncratic productivity and aggregate TFP. Accordingly, there are two layers of conditional saddle paths: individual and aggregate. The individual saddle path has its own cross-sectional implications, which are analyzed separately in Appendix B. Here I focus on the aggregate conditional saddle path under stochastic TFP, following [Krusell and Smith \(1998\)](#).¹⁶ Consistent with the notation in Section 2.2, I denote $K = K(\Phi)$ and $C = C(\Phi, A)$ as aggregate capital and consumption.

The conditional K -nullcline of Definition 5 extends directly to the heterogeneous-household setting: although the endogenous state is now the distribution Φ , net investment depends on Φ only through the aggregates (K, C) , so the locus remains an object in the aggregate phase plane.

Definition 6 (Conditional K -nullcline in the heterogeneous-household setting).

Fix $A \in \{B, G\}$. The conditional K -nullcline is the zero-net-investment locus in the aggregate (K, C) plane,

$$\mathcal{N}(A) := \{(K, C) \in \mathbb{R}_{++} \times \mathbb{R}_{++} : \Delta K(K, C; A) = 0\},$$

where $\Delta K(K, C; A)$ is the one-period change in aggregate capital implied by the aggregate resource constraint at (K, C) under regime A . As in the representative-agent case, $\mathcal{N}(A)$ is distinct from the realized crossings $\{(K_t, C_t) : K_{t+1} = K_t\}$ along any frozen-regime continuation $\Phi_{t+1} = \Gamma_{\text{endo}}(\Phi_t, A)$, which under strict K -monotonicity contain no point at all before the limit.

[Krusell and Smith \(1998\)](#) posits an endogenous law of motion that tracks aggregate

¹⁶In the original model of [Krusell and Smith \(1998\)](#), exogenous individual labor supply co-moves with aggregate TFP. All results are unaffected by this feature, but for expositional brevity I assume labor supply is exogenously fixed.

capital K rather than the full distribution Φ : $\log K' = b_0(A) + b_1(A) \log K$, where b_0 and b_1 are state-dependent coefficients. This formulation embeds two distinct assumptions: (i) aggregate capital K suffices to summarize the endogenous aggregate state, and (ii) K follows a log-linear law of motion.

Since [Krusell and Smith \(1998\)](#), it has been an open question whether K -based forecasting works so well ($R^2 > 0.9999$) because of an accidentally accurate approximation or a deeper structural property.¹⁷ The conditional-saddle framework separates (i) into two questions of very different difficulty. *Within* a frozen-regime spell, K exactly indexes the distributional state on the conditional saddle, so conditioning on K loses no within-spell information; because the coordinate is strictly monotone along a single orbit, this indexing is immediate, and it holds path by path. The economically substantive question is the across-history one at the heart of the [Krusell and Smith \(1998\)](#) puzzle: why approximately the *same* rule $K' = f(K, A)$ works from every distributional state, not merely along one orbit. Proposition 9 addresses this by bounding the across-spell discrepancy under a uniform contraction condition on the frozen-regime dynamics: when that condition holds, the discrepancy decays with shared regime history and vanishes as regime persistence grows. Two caveats keep the claim honest. The contraction condition is maintained rather than derived from primitives, and the bound it delivers is loose in the standard calibration (detailed below); the near-exact accuracy itself is documented directly, by comparing the globally solved K' against the (K, A) -based prediction at spell boundaries (an average discrepancy of 2.5×10^{-4} in $\log K$). The synchronization theorem thus identifies the mechanism that *would* make one rule accurate across histories; the direct check confirms the accuracy is present in this calibration. The same geometry indicates where approximate aggregation should strain: when nonconvexities flatten the nulleline gap or generate multiple K -stationary points—as in lumpy-investment economies where the cross-sectional distribution carries independent forecasting power ([Khan and Thomas, 2008](#); [Bachmann et al., 2013](#))—the monotonicity hypothesis weakens, and the framework predicts degraded scalar-state accuracy rather than assuming it away.

By contrast, (ii) remains a parametric approximation (the true law of motion for K is generally non-Markovian and need not be log-linear), but this is a matter of functional-form choice. The practical implication is twofold. Within a spell, where K is an exact sufficient statistic, accuracy improves through richer functional forms for the law of motion rather than through additional distributional moments. Across spells, where the law of motion for K is

¹⁷[Den Haan \(2010\)](#) shows that the R^2 of the forecasting regression can substantially overstate dynamic accuracy, because one-step-ahead fit masks the accumulation of errors over long simulated paths; his proposed accuracy test compares multi-period dynamic forecasts against the realized path. The across-spell handoff check reported below, comparing the globally solved K' against the (K, A) -based prediction at spell boundaries, is in the spirit of this stricter standard.

not Markovian, additional state variables can in principle sharpen cross-history forecasts; the potential gain is bounded by the across-spell discrepancy, which is small in the calibration examined here.

The following theorem provides the first step toward establishing the exact conditional sufficiency of K . The core insight is simple: if the economy converges along a one-dimensional equilibrium path and some aggregate variable moves strictly in one direction throughout, its current value tells you exactly where the economy is along that path. No two distinct distributional states can share the same capital level, because the economy passes through each level exactly once.

Theorem 1 (Monotone aggregate variable as a coordinate).

Fix (Φ_0, A_0) and $A \in \{B, G\}$. Let $\{\Phi_t\}_{t \geq 0}$ be the frozen-regime continuation $\Phi_{t+1} = \Gamma_{\text{endo}}(\Phi_t, A)$ and let $\Phi^{cs} := \lim_{t \rightarrow \infty} \Phi_t$. Let $e : \mathcal{M}(\Phi_0, A) \rightarrow \mathbb{R}$ be an aggregate equilibrium variable. Suppose that the scalar sequence $\{e(\Phi_t)\}_{t \geq 0}$ is strictly monotone and converges to $e(\Phi^{cs})$. Then e uniquely indexes states on the conditional saddle:

$$e(\Phi) = e(\Phi') \implies \Phi = \Phi' \quad \forall \Phi, \Phi' \in \mathcal{M}(\Phi_0, A). \quad (30)$$

Proof. The plan in words: we label each state on the conditional saddle by the time at which the path reaches it, with the limit point labeled ∞ ; strict monotonicity of $\{e(\Phi_t)\}$ makes these labels unambiguous. Formally, define $\iota_A : \mathbb{N} \cup \{\infty\} \rightarrow \mathcal{M}(\Phi_0, A)$ by $\iota_A(t) := \Phi_t$ for $t \in \mathbb{N}$ and $\iota_A(\infty) := \Phi^{cs}$. By definition of $\mathcal{M}(\Phi_0, A)$ and convergence $\Phi_t \rightarrow \Phi^{cs}$,

$$\mathcal{M}(\Phi_0, A) = \{\Phi_t : t \geq 0\} \cup \{\Phi^{cs}\}, \quad (31)$$

so ι_A is surjective (onto: every state on the conditional saddle is reached by some time index). Strict monotonicity of $\{e(\Phi_t)\}$ implies that $e(\Phi_t) \neq e(\Phi_\tau)$ for $t \neq \tau$, hence $\Phi_t \neq \Phi_\tau$ for $t \neq \tau$; therefore ι_A is injective (one-to-one) on $\mathbb{N}$. Moreover, since $\{e(\Phi_t)\}$ is strictly monotone and converges to $e(\Phi^{cs})$, it also holds that $e(\Phi_t) \neq e(\Phi^{cs})$ for all finite t , so ι_A is injective on $\mathbb{N} \cup \{\infty\}$. Thus ι_A is a bijection (one-to-one and onto) between $\mathbb{N} \cup \{\infty\}$ and $\mathcal{M}(\Phi_0, A)$.

Now let $\Phi, \Phi' \in \mathcal{M}(\Phi_0, A)$. Pick $s, s' \in \mathbb{N} \cup \{\infty\}$ such that $\Phi = \iota_A(s)$ and $\Phi' = \iota_A(s')$. If $e(\Phi) = e(\Phi')$, then $e(\iota_A(s)) = e(\iota_A(s'))$. Since $e \circ \iota_A$ is injective, it follows that $s = s'$, hence $\Phi = \Phi'$. Therefore e is injective on $\mathcal{M}(\Phi_0, A)$, and an inverse map φ_A exists on $e(\mathcal{M}(\Phi_0, A))$. ■

Consequently, there exists a unique inverse map $\varphi_A : e(\mathcal{M}(\Phi_0, A)) \rightarrow \mathcal{M}(\Phi_0, A)$ such that $\varphi_A(e(\Phi)) = \Phi$ on $\mathcal{M}(\Phi_0, A)$, and any equilibrium object restricted to $\mathcal{M}(\Phi_0, A)$ can be written as a (single-valued) function of the scalar coordinate e .

The geometric intuition is that the conditional saddle is a one-dimensional curve of equilibrium states, and e acts as a “progress meter” along this curve: it moves strictly in one direction and never reverses. If two distinct states $\Phi \neq \Phi'$ shared the same value $e(\Phi) = e(\Phi')$, both would be at the same “progress” position, yet their subsequent histories could not merge on a one-dimensional path, a contradiction. Hence e uniquely labels states on the conditional saddle.

The following proposition shows that the conditional K -nullclines coincide with those of the representative-agent model.

Proposition 7 (Conditional K -nullclines of Krusell and Smith (1998)).

The conditional K -nullclines of the heterogeneous-household model are identical to those of the representative-household model and invariant over the initial distribution Φ_0 :

$$C_A^{Knull}(K) = AK^\alpha - \delta K, \text{ for } A \in \{B, G\}. \quad (32)$$

Proof.

As in Proposition 3, the stationary condition $\delta K = I$ and the aggregate resource constraint $Y = C + I$ immediately imply the form of the conditional nullclines. ■

As anticipated after Proposition 1, this is the distribution-invariance instance of that common root, and it is what enables dimension reduction: the boundary governing the sign of capital accumulation, an accounting identity indifferent to *who* saves, does not move as the distribution evolves along the conditional saddle.

The invariance is specific to the neoclassical production structure with exogenous labor supply. Under separable CRRA preferences with wealth effects on labor, labor supply introduces a distributional channel into the resource constraint. The nullcline then acquires distributional dependence, blocking the analytical route to K -sufficiency (though the perfect-foresight diagnostic of Section 4.1.1 remains applicable).

Proposition 8 (K monotonicity and injectivity).

Fix $A \in \{B, G\}$. Under Assumptions 1 and 2, together with unique K -stationarity along the frozen- A continuation (the no-plateau condition of Section 2.2), aggregate capital converges to K_A^{cs} strictly monotonically on $\mathcal{M}(\Phi_0, A)$. Consequently, K is injective on $\mathcal{M}(\Phi_0, A)$.

Proof.

By Proposition 7, the conditional K -nullcline takes the same form as in the RBC benchmark. The proof of strict monotonicity therefore follows that of Proposition 2 verbatim, with the RBC nullcline replaced by $C_A^{Knull}(K) = AK^\alpha - \delta K$. Unique K -stationarity enters in two parts. Locally, it is supplied by the stable-eigenvalue condition. Globally, it is maintained

as a pathwise hypothesis, verifiable by direct inspection of the simulated frozen-regime continuation. The primitive MPC condition (Appendix G, Lemma G.1) ensures the no-side-switching property along the continuation but does not on its own rule out non-steady-state flat steps. Since $\{K_t\}$ is strictly monotone and converges to K_A^{cs} , applying Theorem 1 with $e(\Phi) = K(\Phi)$ yields $K(\Phi) = K(\Phi') \implies \Phi = \Phi'$ for all $\Phi, \Phi' \in \mathcal{M}(\Phi_0, A)$. ■

Proposition 7 makes the K -nullcline *sign-determining*: the graph $C = C_A^{Knull}(K)$ partitions the (K, C) plane into regions with $\Delta K > 0$ and $\Delta K < 0$, via $\Delta K = C_A^{Knull}(K) - C$. Assumption 2 rules out discrete-time jumps across this partition by requiring that the segment connecting (K_t, C_t) and (K_{t+1}, C_{t+1}) does not intersect the nullcline. Together, these two ingredients imply that net investment cannot reverse sign along the conditional saddle. Adding unique K -stationarity (the no-plateau condition) to rule out non-steady-state flat steps, capital dynamics are strictly monotone, and each frozen-regime spell admits an exact one-dimensional representation indexed by K .

Proposition 8 says that, conditional on (Φ_0, A) , the mapping $K : \mathcal{M}(\Phi_0, A) \rightarrow \mathbb{R}$ is one-to-one, so an inverse map $\varphi_{\Phi_0, A}$ exists with $\Phi = \varphi_{\Phi_0, A}(K(\Phi))$ on $\mathcal{M}(\Phi_0, A)$.¹⁸ Any equilibrium object restricted to $\mathcal{M}(\Phi_0, A)$ therefore admits an exact single-valued representation as a function of (K, A) .

To complement this theoretical result, Figure 8 plots the computed conditional saddle paths in the (K, C) phase diagram. The model is solved globally using the repeated transition method, and the dynamics under each frozen aggregate state are simulated for 2,500 periods. Although the true endogenous state of the model is the full distribution Φ rather than K alone, the figure shows that conditional saddle paths are strictly and monotonically ordered in K , providing clear computational support for the sufficiency result.

The tight ergodic alignment in Figure 8 could mask the within-spell structure that Proposition 8 pins down. Figure 9 resolves this by zooming into a narrow K -window: the cloud, line-like at the ergodic scale, is visibly “dusty” once resolved, yet within each individual spell the consecutive (K_t, C_t) points (linked by dashed lines) move monotonically in K . Each spell traces out an exact one-dimensional orbit; ergodic mixing across spells produces the band of points one observes at any given K .

A closer reading of Figure 9 reveals more structure than Proposition 8 alone delivers. The within-spell dashed links are not merely monotone in K but visibly *parallel*, surrounded by the bounded cross-spell dust. This parallelism is the visual signature of *first-order* K -sufficiency: the local slope dC/dK along each conditional saddle is approximately invariant to the starting distributional condition Φ_0 .

¹⁸This is a path-specific notion of sufficiency: for a fixed (Φ_0, A) , K uniquely indexes states on the entire conditional saddle (i.e., globally *along the transition path*).




Figure 8: Conditional saddle paths in [Krusell and Smith \(1998\)](#)

Notes: The figure plots conditional saddle paths for $A = B$ (solid) and $A = G$ (dashed) implied by a canonical heterogeneous household business cycle model ([Krusell and Smith, 1998](#)). The histogram in the background plots the time-series distribution of the aggregate capital stock.




Figure 9: KS98 conditional saddle: zoom-out (left) vs. zoom-in (right)

Notes: Panel (a) shows the full ergodic view; the dashed rectangle marks the zoom window of panel (b). Panel (b) shows a narrow K -window ($\pm 0.15\%$ around the ergodic median); each ergodic observation is plotted as a colored marker (orange $A = B$, blue $A = G$), with consecutive within-spell observations linked by black dashed lines. Each dashed link is monotone in K , illustrating Proposition 8: within a spell, K uniquely indexes the distributional state. The visible cross-spell “dust” is precisely what Proposition 9 bounds.

The two visual features correspond to two orders of coincidence. Strict K -sufficiency would require both zeroth-order coincidence (same C at common K) and first-order coincidence (same slope at common K). The bounded dust in the zoom panel is the (bounded) failure of the former, governed by Proposition 9; the parallelism is the latter holding approximately. Asymptotically this is Proposition 10 at work: near the conditional steady state, all orbits align with a single direction, so their slopes converge to the common value $\psi_A = R_A^{cs} - \lambda_A^s$ regardless of starting point. A formal global version of this slope-coincidence claim would require additional geometric structure and is left to future work.

From spell-wise to global sufficiency Proposition 8 is stated regime-by-regime: K uniquely indexes the conditional saddle $\mathcal{M}(\Phi_0, A)$ under a single frozen regime. The equilibrium path is a concatenation of such one-dimensional segments, each exactly indexed by (K, A) within its own spell. The handoff between spells, however, is governed by the full distributional state Φ_τ at the switch date, and the scalar pair (K_τ, A_τ) need not determine it: different histories can arrive at the same (K_τ, A_τ) carrying different distributions. Spell-wise sufficiency is thus distinct from unconditional global sufficiency, which would require the same K to pin the same Φ regardless of path history. The following proposition bounds this across-spell discrepancy.

The key intuition is as follows. When the economy is observed at a given K , different prior histories may have produced different underlying distributions Φ : a “cloud” of distributions at the same capital level. How tight is this cloud? The answer comes from a synchronization property of contracting random dynamics: when every frozen-regime map contracts, the state today is determined by the regime history alone, up to an error that decays geometrically in the length of history. Two states reached through histories that agree in recent periods must therefore be close, no matter how different their distant pasts. The dispersion in the same- K cloud is thus governed by how far back two histories must be traced before they diverge.

Write $\Gamma_A := \Gamma_{\text{endo}}(\cdot, A)$. The theorem’s hypothesis is collected in the following definition; as discussed below, it is the substantive maintained assumption of this subsection.

Definition 7 (Uniform frozen-regime contraction).

The family $\{\Gamma_A\}_{A \in \{B, G\}}$ is a uniform frozen-regime contraction on $\mathcal{U}$ at rate $\lambda_c \in (0, 1)$ if $\mathcal{U} \subseteq \mathcal{X}$ is closed in the complete metric space $(\mathcal{X}, d)$, forward-invariant under every Γ_A (each Γ_A maps $\mathcal{U}$ into itself), of finite diameter $\text{diam}(\mathcal{U}) := \sup_{\Phi, \Phi' \in \mathcal{U}} d(\Phi, \Phi')$, and

$$d(\Gamma_A \Phi, \Gamma_A \Phi') \leq \lambda_c d(\Phi, \Phi') \quad \text{for all } \Phi, \Phi' \in \mathcal{U} \text{ and all } A.$$

Theorem 2 (Synchronization under frozen-regime contraction).

Let $\{\Gamma_A\}$ be a uniform frozen-regime contraction on $\mathcal{U}$ at rate λ_c (Definition 7). Then:

- (i) (Random fixed point: the history-dependent analogue of a contraction's unique fixed point) For every regime history $\omega = (\dots, A_{t-2}, A_{t-1})$, the long-history limit (the state obtained by applying the history's maps from an ever more distant start)

$$\Phi^*(\omega) := \lim_{T \rightarrow \infty} (\Gamma_{A_{t-1}} \circ \Gamma_{A_{t-2}} \circ \dots \circ \Gamma_{A_{t-T}})(\Phi)$$

exists, is independent of $\Phi \in \mathcal{U}$, and is the unique state in $\mathcal{U}$ consistent with ω ; any trajectory initialized in $\mathcal{U}$ at date $t - T$ satisfies $d(\Phi_t, \Phi^*(\omega)) \leq \lambda_c^T \text{diam}(\mathcal{U})$.

- (ii) (Suffix synchronization) If two histories ω, ω' agree over their most recent T periods, then

$$d(\Phi^*(\omega), \Phi^*(\omega')) \leq \lambda_c^T \text{diam}(\mathcal{U}).$$

Proof.

See Appendix E. ■

Part (i) is the random-fixed-point construction of [Schenk-Hoppé and Schmalfuß \(2001\)](#) specialized to the frozen-regime maps; part (ii) is its consequence for any two states in the same- K cloud. The cloud bound follows by averaging part (ii) over the distribution of the common-suffix length. For the stationary symmetric two-state chain ($\pi_{BB} = \pi_{GG} = \rho_A$), let $q := \rho_A^2 + (1 - \rho_A)^2$ denote the probability that two independent histories agreeing at a date also agree one period earlier.¹⁹

Proposition 9 (Across-spell approximation bound).

Let $\{\Gamma_A\}$ be a uniform frozen-regime contraction on $\mathcal{U}$ at rate λ_c (Definition 7), and let $\Phi, \Phi' \in \mathcal{U}$ be the long-history limits of two regime histories that agree over their most recent $T \geq 1$ periods. Then

$$d(\Phi, \Phi') \leq \lambda_c^T \text{diam}(\mathcal{U}).$$

If, moreover, the two histories are drawn independently from the stationary symmetric chain, conditioned to share the current regime, then $T - 1$ is geometric with success probability $1 - q$ and

$$\mathbb{E}[d(\Phi, \Phi')] \leq \frac{(1 - q) \lambda_c}{1 - q \lambda_c} \text{diam}(\mathcal{U}).$$

¹⁹Stepping backward one period, the two independent histories continue to agree if both persist (ρ_A^2) or both switch ($(1 - \rho_A)^2$); symmetry makes the time-reversed chain (the process describing the history read backward in time) identical to the forward chain. For asymmetric chains, q is replaced by the corresponding backward agreement probability of the time-reversed chain.

Proof.

See Appendix E. ■

The diameter $\text{diam}(\mathcal{U})$ is finite under the borrowing constraint $a' \geq 0$ and compact idiosyncratic support. Unlike a bound stated in terms of the same- K cloud itself, the right-hand side involves only the primitive set $\mathcal{U}$; the proposition is therefore a genuine bound on the cloud rather than a contraction statement about it. The bound vanishes as $\rho_A \rightarrow 1$ (long spells make recent suffixes long) or $\lambda_c \rightarrow 0$ (fast within-spell convergence).

Uniform contraction is a substantive assumption The hypothesis that $\Gamma_{\text{endo}}(\cdot, A)$ contracts uniformly on $\mathcal{U}$ is not innocuous: the map embeds general-equilibrium price feedback, and establishing contraction-type properties for such operators is precisely the difficulty confronted by global existence and uniqueness analyses (Proehl, 2025). Locally, Proposition 10 anchors the contraction rate to the stable eigenvalue $\lambda_A^s \in (0, 1)$; globally, the assumption is maintained, and the sample-based diagnostic below provides an empirical check: a lower-bound estimate, not a certification. In particular, the paper supplies no upper bound on λ_c : even the hypothesis $\lambda_c < 1$ is maintained rather than verified. No tension with the saddle structure arises, however: Γ_{endo} is the equilibrium law of motion after the forward-looking variable has been solved out, so the unstable direction of the (K, C) system is absent from the distributional dynamics.

Ergodic interpretation Theorem 2(i) is what licenses reading Proposition 9 on the ergodic set. Every state observed along a long simulated path is, up to a geometrically small error, the long-history limit of its own realized history. The long-history limit is the state that history alone drives the economy toward, regardless of where it started. Every pair in the same- K cloud is therefore a pair of long-history limits of regime histories. Each such pair inherits the pathwise bound $d(\Phi, \Phi') \leq \lambda_c^T \text{diam}(\mathcal{U})$, with T the realized common-suffix length. This pathwise bound applies at every date (deep within a spell or just after a switch) because the synchronization argument runs over histories rather than over single spells.

The expectation formula, by contrast, is computed under the idealized model of two independent stationary histories, and two caveats attach to it:

- (i) within-path pairs are dependent;
- (ii) conditioning on a common K is a selection.

For these two reasons the closed form serves as a benchmark for the typical scale of the cloud rather than a bound on its realized diameter.

Theorem 2 leaves λ_c unspecified, only requiring that it bounds the per-period contraction on $\mathcal{U}$. The next proposition pins down what λ_c should be near the conditional steady state: it is just the stable eigenvalue of the linearized frozen-regime dynamics, the same object that governs the local saddle slope in Proposition 5. The proposition presupposes the standard regularity of infinite-dimensional linearization (smoothness of the law of motion; a spectral gap, i.e., a local decay rate cleanly separated from all other modes; a stable manifold, i.e., a smooth curve of states containing the tail of the orbit; and agreement of the metric with the norm used by the linearization); Appendix F states these conditions precisely, with plain-language glosses, and the proof is short once they are granted.

Proposition 10 (Local contraction rate on the conditional saddle).

Near the conditional steady state, the frozen-regime dynamics on the conditional saddle contract distances at the geometric rate $|\lambda_A^s|$, where $\lambda_A^s \in (0, 1)$ is the stable eigenvalue of the linearized frozen-regime dynamics; in aggregate-capital units,

$$K_{t+1} - K_A^{cs} = \lambda_A^s (K_t - K_A^{cs}) + o(|K_t - K_A^{cs}|).$$

Proof.

See Appendix F. ■

The proposition is stated locally and along the conditional saddle; the full statement on the distributional state space, its regularity conditions (smoothness of the law of motion near the conditional steady state, a well-defined stable decay rate separated from all other modes, a smooth curve containing the orbit tail, and agreement of the metric with the local norm (distances measured by d match those used by the linearization near the steady state)), and the proof are in Appendix F. Directions orthogonal to the saddle can contract at different rates, so the global contraction rate λ_c on $\mathcal{U}$ in Theorem 2 could in principle exceed $|\lambda_A^s|$. What the proposition does deliver is the right primitive anchor: the same eigenvalue that drives the slope-ratio diagnostic (Proposition 5) disciplines the local contraction entering Theorem 2 and Proposition 9: a single primitive underlies both. A certified uniform upper bound on λ_c over $\mathcal{U}$ is left to future work; the diagnostic below provides a lower-bound estimate.

A sample-based diagnostic for λ_c The global uniform λ_c in Theorem 2 bounds contraction across a cloud of distributional states at the same K reached from different histories, which need not lie on a single conditional saddle. In plain terms, the diagnostic developed below can show that the map contracts no faster than some observed rate; it cannot certify that the map contracts at least that fast for every pair of states, and the latter is what the bound

needs. The across-spell bound uses λ_c as an *upper* bound on the per-step contraction rate, so a certified verification of Proposition 9 requires a certified upper bound on λ_c . Obtaining such an upper bound requires either a Lipschitz constant for Γ_{endo} on $\mathcal{U}$ (model-specific), or a statistical confidence bound layered onto sample estimates (standard but tedious in function spaces). A certified upper bound is beyond the scope of this paper and is left to future work. Instead, I propose a sample-based *diagnostic* that provides a lower bound on λ_c and a reality check against the local anchor $|\lambda_A^s|$.

The diagnostic proceeds in three steps. (i) Draw N pairs of distributional states $(\Phi^{(i)}, \Phi'^{(i)})$, $i = 1, \dots, N$, from the simulated equilibrium on $\mathcal{U}$, sampling the across-history cloud at common K values. (ii) Compute empirical contraction ratios

$$r_i := \frac{d(\Gamma_{\text{endo}}(\Phi^{(i)}, A), \Gamma_{\text{endo}}(\Phi'^{(i)}, A))}{d(\Phi^{(i)}, \Phi'^{(i)})}.$$

(iii) Report $\hat{\lambda}_c^N := \max_i r_i$. By construction, $\hat{\lambda}_c^N \leq \lambda_c$ (the sample max cannot exceed the true supremum), so $\hat{\lambda}_c^N$ is a *sample-based lower bound*, and substituting it for λ_c in the across-spell formula gives an *optimistic* approximation. Proposition 10 supplies the along-saddle anchor $|\lambda_A^s|$ as a reality check. When off-saddle dynamics contract at comparable rates, $|\lambda_A^s|$ and $\hat{\lambda}_c^N$ should agree closely, indirect evidence that the true λ_c is not much larger.

Under the Krusell and Smith (1998) calibration, quarterly regime persistence is $\rho_A = \pi_{AA} = 0.875$ (average 8-quarter spells) and the sample diagnostic yields $\hat{\lambda}_c^N \approx 0.964$, with the local anchor $|\lambda_A^s| \approx 0.963$ agreeing to two decimals. The backward agreement probability is $q = \rho_A^2 + (1 - \rho_A)^2 \approx 0.781$, and substituting $\hat{\lambda}_c^N$ into the benchmark formula of Proposition 9 gives $\mathbb{E}[\lambda_c^T] \approx 0.85$ as the typical contraction factor multiplying $\text{diam}(\mathcal{U})$ under the independent-histories benchmark, an optimistic estimate on both counts. The direct quantitative check, comparing the globally solved K' at each spell boundary against the prediction from tracking only (K, A) , gives an average discrepancy of 2.5×10^{-4} in log terms (1.3% of one standard deviation of $\ln K$). This measures a different object from the theorem's uniform distributional bound above, a mean forecast error in K' rather than a worst-case distance between distributions, so the two are not directly comparable; it separately documents that the across-spell error is small in this calibration.

State reduction versus functional-form approximation The Krusell and Smith (1998) approach is often labeled near-rational (Akerlof and Yellen, 1985), but this conflates two distinct sources of approximation that the conditional-saddle structure separates. Within a given orbit, conditioning on K is exact, since no two states on the same conditional saddle share the same capital level; approximation there concerns the functional form of the law

of motion rather than the choice of coordinate. Across histories, both functional-form and state-reduction errors remain, because different histories can leave different distributions behind the same capital stock, although the direct check finds the latter small in this calibration (Proposition 9). The framework identifies this cross-history approximation rather than assuming it away.

Connection to block recursive equilibrium The conditional sufficiency of K is closely related to *block recursive equilibrium* (BRE), in which aggregate prices and quantities can be determined from a low-dimensional state without tracking the full distribution (Menzio and Shi, 2010, 2011; Schaal, 2017). In the directed-search environments of Menzio and Shi (2011) and Schaal (2017), block recursivity holds exactly by construction: market segmentation severs the feedback from the distribution to individual decision problems. Here, by contrast, the structure is not built into the environment but emerges spell-wise from the saddle geometry. Within each frozen-regime spell, the conditional saddle framework delivers exactly this structure: K pins down all aggregate equilibrium objects, and individual policies are determined given (K, A) . Spell-wise K -sufficiency thus establishes a conditional, path-specific form of block recursivity: along a given conditional saddle $\mathcal{M}(\Phi_0, A)$, aggregate prices, quantities, and the distributional state Φ itself are recoverable from K . This is weaker than conventional block recursivity, which would recover Φ from (K, A) alone; here the recovery also requires the identity of the orbit, which is fixed by the full history that generated the current state, not by recent history alone; different histories can reach the same (K, A) with different distributions, and a shared *recent* history yields only approximate agreement, at the rate bounded in Proposition 9. Whether the equilibrium is globally block recursive depends on the across-spell discrepancy bounded in Proposition 9. Under the Krusell and Smith (1998) calibration, this discrepancy is quantitatively negligible.

Parametric form of the conditional saddle paths Despite the indexing function K , the specific form of the law of motion remains undetermined. The repeated transition method (Lee, 2026a) exploits the recurrence of equilibrium allocations along the conditional saddle path; Theorem 1 and Proposition 8 provide the theoretical foundation for implementing the RTM using K as the conditional sufficient statistic in heterogeneous-agent models.

Verifying monotonicity in practice The sufficiency results above rely on strict monotonicity of K along frozen-regime continuations (Assumption 2 together with unique K -stationarity). How should a researcher verify this in a given model? The following two subsections develop complementary strategies. Section 4.1.1 provides a *perfect-foresight*

diagnostic: one solves the deterministic ($\mu = 0$) transition dynamics and checks whether K converges monotonically; combined with a no-plateau condition and uniform spectral regularity (at every belief weight, a well-defined local decay rate separated from all other modes) along the homotopy, this transports monotonicity from $\mu = 0$ to $\mu = 1$ (rational expectations). Section 4.1.2 provides an independent route from static primitives (bounds on household MPCs and price feedback) that does not require solving any transition path. The concavity of the consumption function under income uncertainty (Carroll and Kimball, 1996) grounds such MPC bounds in theory.

4.1.1 Perfect-foresight diagnostic for monotonicity

This section shows that K -monotonicity under rational expectations can be verified from the deterministic perfect-foresight model, paired with two additional regularity conditions: unique K -stationarity (no non-steady-state flat steps) and uniform spectral regularity along the homotopy. The key insight is that the K -nullcline, the “obstacle” that the path must not cross, is invariant over beliefs (Propositions 3 and 7), so monotonicity can be transported continuously from perfect foresight to rational expectations once the regularity conditions are in place.

Belief family For $\mu \in [0, 1]$, define the belief transition matrix $\Pi(\mu) := (1 - \mu)I + \mu\Pi$, where I is the identity and Π is the true Markov matrix for $A \in \{B, G\}$. At $\mu = 0$, agents assign probability one to the current regime persisting (perfect foresight); at $\mu = 1$, agents hold rational expectations. For each μ , the μ -RCE is the recursive competitive equilibrium under $A' \sim \Pi(\mu)(\cdot|A)$ with the same technology and market structure. Fix an initial state Φ_0 and a regime $A \in \{B, G\}$. For each μ , denote the frozen-regime continuation from Φ_0 under the μ -RCE by $\{(K_t(\mu), C_t(\mu))\}_{t \geq 0}$, with conditional steady state $K_A^{cs}(\mu)$. The same Φ_0 is held fixed across the homotopy: only beliefs vary with μ , not the initial condition.

Under standard regularity, the frozen-regime continuation and conditional steady state are continuous in μ .²⁰ The question is whether K -monotonicity, which may hold at $\mu = 0$ (perfect foresight), survives the passage to $\mu = 1$ (rational expectations). To formalize this, define the *monotonicity set*

$$\mathcal{S} := \{\mu \in [0, 1] : \{K_t(\mu)\}_{t \geq 0} \text{ is strictly monotone along the frozen-}A \text{ continuation}\}.$$

²⁰Since $\Pi(\mu)$ enters the Bellman equation linearly, the recursive equilibrium operator is continuous in μ . The result follows by the parametric contraction mapping theorem (the fixed point of a contraction varies continuously with a parameter) and continuous dependence of the market-clearing fixed point. See Appendix H.

The next result anchors the verification on a single *deterministic* check—monotonic convergence of the perfect-foresight transition path, a textbook property of the Ramsey model—paired with two regularity conditions along the homotopy: a no-plateau condition (unique K -stationarity) and uniform spectral regularity. Under these, monotonicity transports from perfect foresight to rational expectations. The intuition is that the K -nullcline is the boundary between the rising- K and falling- K regions of the phase diagram, and it does not move as beliefs vary from perfect foresight to rational expectations (Propositions 3 and 7). Suppose the perfect-foresight path stays on one side of this immovable boundary. Turning on rational expectations smoothly cannot push the path across it without violating the regularity conditions. To see why, suppose the path did cross. By continuity, the nullcline gap would then have to pass through zero at a non-steady-state capital level (the intermediate-value step). But the no-plateau condition (unique K -stationarity) rules out exactly such a zero.

Definition 8 (Regular belief homotopy).

Fix (Φ_0, A) . The belief homotopy is regular if, over $\mu \in [0, 1]$:

- (a) *(Continuity) the family $\{(\Phi_t(\mu))_{t \geq 0}, \Phi_A^{cs}(\mu)\}$ is jointly continuous in μ ; and*
- (b) *(Uniform spectral regularity) at every μ , the linearized frozen-regime operator at $\Phi_A^{cs}(\mu)$ has a simple (multiplicity-one) stable eigenvalue $\lambda_A^s(\mu) \in (0, 1)$, separated from the rest of the spectrum by a spectral gap bounded away from zero uniformly in μ . In words: at every belief weight, the local decay rate is well defined and stays clearly separated from all other modes.*

Condition (a) is the continuity result of Appendix H; condition (b) imposes the spectral regularity of Proposition 10 (Appendix F) at every μ , uniformly along the homotopy.

Proposition 11 (Perfect-foresight diagnostic).

Fix (Φ_0, A) and suppose: (i) the perfect-foresight ($\mu = 0$) transition path from Φ_0 is strictly K -monotone; (ii) $\Delta K_t = 0$ only at $\Phi_t = \Phi_A^{cs}(\mu)$, for every $\mu \in [0, 1]$; and (iii) the belief homotopy is regular. Then aggregate capital is strictly K -monotone along the rational-expectations ($\mu = 1$) frozen-regime continuation.

Proof.

In words: monotonicity holds at perfect foresight, persists under small changes in the belief weight, and survives passage to limiting belief weights; so there can be no first belief weight at which it fails, and it extends to all of $[0, 1]$. Formally, since the interval $[0, 1]$ is connected—it cannot be split into two separate, non-touching pieces—the only subset that is simultaneously nonempty, open (robust to small changes in μ), and closed (preserved under limits in μ) is $[0, 1]$ itself. It therefore suffices to show $\mathcal{S}$ has these three properties.

Nonempty. $0 \in \mathcal{S}$ by assumption.

Open (monotonicity persists under small changes in the belief weight). Fix $\mu^* \in \mathcal{S}$ and define the nullcline gap $H_t(\mu) := C_t(\mu) - C_A^{Knull}(K_t(\mu))$, so that $\Delta K_t(\mu) = -H_t(\mu)$. At μ^* , strict monotonicity of K implies $H_t(\mu^*)$ has constant nonzero sign for all pre-limit t (every date at which capital is still moving). By Proposition 7, the nullcline C_A^{Knull} is independent of μ . Combined with continuous dependence of $(K_t(\mu), C_t(\mu))$ on μ , for any finite T there exists $\varepsilon > 0$ such that $|\mu - \mu^*| < \varepsilon$ implies $H_t(\mu) H_t(\mu^*) > 0$ for all $t \leq T$.

For the tail: monotone convergence at μ^* implies the stable eigenvalue of the linearized frozen-regime dynamics at $K_A^{cs}(\mu^*)$ is real and lies in $(0, 1)$; otherwise local convergence would be oscillatory. By continuity of the eigenvalue in μ , this property persists for nearby μ . The two steps fit together as follows. Choose T large enough that for every $t > T$ each nearby path has entered the neighborhood where the eigenvalue argument applies, and take the ε associated with that T . The finite-window step then delivers strict monotonicity for $t \leq T$ and the eigenvalue step delivers it for $t > T$, so monotonicity holds at every date and nearby μ remain in $\mathcal{S}$. See Appendix I for details.

Closed (monotonicity survives passage to a limiting belief weight). Let $\mu_n \rightarrow \mu^*$ with $\mu_n \in \mathcal{S}$. By continuity, $H_t(\mu_n) \rightarrow H_t(\mu^*)$ for each t . Since each $H_t(\mu_n)$ has constant nonzero sign, the limit inherits constant weak sign (say $H_t(\mu^*) \geq 0$ for all t), so the K -path is weakly monotone.

It remains to show strict inequality. By the eigenvalue argument from the Open step, $H_t(\mu^*) > 0$ for all sufficiently large t , so any zero of H must occur at a finite t^* with $K_{t^*}(\mu^*) \neq K_A^{cs}(\mu^*)$. $H_{t^*} = 0$ means $\Delta K_{t^*} = 0$: the point (K_{t^*}, C_{t^*}) lies on the distribution-free K -nullcline (Proposition 7) at a non-steady-state capital level. This contradicts *unique K-stationarity*: the property that $\Delta K_t = 0$ along the convergent path only at the conditional steady state. In the representative-agent model, K is Markov under the frozen regime, so $\Delta K_{t^*} = 0$ makes K_{t^*} a fixed point; uniqueness of the conditional steady state (Assumption 1(i)) yields the contradiction directly. In the heterogeneous-agent model, the eigenvalue condition $\lambda_A^s(\mu^*) \in (0, 1)$ establishes unique K -stationarity locally near $K_A^{cs}(\mu^*)$; globally, unique K -stationarity is maintained as part of hypothesis (ii) of the proposition and verified pathwise by direct inspection of the simulated continuation. So $H_t(\mu^*) > 0$ for all t and $\mu^* \in \mathcal{S}$. ■

Strict K -monotonicity is equivalent to the endpoint no-side-switching condition $H_t H_{t+1} > 0$ for every pre-limit t , where $H_t := C_t - C_A^{Knull}(K_t)$ is the nullcline gap. The surrounding conditions divide the remaining work as follows.

- (*Curvature margin*) Combined with the curvature-margin check of Lemma G.2 (Appendix G), the proposition delivers the full line-segment condition of Assumption 2.
- (*Unique K -stationarity*) Condition (ii) is what rules out non-steady-state flat steps. It is automatic in the representative-agent model (uniqueness of the conditional steady state, Assumption 1(i)). In the heterogeneous-agent model, it holds locally via the stable-eigenvalue condition (Proposition 10) and is maintained globally as a pathwise condition, directly verifiable along the simulated continuation.
- (*MPC primitives*) The bounded price-income MPC of Assumption 3 and the path-level finite-difference bound of Assumption 4 (Appendix G), which caps the period-to-period movement of aggregate consumption, enter the primitive inequality of Lemma G.1; when that inequality is verified along the continuation, side-switching of the nullcline gap is ruled out. Even then, however, these conditions do not rule out non-steady-state flat steps.

Discussion Proposition 11 reduces verification to a perfect-foresight monotonicity check, paired with the no-plateau and uniform-regularity conditions in hypotheses (ii) and (iii). The PF check is a textbook property of the Ramsey saddle in the representative-agent case; in the heterogeneous-agent case, monotone convergence of aggregate capital along computed Bewley–Aiyagari transitions under fixed TFP is a robust numerical regularity (Lee, 2026a), though not a general theorem. The proof connects $\mu = 0$ to $\mu = 1$ through a continuous path of equilibria, using nullcline invariance (Proposition 7) as the common “obstacle.” The heterogeneous-agent case requires unique K -stationarity at every $\mu \in [0, 1]$ to rule out non-steady-state flat steps, and uniform spectral regularity (Proposition 10, Appendix F) to control the tail along the homotopy. These are maintained as hypotheses and admit direct numerical verification along the simulated continuation; the MPC-based primitives of Section 4.1.2 additionally rule out endpoint side-switching but do not by themselves establish unique K -stationarity.

The diagnostic extends to any continuous aggregate variable e satisfying unique stationarity ($e(\Phi_{t+1}) = e(\Phi_t)$ only at $\Phi_t = \Phi^{\text{cs}}$), even when the nullcline is not distribution-free. This matters for the extensions in Section 4.1.3. The CRRA-separable labor case makes the nullcline distribution-dependent; at $\mu = 0$, however, the aggregate dynamics reduce to a single deterministic transition path, so monotonicity needs to be checked only along that path.

4.1.2 Primitive conditions for Assumption 2

An independent verification route works directly under rational expectations ($\mu = 1$). It uses static allocation primitives (bounds on household marginal propensities to consume and general-equilibrium price feedback) rather than solving the full stochastic equilibrium. The check is conditional on a realized frozen-regime continuation (or certified envelope bounds on the relevant moments along it, i.e., verified upper and lower bounds that bracket those moments). Verification proceeds in two steps (Appendix G). Lemma G.1 gives a primitive condition on MPCs and GE price feedback that rules out endpoint side-switching of the nullcline gap: $H_t H_{t+1} > 0$ for every t with $K_{t+1} \neq K_t$. Lemma G.2 then upgrades endpoint no-side-switching to the full line-segment no-crossing of Assumption 2 through a curvature-margin check: the nullcline gap at the endpoints must exceed $M(K_{t+1} - K_t)^2/8$, where M bounds the curvature of the nullcline on the traversed interval. Both steps admit direct numerical verification along the simulated continuation.

Appendix G presents the primitive MPC condition as a three-level hierarchy. Level 1 covers log utility with Cobb–Douglas production and bounded MPCs; it gives an explicit pointwise bound at the conditional steady state. Under the [Krusell and Smith \(1998\)](#) calibration, the inequality holds with substantial slack (LHS ≈ 0.04 , threshold 1). Level 2 extends Level 1 numerically to CRRA utility with σ in the empirically relevant range. Level 3 is the perfect-foresight diagnostic of Proposition 11, a numerical fallback when closed-form primitives cannot be verified. The hierarchy has real bite: Appendix G illustrates a non-concave production economy where the primitive condition cannot rule out segment crossing, confirming that Assumption 2 is a substantive restriction.

The hierarchy culminates in a theorem with two complementary parts. Near the conditional steady state, no MPC bound and no pathwise condition are needed: local spectral regularity of the equilibrium law of motion alone delivers exact K -sufficiency. Away from the steady state, the log–Cobb–Douglas structure converts the verification chain (no-side-switching, segment no-crossing, and the relevant moment bounds) into explicit inequalities. Two definitions organize the theorem’s hypotheses, and neither asserts anything. The first names the region on which the inequalities hold: an interval qualifies whenever its two endpoint bounds are satisfied, which can be evaluated directly from calibrated parameters. The second names the pathwise conditions checked on a computed continuation. The theorem then states what passing these checks certifies.

Definition 9 (K -monotonicity interval).

Fix $A \in \{B, G\}$ and constants $\bar{m} \in (0, 1]$, $\bar{Z} > 0$, $\bar{a} > 0$. Define

$$\begin{aligned} B_A(K) &:= \bar{m} \alpha (1 - \alpha) A [K^{\alpha-1} \bar{Z} + K^{\alpha-2} \bar{a}], \\ N_A(K) &:= \alpha A K^{\alpha-1} - \delta, \\ M_A(K) &:= \alpha (1 - \alpha) A K^{\alpha-2}. \end{aligned}$$

An interval $[\underline{K}, \bar{K}]$ with $0 < \underline{K}$ and $\bar{K} \leq (A/\delta)^{1/(1-\alpha)}$ is a K -monotonicity interval with margin $\kappa \in (0, 1)$ if

$$B_A(\underline{K}) - N_A(\bar{K}) \leq 1 - \kappa \quad \text{and} \quad 2f(\bar{K}; A) + (1 - \delta)\bar{K} < \frac{8\kappa}{M_A(\underline{K})}.$$

The definition is read as follows. The constants are supplied by Assumptions 3–4 (Appendix G): $\bar{m}$ is the MPC bound, and $\bar{Z}$, $\bar{a}$ are the path bounds on mean productivity and mean assets. Each function has a transparent role: B_A is the MPC–price-feedback term, N_A the nullcline slope, and M_A the nullcline curvature, while the cap $\bar{K} \leq (A/\delta)^{1/(1-\alpha)}$ keeps the nullcline nonnegative. Because all four objects are monotone in K — B_A , N_A , M_A decrease and $2f + (1 - \delta)K$ increases, by inspection of the Cobb–Douglas forms—the two endpoint inequalities control the corresponding quantities over the whole interval; this is what makes the interval checkable from its endpoints alone. The inequalities do their work in the proof of Theorem 3(b): the first supplies the no-side-switching margin, and the second the curvature margin that upgrades endpoint sign preservation to segment no-crossing.

The second definition collects the pathwise conditions.

Definition 10 (K -admissible continuation).

Fix $A \in \{B, G\}$, a K -monotonicity interval $\mathcal{K} = [\underline{K}, \bar{K}]$, and a neighborhood $\mathcal{V}_A$ of Φ_A^{cs} . A frozen- A continuation is K -admissible on $(\mathcal{K}, \mathcal{V}_A)$ if

- (i) its K -path, together with the straight segments connecting consecutive capital levels, remains in $\mathcal{K}$; and
- (ii) every K -stationary state of the continuation with capital in $\mathcal{K}$ belongs to $\mathcal{V}_A$.

Part (b) of the theorem below is a *verification theorem*: its hypotheses are explicit checks on a computed continuation, namely admissibility conditions (i) and (ii), rather than conditions on primitives, and passing the checks certifies exactness along the verified continuation.

Theorem 3 (Within-spell K -sufficiency under log utility and Cobb–Douglas).

Fix $A \in \{B, G\}$ and a conditional saddle $\mathcal{M}(\Phi_0, A)$ in the heterogeneous-household economy of this section.

- (a) Suppose the regularity conditions of Proposition 10 hold at Φ_A^{cs} with stable eigenvalue $\lambda_A^s \in (0, 1)$. Then there exists a neighborhood $\mathcal{V}_A$ of Φ_A^{cs} such that, along the frozen- A continuation started from any $\Phi'_0 \in \mathcal{M}(\Phi_0, A) \cap \mathcal{V}_A$: (i) $\{K_t\}$ is strictly monotone and converges to K_A^{cs} without crossing it; (ii) $\Delta K_t = 0$ only at $\Phi_t = \Phi_A^{cs}$; and (iii) K is an exact conditional sufficient statistic on $\mathcal{M}(\Phi_0, A) \cap \mathcal{V}_A$.
- (b) Suppose, in addition, that Assumptions 3 and 4 hold, and let $\mathcal{K}$ be a K -monotonicity interval with margin $\kappa \in (0, 1)$. Then along every continuation that is K -admissible on $(\mathcal{K}, \mathcal{V}_A)$, with $\mathcal{V}_A$ as in part (a), Assumption 2 holds, $\{K_t\}$ is strictly monotone until convergence, and K is an exact conditional sufficient statistic on the entire continuation.

Proof.

See Appendix G. ■

Numerical verification under the Krusell and Smith (1998) calibration The ergodic range $[\underline{K}, \overline{K}] = [34.6, 38.5]$ is a K -monotonicity interval with margin $\kappa = 0.95$: $B_A(\underline{K}) - N_A(\overline{K}) \approx 0.04 < 0.05$, and the curvature inequality holds with a margin exceeding two orders of magnitude. The productivity bound $\bar{Z} \leq z_{\max}$ holds a priori because the productivity process has bounded (compact) support; $\bar{a}$ and the ergodic range are read off the computed equilibrium, so the check validates the computed continuation rather than replacing computation. Replacing $\bar{a}$ by the a priori feasibility cap $(A/\delta)^{1/(1-\alpha)} \approx 319$ still yields $B_A(\underline{K}) - N_A(\overline{K}) \approx 0.23 < 1$, so the inequalities can still be verified from calibrated parameters alone.

Three clarifications make precise what the theorem does and does not claim.

1. *Part (a)'s hypotheses are regularity conditions, not primitives.* They are not primitives in the preferences-and-technology sense: they are local spectral regularity conditions on the equilibrium law of motion (a simple—that is, non-repeated—real stable eigenvalue with a smooth curve containing the orbit tail; Appendix F). What part (a) establishes is that *given* this standard regularity, the entire pathwise apparatus—MPC bounds, segment-crossing restrictions, plateau conditions—is unnecessary near the conditional steady state. Note also that proximity in part (a) is measured in the distributional state space. Membership of K_0 in $K(\mathcal{V}_A)$ alone is not sufficient, since distinct distributions can share the same capital.
2. *Part (a) speaks to the late phase of a spell.* It applies once the continuation has entered the (unquantified) linearization neighborhood. Coverage of the entire ergodic range, including spell starts inherited from the other regime, is the role of part (b).

3. *Part (b)'s region is explicit but not assumption-free.* Assumption 4 is a maintained path-level bound whose derivation from the policy-level Assumption 3 remains open (Appendix G). Admissibility condition (ii) of Definition 10 (no plateau outside the local state-space neighborhood) is the single residual condition separating this result from a fully unconditional global theorem. It is stated as a hypothesis rather than buried in the verification chain, and Appendix G discusses why closing it appears to require new machinery for the basin of attraction of the conditional steady state (the set of states from which the frozen-regime dynamics converge to it).

4.1.3 Extensions of the dimension-reduction result

Theorem 1 provides verifiable conditions for dimension reduction: identify the K -nullcline and verify monotonicity, using either the perfect-foresight diagnostic (Proposition 11) or the primitive MPC conditions (Appendix G). This subsection applies these tools to several extensions.

Economies with multiple endogenous states Theorem 1 applies to an economy with a multivariate (distributional) endogenous state. As long as there is an aggregate equilibrium variable e that strictly monotonically converges to the conditional steady-state level, e is a conditional sufficient statistic on each conditional saddle path.

For example, consider the following model that extends the heterogeneous-agent model above by adding endogenous bond holding. The corresponding budget constraint is:

$$c + a' + q(\Phi, A)b' = a(1 + r(\Phi, A)) + b + w(\Phi, A)z \quad (33)$$

where b is bond holding and q is the bond price competitively determined by

$$\int b'(a, z; \Phi, A)d\Phi = 0. \quad (34)$$

The key observation is that Proposition 7 continues to hold: the conditional K -nullcline remains $C_A^{Knull}(K) = AK^\alpha - \delta K$, determined solely by the aggregate resource constraint. Since bonds are in zero net supply, their inclusion does not alter the sign-determining property of the nullcline. Thus, the monotonicity argument of Proposition 8 applies, and K remains a conditional sufficient statistic, provided the no-crossing and unique- K -stationarity conditions of that proposition are verified along the continuation. Lee (2026a) confirmed this prediction computationally: equilibrium allocations are strictly monotonically sorted along K in the globally solved model, consistent with Theorem 1.

Multiple predetermined aggregate states When the economy features multiple predetermined aggregate states (e.g., liquid and illiquid capital in a two-asset HANK model), the conditional saddle is generically higher-dimensional, and Theorem 1 requires a monotone *scalar* coordinate along this higher-dimensional object. Existence of such a coordinate is model-specific and cannot be guaranteed by the K -nullcline argument alone. The perfect-foresight diagnostic (Proposition 11) remains applicable provided one can identify a monotone aggregate variable along the deterministic transition path. Establishing general sufficient conditions for dimension reduction on multi-dimensional conditional saddles is an open question.

Models with endogenous labor supply With GHH preferences $u(c, l_H) = \frac{1}{1-\sigma}(c - \frac{\eta}{1+1/\chi} l_H^{1+1/\chi})^{1-\sigma}$, labor supply depends only on wages and productivity, eliminating wealth effects. The conditional K -nullcline depends on technology parameters and the stationary productivity distribution but not on the wealth distribution (Appendix J), so the monotonicity argument of Theorem 1 applies and K remains a conditional sufficient statistic along each conditional saddle path, again provided strict monotonicity and the no-plateau condition are verified along the continuation.

Under separable CRRA utility $u(c, l_H) = \frac{c^{1-\sigma}}{1-\sigma} - \frac{\eta}{1+1/\chi} l_H^{1+1/\chi}$, individual labor supply depends on consumption through the wealth effect, so the K -nullcline acquires distributional dependence. Strict monotonicity cannot be verified analytically in this case, and establishing it formally under CRRA-separable preferences remains an open question. The perfect-foresight diagnostic of Proposition 11 remains applicable since at $\mu = 0$ the aggregate dynamics follow a single deterministic transition path, and computational experiments under standard calibrations confirm strictly monotone convergence along it (Lee, 2026a). The derivations for both specifications are in Appendix J.

The boundary of dimension reduction The extensions above map where the dimension-reduction logic carries over; the framework is equally informative about where it stops. In the lumpy-investment economy of Khan and Thomas (2008), nullcline invariance fails (wealth effects on labor supply and threshold-driven aggregate investment make the resource constraint distribution-dependent), so the analytical route to K -sufficiency is blocked, and the within-spell monotonicity prediction of Proposition 8 becomes a testable diagnostic rather than a theorem. Appendix M runs the test across three economies: the Khan and Thomas (2008) model in general equilibrium, its partial-equilibrium variant with prices fixed, and the Krusell and Smith (1998) benchmark. The diagnostic is clean, and scalar-state forecasting accurate, in both general-equilibrium economies, rationalizing the finding of Khan

and Thomas (2008) that approximate aggregation survives lumpy investment. Removing the price feedback breaks the monotonicity and the accuracy together: dynamic forecast errors rise thirtyfold even as the one-step R^2 stays at 0.998, echoing Den Haan (2010), and spells flagged by the diagnostic carry the largest errors. Where the diagnostic is clean, scalar-state methods are on solid ground; where it degrades, the cross-sectional distribution carries independent forecasting power (Bachmann et al., 2013).

4.2 Search and matching: comparative statics

In this section, I demonstrate how the conditional saddle framework can be utilized for visualizing comparative statics of stochastic dynamic equilibrium. The conditional saddle framework applies to any recursive competitive equilibrium with a predetermined endogenous state, not only neoclassical growth models. This section applies it to a Diamond–Mortensen–Pissarides search-and-matching economy, where the predetermined state is unemployment. Conditional saddle paths organize the joint dynamics of unemployment and vacancies: the regime-specific vacancy loci slope upward, and the downward-sloping Beveridge-curve pattern in (u, v) data emerges as a loop traced by transitions across them. The slope-asymmetry characterization of Lemma 1 and Proposition 5 acquires a concrete labor-market interpretation: the surplus compression that the fundamental-surplus analysis of Ljungqvist and Sargent (2017) places at the heart of the Shimer puzzle manifests as differential saddle steepness. The baseline model produces conditional saddles whose slope asymmetry generates state dependence in vacancy creation; comparative statics on unemployment benefits reveal a non-monotone relationship with volatility, traceable to convergence of the regime-specific conditional steady states at high benefit levels. Benefit-policy extensions are collected in Appendix K.

The DMP environment Consider a discrete-time stochastic search-and-matching economy in the Diamond–Mortensen–Pissarides tradition.²¹ Aggregate productivity follows a Markov chain $z \in \{G, B, \dots\}$; for exposition I focus on the two extreme states $z \in \{G, B\}$. The aggregate state is $S = [u_{-1}, z]$, where unemployment $u_{-1} = 1 - n_{-1}$ is the sole predetermined endogenous state. Employment evolves according to

$$n = (1 - \lambda) n_{-1} + q(\theta) v, \quad (35)$$

²¹The full model specification and calibration are in Appendix K. I adopt a standard formulation following Shimer (2005) with Nash-bargained wages, exogenous separation, and Cobb–Douglas matching. The computation uses an 11-state Tauchen discretization of log TFP; for expositional simplicity, I present the two extreme states as G and B .

where $\lambda \in (0, 1)$ is the exogenous separation rate, $\theta = v/u$ is labor market tightness (with $u = 1 - n$ the unemployment pool determined jointly with current matching, while the predetermined state is inherited unemployment $u_{-1} = 1 - n_{-1}$), and $q(\theta) = m_0 \theta^{-\xi}$ is the vacancy-filling probability implied by a Cobb–Douglas matching function. Market tightness $\theta(S)$, and hence vacancies $v(S) = \theta(S) \cdot u(S)$, is the jump variable. It is pinned down by the free-entry condition for vacancy posting:

$$\frac{\kappa}{q(S)} = \beta \mathbb{E} \left[\left(\frac{c(S)}{c(S')} \right)^\sigma \left(z' - w(S') + (1 - \lambda) \frac{\kappa}{q(S')} \right) \middle| S \right], \quad (36)$$

where κ is the vacancy posting cost, $w(S)$ is the Nash-bargained wage, and $c(S) = n(S)z - \kappa v(S) + (1 - n(S))b$ is aggregate consumption from the resource constraint. The timing convention is hybrid: matches formed in the current period enter employment and production immediately, while free entry values a vacancy from the following period onward (Appendix K details this convention).

The conditional saddle path $\mathcal{M}(u_0, z)$ is the orbit $\{(u_t, v_t)\}_{t \geq 0}$ generated by (35) and the equilibrium vacancy function $v(\cdot, z)$ under frozen productivity z . Its limit point is $(u_z^{\text{cs}}, v_z^{\text{cs}})$.

Slope asymmetry and the Beveridge curve Figure 10 plots the conditional saddle paths in (u_{-1}, v) space. Both loci slope upward (higher inherited unemployment raises equilibrium vacancies), so state dependence is sign-preserving, as in the RBC benchmark. However, the slopes differ markedly. Under G , matches are profitable and the saddle is steep; under B , the flow surplus is low and the saddle is nearly flat.

The mechanism is the free-entry condition (36). Each additional unemployed worker represents a potential match, but the value of that match depends on productivity. Under high z , the flow surplus $z - w$ is large, so firms respond aggressively to a marginal increase in the unemployment pool, producing a steep vacancy locus. Under low z , the surplus is compressed—with the outside option b fixed, the Nash wage falls less than one-for-one with productivity, squeezing $z - w$ (Hagedorn and Manovskii, 2008; Ljungqvist and Sargent, 2017)—so vacancy posting barely responds, producing a nearly flat locus. The slope asymmetry is the visual signature of the surplus compression that underlies the unemployment-volatility puzzle of Shimer (2005).

This asymmetry generates the Beveridge curve “loop” observed in empirical (u, v) data. With the economy near the G steady state, a negative shock switches it onto the flat B saddle. Vacancies collapse on impact and barely adjust as unemployment accumulates along the flat path. Upon recovery, the economy jumps back onto the steep G saddle at high inherited u_{-1} , driving a sharp vacancy spike and rapid job creation. The counterclockwise

rotation through recessions and recoveries is thus the across-saddle jump dynamics of the stochastic equilibrium.




Figure 10: Conditional saddle paths in the DMP model, constant b

Notes: Conditional saddle paths in inherited-unemployment–vacancy (u_{-1}, v) space under frozen productivity $z \in \{G, B\}$. Stars mark conditional steady states; arrows indicate the direction of convergence. Extensions with stochastic and unemployment-dependent benefits are in Appendix K.

Comparative statics of the benefit level Ergodic unemployment volatility is hump-shaped in the unemployment–benefit level b , with an interior peak. The operative geometric diagnostic is the gap between regime-specific conditional steady states, $u_B^{\text{cs}}(b) - u_G^{\text{cs}}(b)$, not slope asymmetry alone. At high b , u_B^{cs} saturates while u_G^{cs} keeps rising; the steady-state gap collapses, and the synchronization contraction of Theorem 2(i) anchors each regime’s ergodic states near its conditional steady state, so the collapsing gap compresses the range over which unemployment can fluctuate between regimes.²² The mechanism is the asymmetric surplus compression implicit in the Nash wage: the same fundamental-surplus channel through which Ljungqvist and Sargent (2017) show the benefit level governs amplification across DMP variants; the saddle geometry adds the steady-state-gap diagnostic that makes the non-monotonicity visible. Robustness to wage-setting rules and matching elasticities is documented in Appendix K.

How does the benefit level b reshape the conditional saddle geometry, and what does this imply for unemployment volatility? The Nash wage $w = (1 - \eta)b + \eta(z + \kappa\theta)$ is increasing in b , so raising b compresses the match surplus $z - w$. The compression is proportionally more severe in the B regime, where $z - b$ is already small. Figure 11a reports the conditional saddle paths from solving the DMP model for $b \in \{0.10, 0.40, 0.70, 0.90\}$. At every benefit level, the G saddle is steeper than the B saddle; as b rises, the deepening surplus compression

²²The theorem’s uniform-contraction hypothesis is maintained for the scalar-state DMP frozen-regime map, where it is directly checkable.

flattens both loci, leaving the B saddle nearly flat, while pushing the conditional steady states toward higher unemployment. One might expect the deepening compression to raise volatility monotonically, but the effect is *hump-shaped* (Figure 11b): volatility peaks at $b \approx 0.55$ ($\text{std}(\log u) \approx 7.6\%$, with a nearly flat top through $b \approx 0.60$) and then *declines*. The exact peak location is model- and calibration-specific. Two clarifications guard against misreading. First, the result is a ceteris-paribus comparative static: the amplification in Hagedorn and Manovskii (2008) at $b \approx 0.955$ operates through their joint recalibration of the bargaining weight and the vacancy-posting cost, not through the benefit level alone, so the hump here neither contradicts nor overturns their calibration. Second, holding all other parameters fixed, the saddle geometry explains why b alone is hump-shaped: at high b , conditional steady states converge despite large slope gaps, compressing the range over which unemployment can fluctuate.²³

The saddle framework diagnoses this hump by revealing that what governs volatility is the *gap between conditional steady states*, not slope asymmetry alone. The right subplot of Figure 11b plots conditional steady-state unemployment under frozen good (u_G^{cs}) and bad (u_B^{cs}) TFP for each b . As b rises, u_G^{cs} increases steadily while u_B^{cs} rises faster, widening the gap from under 1 p.p. at $b = 0.05$ to about 2.2 p.p. at $b \approx 0.60$. Beyond that point, u_B^{cs} saturates. Agents expect future TFP recovery, so vacancy posting persists even under bad productivity, capping the bad-state conditional steady state near 6.8%. Meanwhile, u_G^{cs} continues to rise as higher b compresses the surplus even under good TFP. The two conditional steady states therefore *converge* at high b : the saddles remain differentially tilted but sit close together, leaving little room for unemployment to fluctuate between regimes. Both the gap $u_B^{\text{cs}} - u_G^{\text{cs}}$ and simulated volatility are hump-shaped in b , peaking near the same value. The conditional steady-state gap, not the slope difference, is the operative geometric diagnostic for unemployment volatility.

Sensitivity to wage setting and matching elasticity The hump-shaped relationship between b and unemployment volatility is robust across wage-setting rules and for matching elasticities up to the baseline value (Figure K.2). Under Nash bargaining with matching elasticities $\xi \in \{0.40, 0.50, 0.635, 0.80\}$, the hump persists for $\xi \leq 0.635$, with the peak shifting leftward as ξ increases: from $b \approx 0.80$ at $\xi = 0.40$ to $b \approx 0.55$ at the baseline $\xi = 0.635$. Along this sweep the Hosios condition $\eta = \xi$ is maintained, so raising ξ both dampens the tightness response to surplus differences (the free-entry elasticity of tightness with respect to the match value is $1/\xi$) and raises the worker’s bargaining weight, compressing

²³In the parameterization here, simulated volatility at $b = 0.955$ is about 4.2%, compared with 6.8% at the baseline $b = 0.40$. The saddle diagnostic identifies steady-state convergence as the mechanism.




Figure 11: Comparative statics of b

Notes: Panel (a): conditional saddle paths in (u_{-1}, v) space for $b \in \{0.10, 0.40, 0.70, 0.90\}$; stars mark conditional steady states. Panel (b), left: $\text{std}(\log u)$ over 19 values of b ; the star marks the peak ($b \approx 0.55$), the square marks the baseline ($b = 0.40$), and the dash-dotted vertical line marks the [Hagedorn and Manovskii \(2008\)](#) replacement rate ($b = 0.955$). Panel (b), right: conditional steady-state unemployment under frozen good (u_G^{cs} , green) and bad (u_B^{cs} , red) TFP, extracted from the stochastic policy functions. Both the gap $u_B^{\text{cs}} - u_G^{\text{cs}}$ and simulated volatility are hump-shaped in b .

the firm's surplus share; the conditional steady states therefore converge at lower benefit levels, moving the volatility peak left. At $\xi = 0.80$, volatility is nearly flat in b : the matching function is sufficiently elastic that conditional steady states converge over the entire range. Under *present-wage* bargaining, $w = (1 - \eta)b + \eta z$, which eliminates the congestion externality feedback from tightness to wages, the hump shape persists with peak volatility at $b \approx 0.93$. The non-monotone pattern is driven by surplus geometry rather than the specific wage-

feedback channel.

5 Concluding remarks

Conditional saddle paths, equilibrium trajectories under frozen exogenous regimes, bring the geometric clarity of deterministic phase diagrams to stochastic environments. The slopes of these objects encode state dependence: differential tilt across saddles is the geometric signature of asymmetric shock responses. Their monotonicity delivers dimension reduction: a strictly monotone aggregate variable exactly indexes the equilibrium state within each frozen-regime spell. Persistent spells and exact within-spell sufficiency, together with a synchronization property under a maintained contraction condition (the economy forgets its regime history at a geometric rate), identify the mechanism by which the conditional result would translate into near-global accuracy. The bound this mechanism delivers is loose in the calibration examined, and the near-exact accuracy is documented directly rather than through the bound; together they give a structural account of why scalar-state methods work.

The paper illustrates the framework through three applications. The canonical RBC model establishes the basic geometry and reveals that state-dependent response shapes arise from distance along the saddle even when slopes are nearly equal. The heterogeneous-household model of [Krusell and Smith \(1998\)](#) demonstrates that strict monotonicity along conditional saddles delivers exact within-spell dimension reduction: locally under standard regularity, and globally on an explicit region whose conditions can be checked from calibrated parameters, with one residual assumption stated rather than hidden; the across-spell error is bounded, under the maintained uniform-contraction condition, by how quickly the economy forgets its regime history. A non-concave production illustration shows the no-crossing condition is a genuine geometric restriction, and a lumpy-investment experiment confirms the diagnostic tracks where scalar-state accuracy degrades. The search-and-matching application shows the framework's reach beyond neoclassical growth: the Beveridge-curve loop emerges from transitions across regime-specific saddle loci, surplus compression under low productivity manifests as differential saddle steepness, and the saddle geometry provides a direct diagnostic for the non-monotone relationship between benefit generosity and unemployment volatility.

More broadly, conditional saddles offer a complementary lens on stochastic models. They provide a language for interpreting nonlinear dynamics, for assessing when scalar-state reductions are exact on the equilibrium path rather than merely accurate, and for characterizing across-path variation as a structural feature of the geometry. A promising direction for future work is to use these geometric objects to sharpen empirical restrictions on state dependence

and to discipline the design of state-contingent stabilization policies.

References

- ABEL, A. B. AND J. C. EBERLY (1996): “Optimal Investment with Costly Reversibility,” *The Review of Economic Studies*, 63, 581–593.
- ACHDOU, Y., J. HAN, J.-M. LASRY, P.-L. LIONS, AND B. MOLL (2021): “Income and Wealth Distribution in Macroeconomics: A Continuous-Time Approach,” *The Review of Economic Studies*, 89, 45–86.
- AHN, S., G. KAPLAN, B. MOLL, T. WINBERRY, AND C. WOLF (2018): “When Inequality Matters for Macro and Macro Matters for Inequality,” *NBER Macroeconomics Annual*, 32, 1–75.
- AIYAGARI, S. R. (1994): “Uninsured Idiosyncratic Risk and Aggregate Saving,” *The Quarterly Journal of Economics*, 109, 659–684.
- AKERLOF, G. A. AND J. L. YELLEN (1985): “Can Small Deviations from Rationality Make Significant Differences to Economic Equilibria?” *American Economic Review*, 75, 708–720.
- ANDREASEN, M. M., J. FERNÁNDEZ-VILLAYERDE, AND J. F. RUBIO-RAMÍREZ (2017): “The Pruned State-Space System for Non-Linear DSGE Models: Theory and Empirical Applications,” *The Review of Economic Studies*, 85, 1–49.
- ARNOLD, L. (1998): *Random Dynamical Systems*, Springer Monographs in Mathematics, Berlin, Heidelberg: Springer Berlin Heidelberg.
- AUCLERT, A., B. BARDÓCZY, M. ROGNLIE, AND L. STRAUB (2021): “Using the Sequence-Space Jacobian to Solve and Estimate Heterogeneous-Agent Models,” *Econometrica*, 89, 2375–2408.
- AUERBACH, A. J. AND Y. GORODNICHENKO (2012): “Measuring the Output Responses to Fiscal Policy,” *American Economic Journal: Economic Policy*, 4, 1–27.
- AZINOVIC, M., L. GAEGAUF, AND S. SCHEIDEGGER (2022): “DEEP EQUILIBRIUM NETS,” *International Economic Review*, 63, 1471–1525.
- AZINOVIC-YANG, M. AND J. ŽEMLIČKA (2025): “Deep Learning in the Sequence Space,” Working paper, arXiv:2509.13623.
- BACHMANN, R., R. J. CABALLERO, AND E. M. R. A. ENGEL (2013): “Aggregate Implications of Lumpy Investment: New Evidence and a DSGE Model,” *American Economic Journal: Macroeconomics*, 5, 29–67.
- BALEY, I. AND A. BLANCO (2019): “Firm Uncertainty Cycles and the Propagation of Nominal Shocks,” *American Economic Journal: Macroeconomics*, 11, 276–337.

- BARRO, R. J. AND X. SALA-I-MARTIN (1992): “Convergence,” *Journal of Political Economy*, 100, 223–251.
- BASU, S. AND B. BUNDICK (2017): “Uncertainty Shocks in a Model of Effective Demand,” *Econometrica*, 85, 937–958.
- BAYER, C. AND R. LUETTICKE (2020): “Solving discrete time heterogeneous agent models with aggregate risk and many idiosyncratic states by perturbation,” *Quantitative Economics*, 11, 1253–1288.
- BERGER, D., K. MILBRADT, F. TOURRE, AND J. VAVRA (2021): “Mortgage Prepayment and Path-Dependent Effects of Monetary Policy,” *American Economic Review*, 111, 2829–78.
- BERGER, D. AND J. VAVRA (2015): “Consumption Dynamics During Recessions: Consumption Dynamics During Recessions,” *Econometrica*, 83, 101–154.
- BEWLEY, T. (1986): “Stationary Monetary Equilibrium with a Continuum of Independently Fluctuating Consumers,” *Contributions to Mathematical Economics in Honor of Gérard Debreu*, 79–102.
- BHANDARI, A., T. BOURANY, D. EVANS, AND M. GOLOSOV (2023): “A Perturbational Approach for Approximating Heterogeneous Agent Models,” NBER Working Paper 31744.
- BILAL, A. (2023): “Solving Heterogeneous Agent Models with the Master Equation,” Working Paper 31103, National Bureau of Economic Research.
- BLANCHARD, O. J. AND C. M. KAHN (1980): “The Solution of Linear Difference Models under Rational Expectations,” *Econometrica*, 48, 1305–1311.
- BLOOM, N., M. FLOETOTTO, N. JAIMOVICH, I. SAPORTA-EKSTEN, AND S. J. TERRY (2018): “Really Uncertain Business Cycles,” *Econometrica*, 86, 1031–1065.
- BOPPART, T., P. KRUSELL, AND K. MITMAN (2018): “Exploiting MIT shocks in heterogeneous-agent economies: the impulse response as a numerical derivative,” *Journal of Economic Dynamics and Control*, 89, 68–92.
- BOURANY, T. (2018): “Wealth Distribution over the Business Cycle: A Mean-Field Game with Common Noise,” Master’s thesis, UPMC-Sorbonne Université, supervised by Yves Achdou, Paris-Diderot University (LJLL).
- BRUMM, J. AND S. SCHEIDEGGER (2017): “Using Adaptive Sparse Grids to Solve High-Dimensional Dynamic Models,” *Econometrica*, 85, 1575–1612.
- BUERA, F. J., J. KABOSKI, M. MESTIERI, AND D. G. O’CONNOR (2024): “The Stable Transformation Path,” Working paper.
- CAO, D. (2020): “Recursive equilibrium in Krusell and Smith (1998),” *Journal of Economic Theory*, 186, 104978.

- CARROLL, C. D. (1997): “Buffer-Stock Saving and the Life Cycle/Permanent Income Hypothesis,” *The Quarterly Journal of Economics*, 112, 1–55.
- CARROLL, C. D. AND M. S. KIMBALL (1996): “On the Concavity of the Consumption Function,” *Econometrica*, 64, 981–992.
- CASS, D. (1965): “Optimum Growth in an Aggregative Model of Capital Accumulation,” *The Review of Economic Studies*, 32, 233–240.
- CHILDERS, R. L. (2018): “Solution of Rational Expectations Models with Function Valued States,” Working paper.
- COOPER, R. W. AND J. C. HALTIWANGER (2006): “On the Nature of Capital Adjustment Costs,” *The Review of Economic Studies*, 73, 611–633.
- DEATON, A. (1991): “Saving and Liquidity Constraints,” *Econometrica*, 59, 1221–1248.
- DECHERT, W. D. AND K. NISHIMURA (1983): “A Complete Characterization of Optimal Growth Paths in an Aggregated Model with a Non-Concave Production Function,” *Journal of Economic Theory*, 31, 332–354.
- DEN HAAN, W. J. (1996): “Heterogeneity, Aggregate Uncertainty, and the Short-Term Interest Rate,” *Journal of Business & Economic Statistics*, 14, 399–411.
- (2010): “Assessing the accuracy of the aggregate law of motion in models with heterogeneous agents,” *Journal of Economic Dynamics and Control*, 34, 79–99.
- DEN HAAN, W. J. AND A. MARCET (1990): “Solving the Stochastic Growth Model by Parameterizing Expectations,” *Journal of Business & Economic Statistics*, 8, 31–34.
- DEN HAAN, W. J. AND P. RENDAHL (2010): “Solving the incomplete markets model with aggregate uncertainty using explicit aggregation,” *Journal of Economic Dynamics and Control*, 34, 69–78.
- FERNÁNDEZ-VILLAYERDE, J., S. HURTADO, AND G. NUÑO (2023): “Financial Frictions and the Wealth Distribution,” *Econometrica*, 91, 869–901.
- GALI, J. (2008): *Monetary Policy, Inflation, and the Business Cycle: An Introduction to the New Keynesian Framework and Its Applications*, Princeton University Press.
- GORMAN, W. M. (1961): “On a Class of Preference Fields,” *Metroeconomica*, 13, 53–56.
- HAGEDORN, M. AND I. MANOVSKII (2008): “The Cyclical Behavior of Equilibrium Unemployment and Vacancies Revisited,” *American Economic Review*, 98, 1692–1706.
- HAN, J., Y. YANG, AND W. E (2026): “DeepHAM: A Global Solution Method for Heterogeneous Agent Models with Aggregate Shocks,” *Quantitative Economics*, 17, 297–341.

- HUGGETT, M. (1993): “The Risk-Free Rate in Heterogeneous-Agent Incomplete-Insurance Economies,” *Journal of Economic Dynamics and Control*, 17, 953–969.
- JUDD, K. L., L. MALIAR, AND S. MALIAR (2011): “Numerically stable and accurate stochastic simulation approaches for solving dynamic economic models: Approaches for solving dynamic models,” *Quantitative Economics*, 2, 173–210.
- KAMIHIGASHI, T. (2003): “Necessity of transversality conditions for stochastic problems,” *Journal of Economic Theory*, 109, 140–149.
- (2005): “Necessity of the transversality condition for stochastic models with bounded or CRRA utility,” *Journal of Economic Dynamics and Control*, 29, 1313–1329.
- KAPLAN, G., B. MOLL, AND G. L. VIOLANTE (2018): “Monetary Policy According to HANK,” *American Economic Review*, 108, 697–743.
- KAPLAN, G. AND G. L. VIOLANTE (2014): “A Model of the Consumption Response to Fiscal Stimulus Payments,” *Econometrica*, 82, 1199–1239.
- KHAN, A. AND J. K. THOMAS (2008): “Idiosyncratic Shocks and the Role of Non-convexities in Plant and Aggregate Investment Dynamics,” *Econometrica*, 76, 395–436.
- KOOP, G., M. PESARAN, AND S. M. POTTER (1996): “Impulse response analysis in nonlinear multivariate models,” *Journal of Econometrics*, 74, 119–147.
- KOOPMANS, T. (1963): “On the Concept of Optimal Economic Growth,” Cowles Foundation Discussion Papers 163, Cowles Foundation for Research in Economics, Yale University.
- KRUSELL, P. AND A. A. SMITH, JR. (1998): “Income and Wealth Heterogeneity in the Macroeconomy,” *Journal of Political Economy*, 106, 867–896.
- LEE, H. (2026a): “A Global Dynamic Nonlinear Solution Framework and the Repeated Transition Method,” Working paper.
- (2026b): “Striking While the Iron Is Cold: Fragility after a Surge of Lumpy Investments,” Working paper.
- LEVINTAL, O. (2018): “Taylor Projection: A New Solution Method for Dynamic General Equilibrium Models,” *International Economic Review*, 59, 1345–1373.
- LJUNGQVIST, L. AND T. J. SARGENT (2017): “The Fundamental Surplus,” *American Economic Review*, 107, 2630–2665.
- MALIAR, L. AND S. MALIAR (2015): “Merging simulation and projection approaches to solve high-dimensional problems with an application to a new Keynesian model,” *Quantitative Economics*, 6, 1–47.
- MALIAR, L., S. MALIAR, AND P. WINANT (2021): “Deep learning for solving dynamic economic models,” *Journal of Monetary Economics*, 122, 76–101.

- MARCET, A. (1988): “Solving nonlinear stochastic models by parameterizing expectations,” Manuscript. Pittsburgh: Carnegie Mellon Univ.
- MELCANGI, D. (2024): “Firms’ precautionary savings and employment during a credit crisis,” *American Economic Journal: Macroeconomics*, 16, 356–386.
- MENZIO, G. AND S. SHI (2010): “Block Recursive Equilibria for Stochastic Models of Search on the Job,” *Journal of Economic Theory*, 145, 1453–1494.
- (2011): “Efficient Search on the Job and the Business Cycle,” *Journal of Political Economy*, 119, 468–510.
- MIAO, J. (2006): “Competitive equilibria of economies with a continuum of consumers and aggregate shocks,” *Journal of Economic Theory*, 128, 274–298.
- NEFTCI, S. N. (1984): “Are Economic Time Series Asymmetric over the Business Cycle?” *Journal of Political Economy*, 92, 307–328.
- PETROSKY-NADEAU, N., L. ZHANG, AND L.-A. KUEHN (2018): “Endogenous Disasters,” *American Economic Review*, 108, 2212–45.
- PIZZINELLI, C., K. THEODORIDIS, AND F. ZANETTI (2020): “STATE DEPENDENCE IN LABOR MARKET FLUCTUATIONS,” *International Economic Review*, 61, 1027–1072.
- PROEHL, E. (2024): “Approximating Equilibria with Ex-Post Heterogeneity and Aggregate Risk,” Working paper.
- (2025): “Existence and Uniqueness of Recursive Equilibria with Aggregate and Idiosyncratic Risk,” Working paper.
- RAMSEY, F. P. (1928): “A Mathematical Theory of Saving,” *The Economic Journal*, 38, 543–559.
- REDDY, P. B., J. M. SCHUMACHER, AND J. C. ENGWERDA (2020): “A Bendixson Criterion for Optimal Control of a Class of Hybrid Systems,” *SIAM Journal on Control and Optimization*, 58, 2219–2252.
- REITER, M. (2010): “Solving the incomplete markets model with aggregate uncertainty by backward induction,” *Journal of Economic Dynamics and Control*, 34, 28–35.
- ROGNLIE, M., A. SHLEIFER, AND A. SIMSEK (2018): “Investment Hangover and the Great Recession,” *American Economic Journal: Macroeconomics*, 10, 113–153.
- SCHAAL, E. (2017): “Uncertainty and Unemployment,” *Econometrica*, 85, 1675–1721.
- SCHENK-HOPPÉ, K. R. (1998): “Random attractors—general properties, existence and applications to stochastic bifurcation theory,” *Discrete and Continuous Dynamical Systems*, 4, 99–130.

- (2001): “Random Dynamical Systems in Economics,” *Stochastics and Dynamics*, 1, 63–83.
- SCHENK-HOPPÉ, K. R. AND B. SCHMALFUSS (2001): “Random fixed points in a stochastic Solow growth model,” *Journal of Mathematical Economics*, 36, 19–30.
- SHIMER, R. (2005): “The Cyclical Behavior of Equilibrium Unemployment and Vacancies,” *American Economic Review*, 95, 25–49.
- SKIBA, A. K. (1978): “Optimal Growth with a Convex–Concave Production Function,” *Econometrica*, 46, 527–539.
- SOLOW, R. M. (1956): “A Contribution to the Theory of Economic Growth,” *The Quarterly Journal of Economics*, 70, 65–94.
- SWAN, T. W. (1956): “ECONOMIC GROWTH and CAPITAL ACCUMULATION,” *Economic Record*, 32, 334–361.
- TENREYRO, S. AND G. THWAITES (2016): “Pushing on a String: US Monetary Policy Is Less Powerful in Recessions,” *American Economic Journal: Macroeconomics*, 8, 43–74.
- VAVRA, J. (2014): “INFLATION DYNAMICS AND TIME-VARYING VOLATILITY: NEW EVIDENCE AND AN SS INTERPRETATION,” *The Quarterly Journal of Economics*, 129, 215–258.
- WAGENER, F. (2003): “Skiba Points and Heteroclinic Bifurcations, with Applications to the Shallow Lake System,” *Journal of Economic Dynamics and Control*, 27, 1533–1561.
- WALSH, K. J. AND E. YOUNG (2024): “Equilibrium Multiplicity in Aiyagari and Krusell-Smith,” Working paper.
- WINBERRY, T. (2018): “A method for solving and estimating heterogeneous agent macro models,” *Quantitative Economics*, 9, 1123–1151.
- (2021): “Lumpy Investment, Business Cycles, and Stimulus Policy,” *American Economic Review*, 111, 364–396.
- YANNAKOPOULOS, A. N. (2011): “Stochastic Saddle Paths and Economic Theory,” in *Dynamics, Games and Science II*, ed. by M. M. Peixoto, A. A. Pinto, and D. A. Rand, Berlin, Heidelberg: Springer Berlin Heidelberg, 735–752.
- ZHANG, L. (2005): “The Value Premium,” *The Journal of Finance*, 60, 67–103.